





Introduction to Kafka Security


Data Flow in Kafka



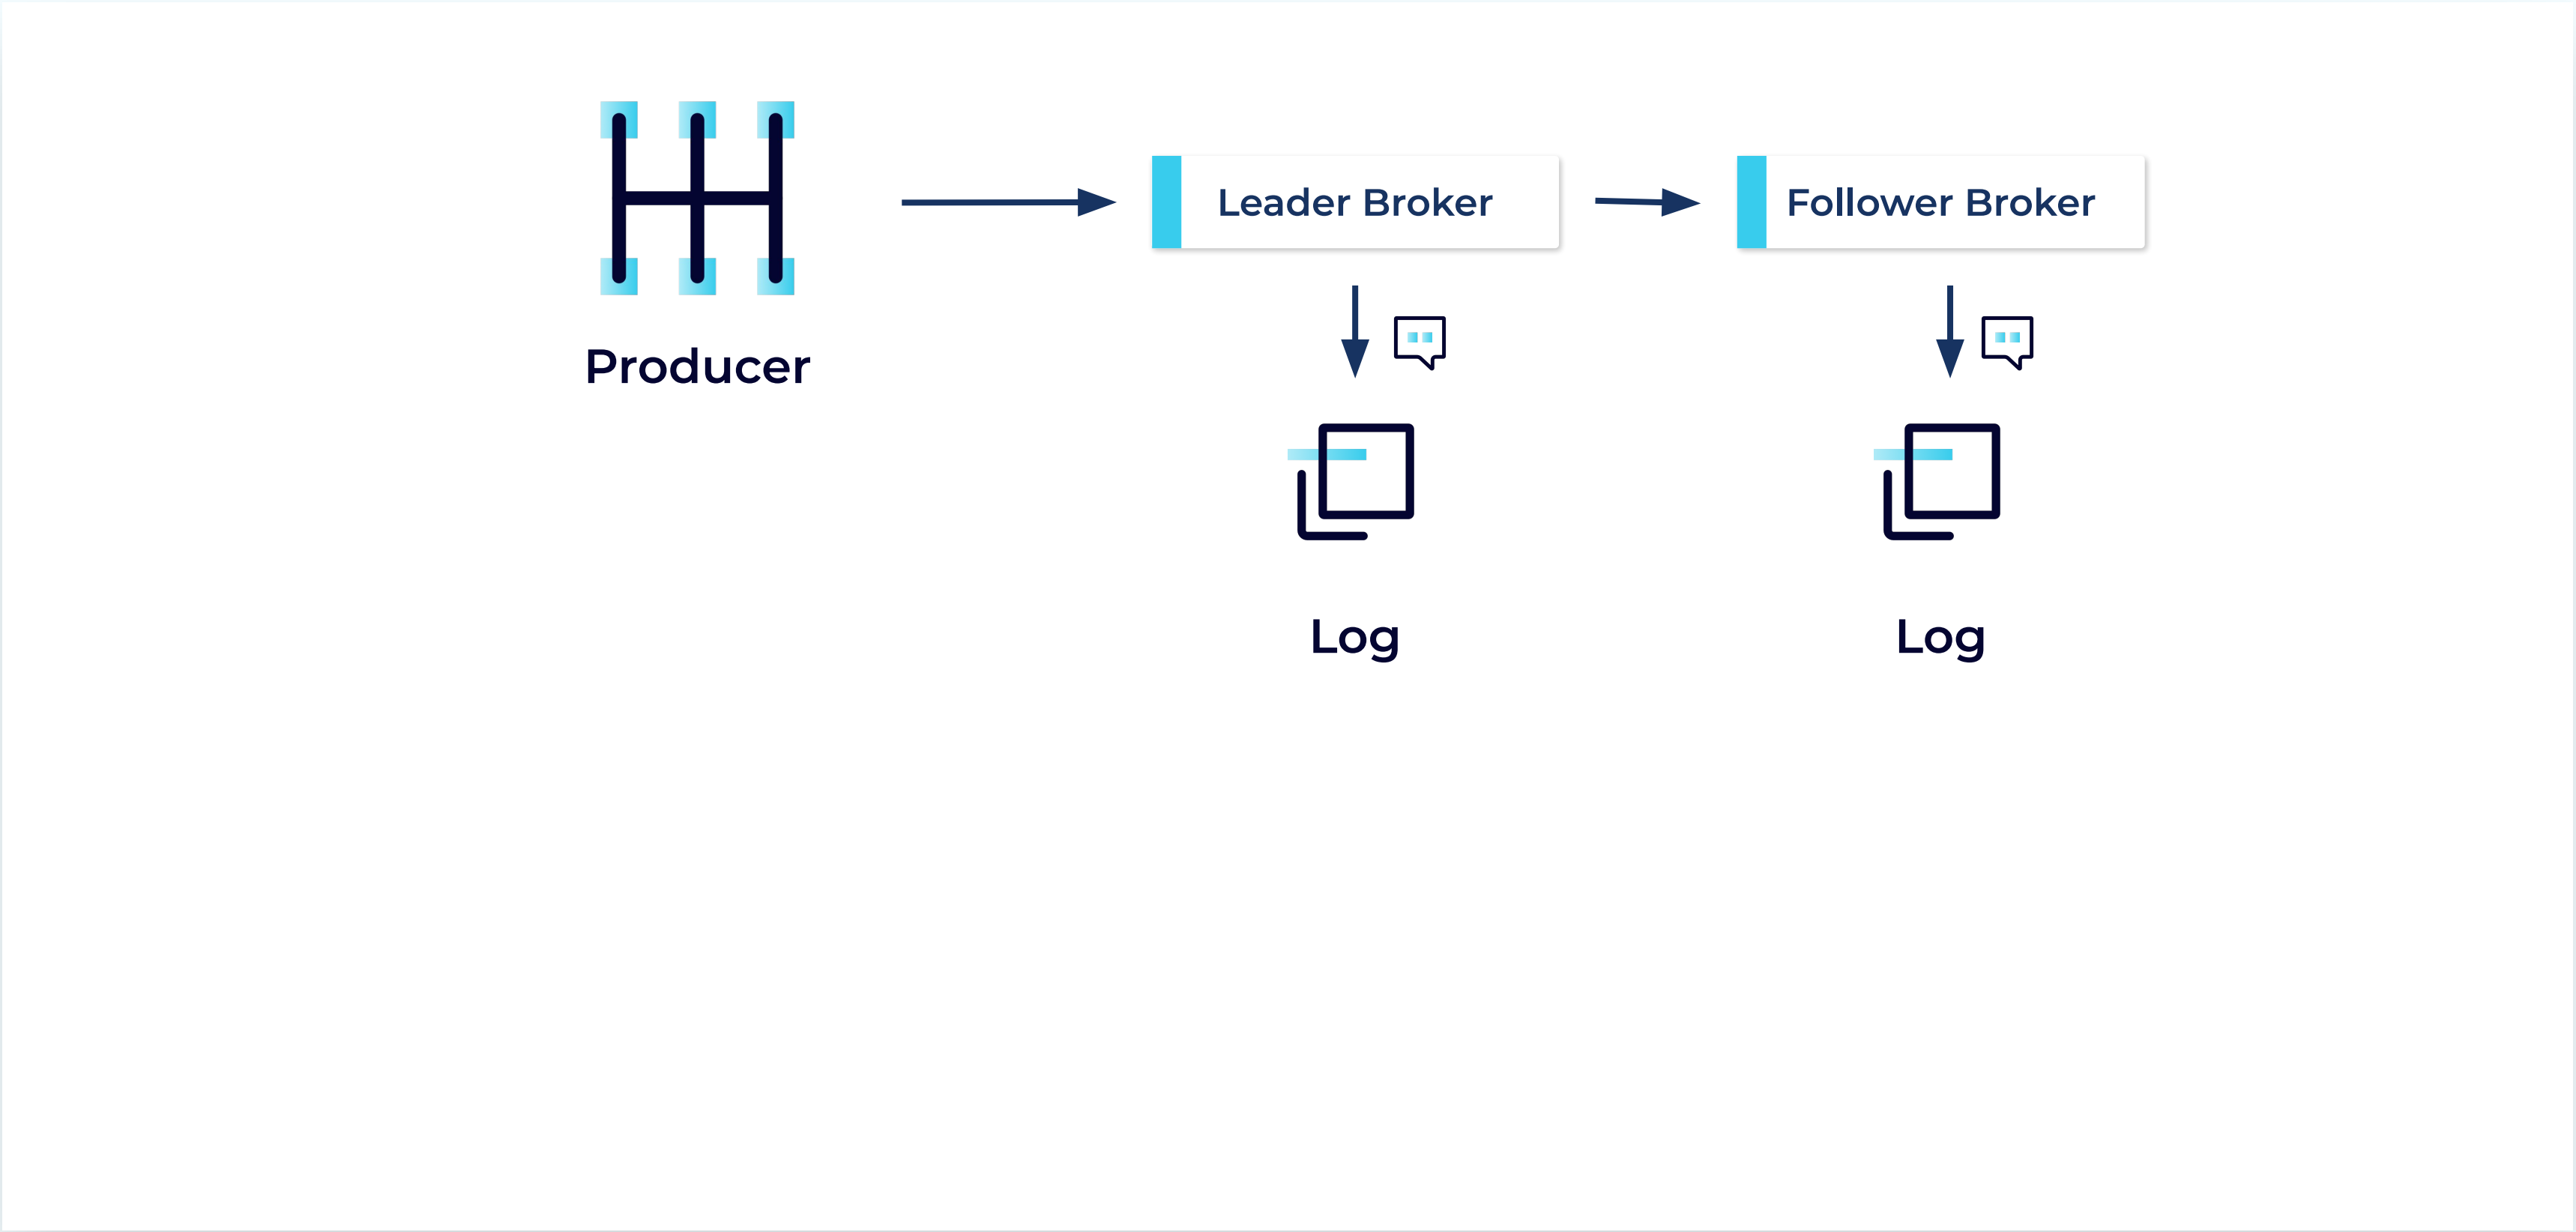
Data Flow in Kafka



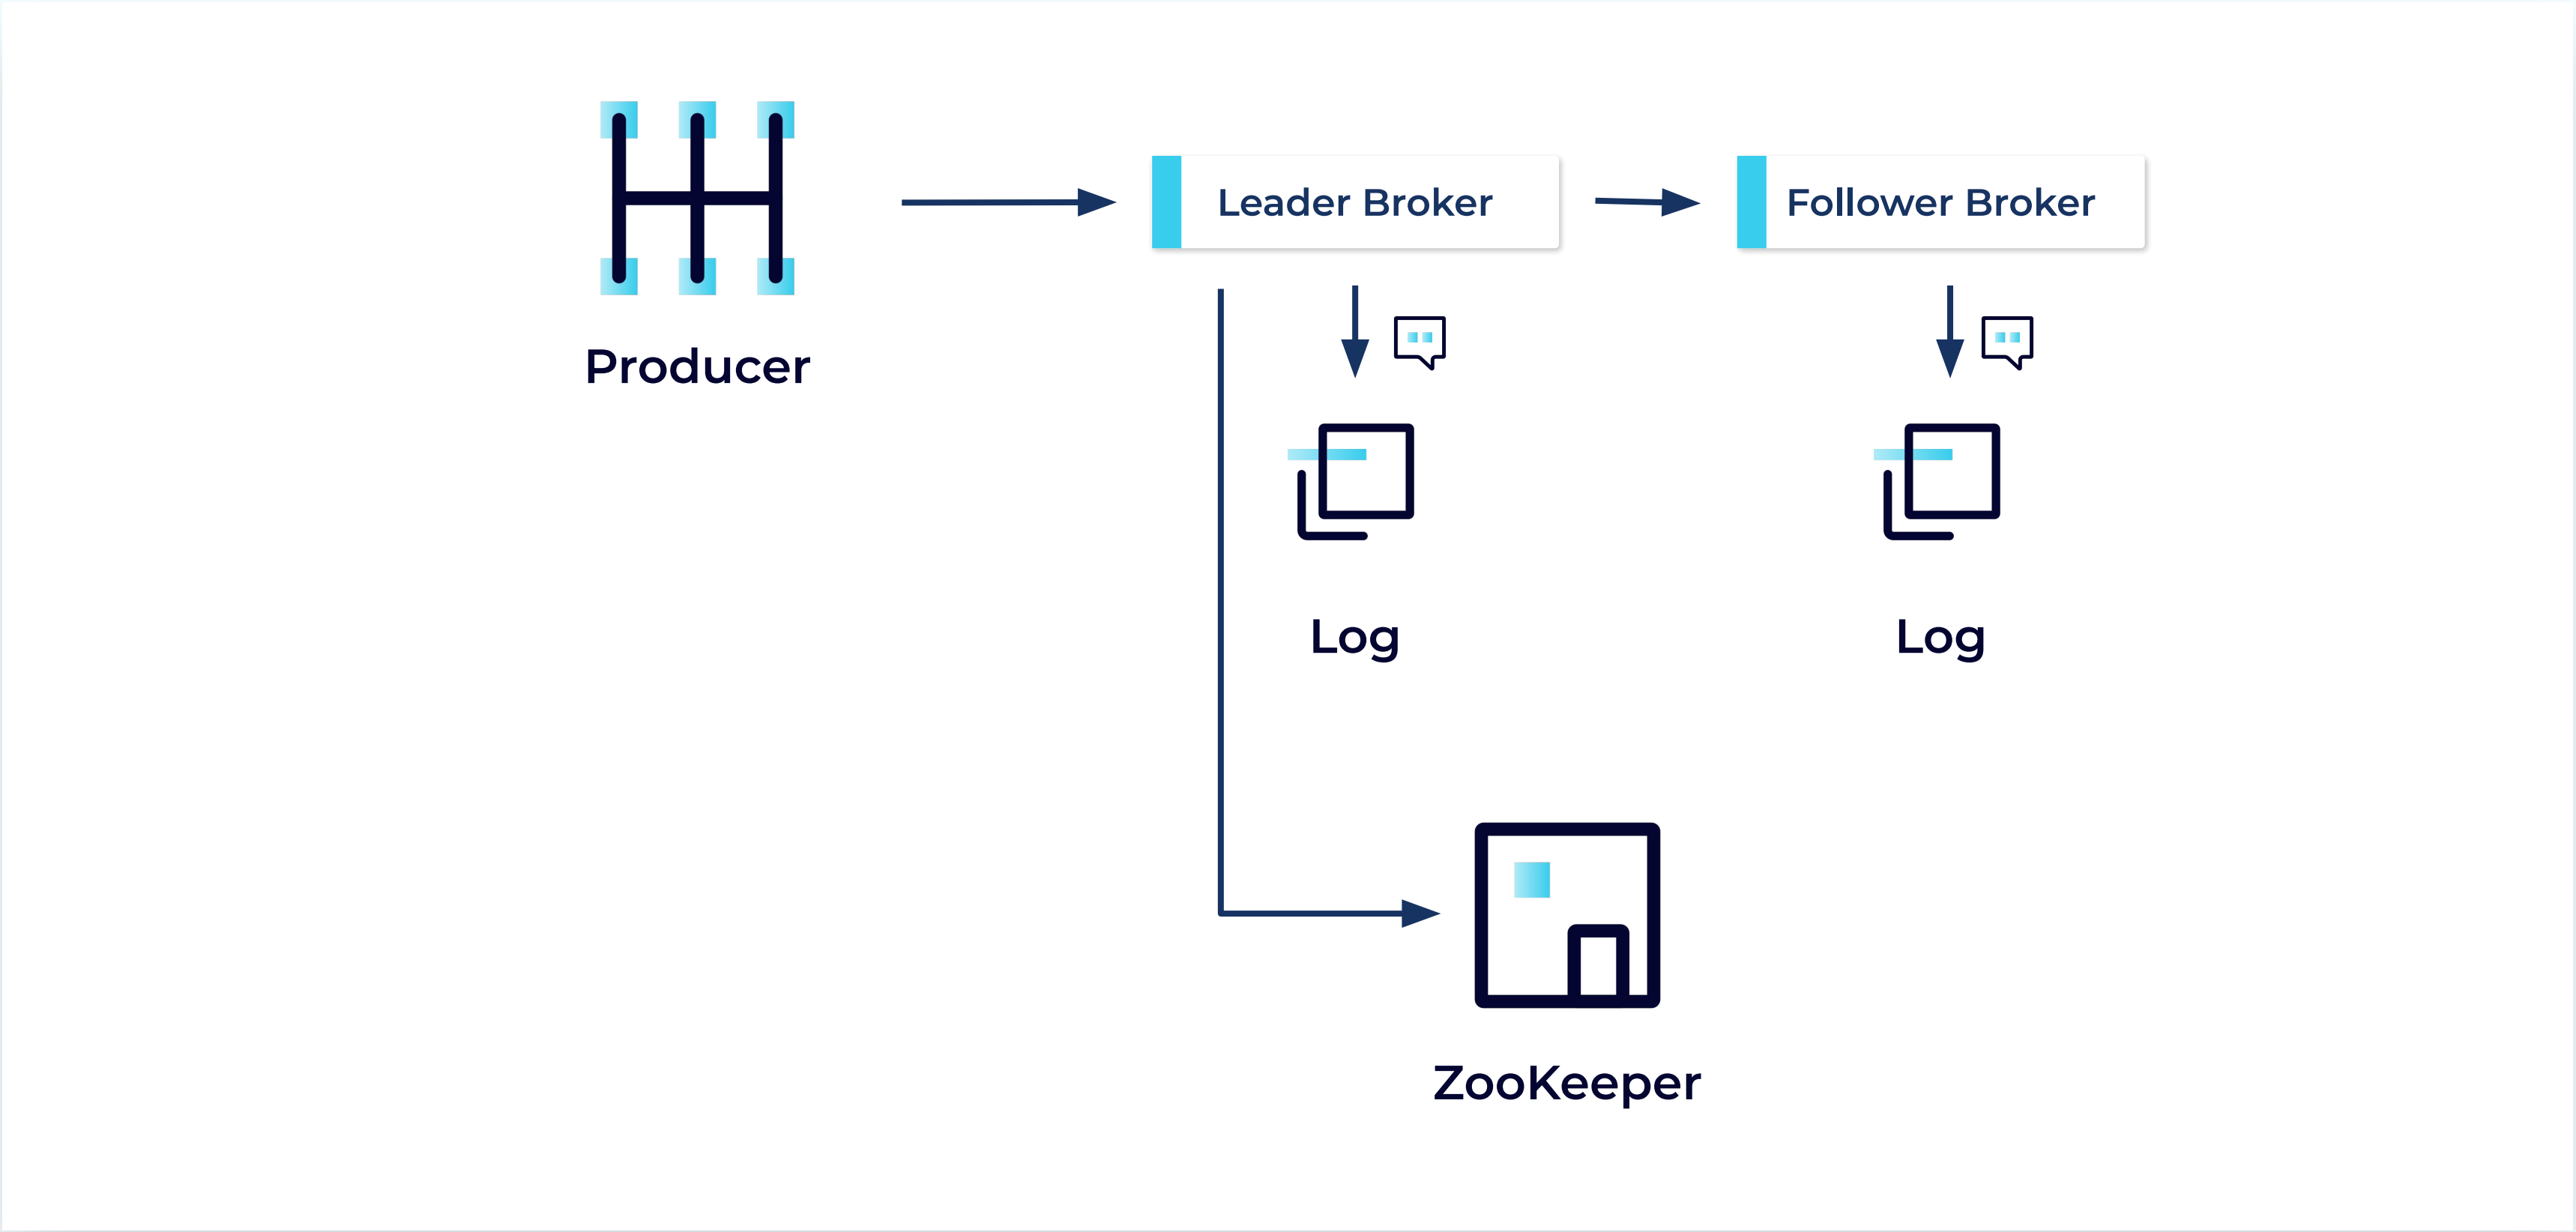
Data Flow in Kafka



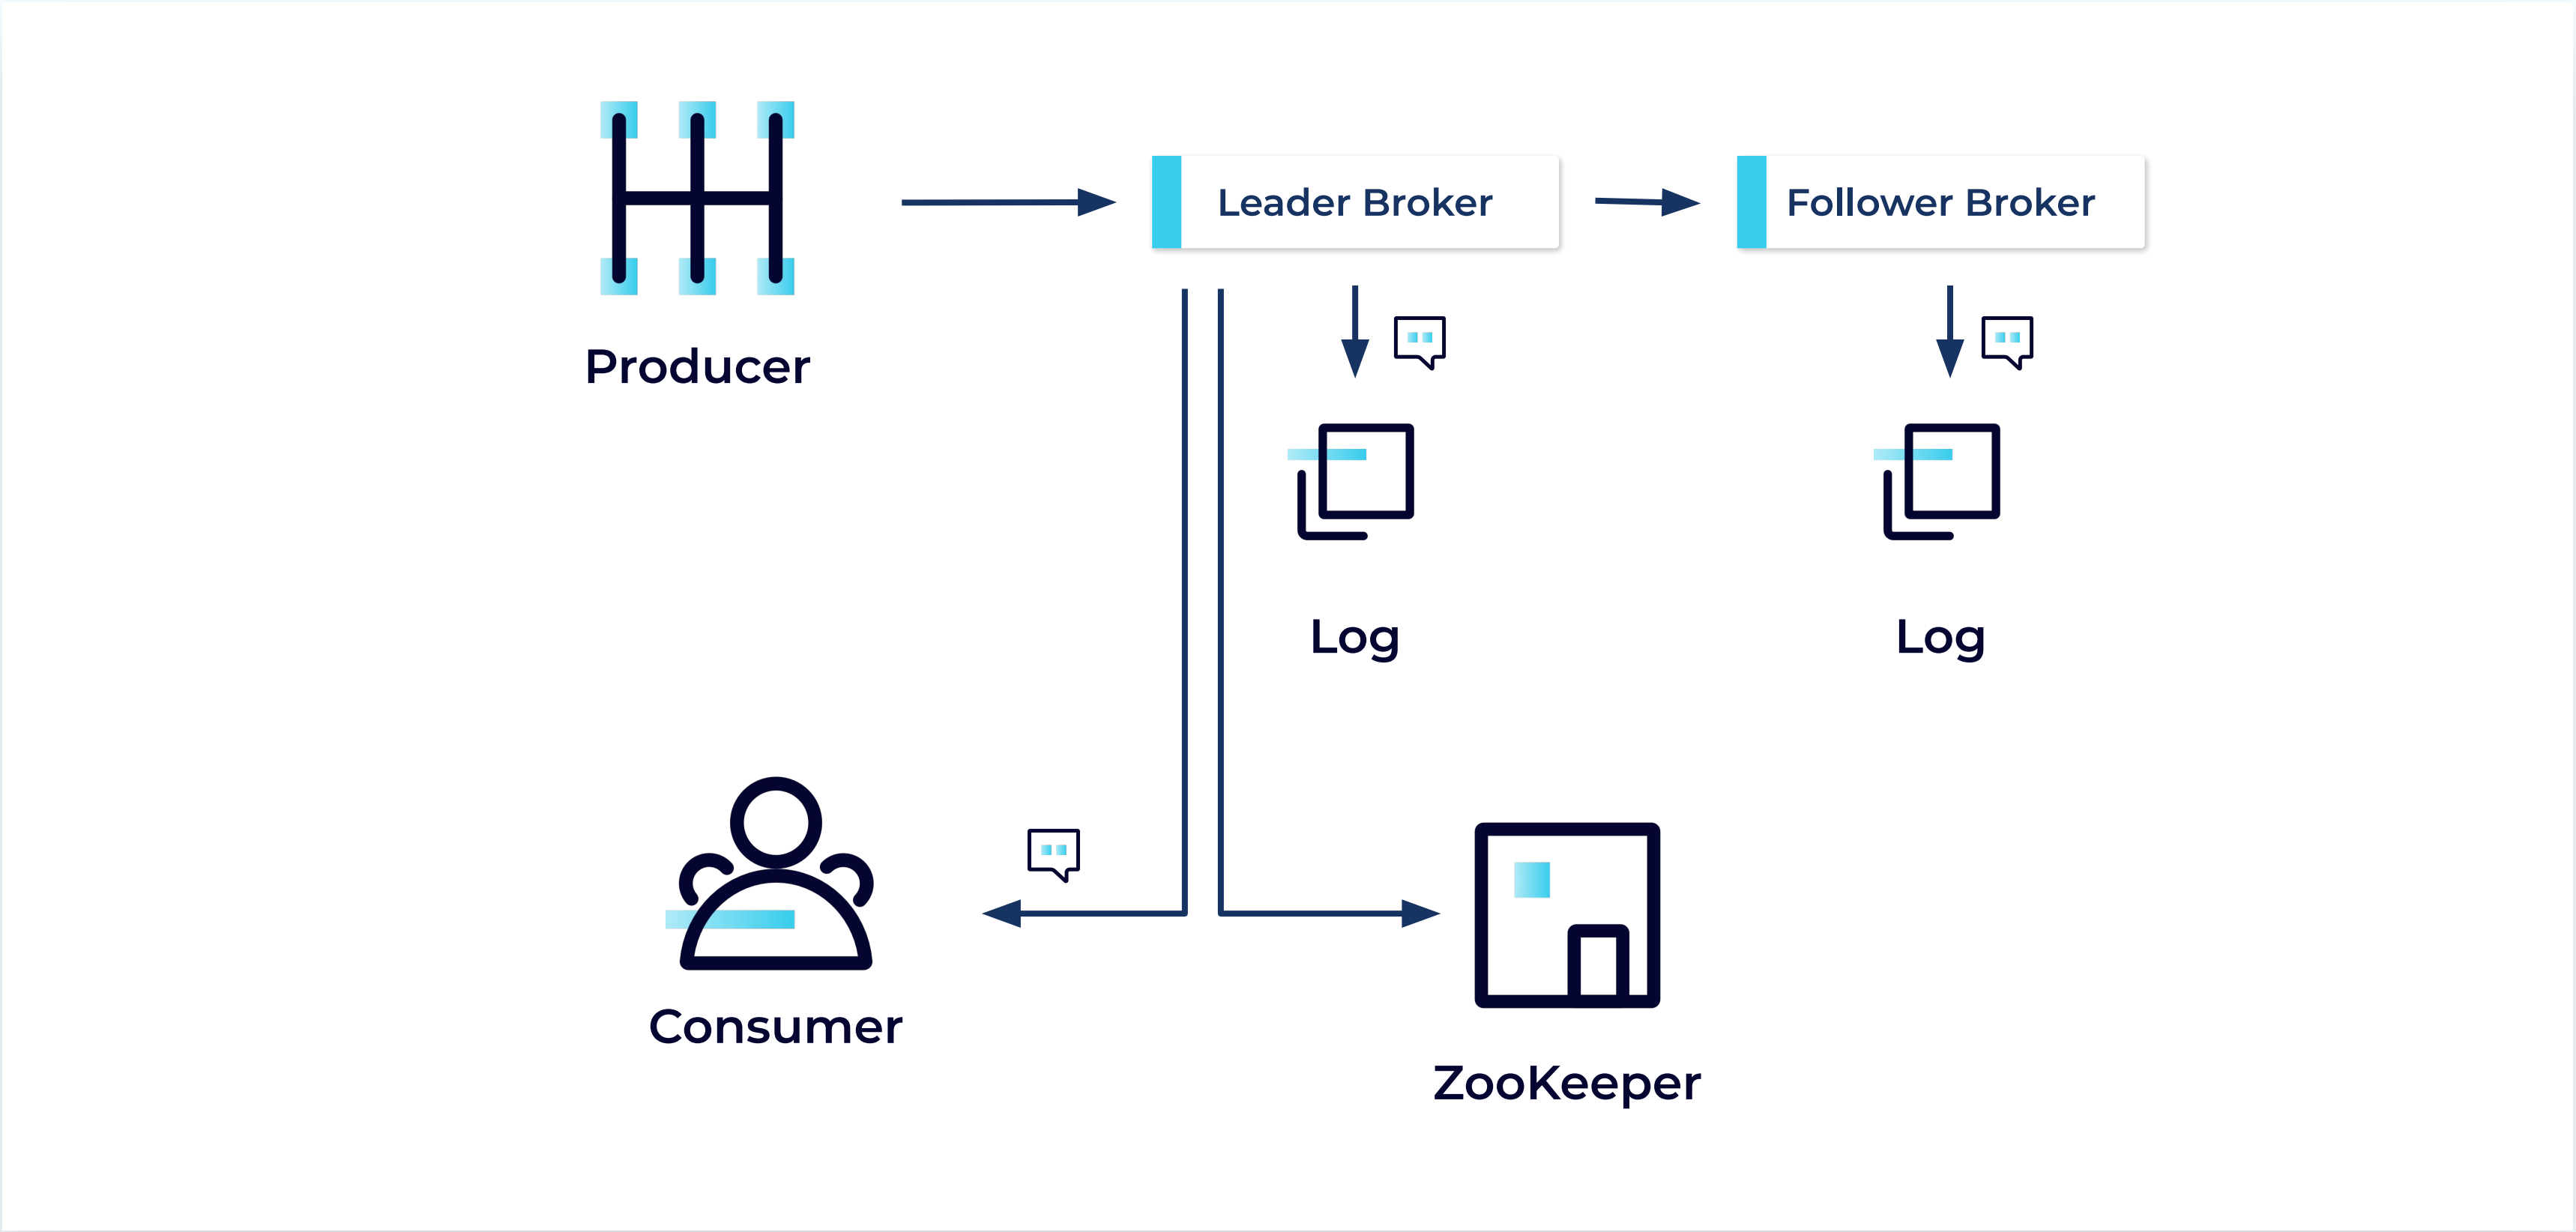
Data Flow in Kafka



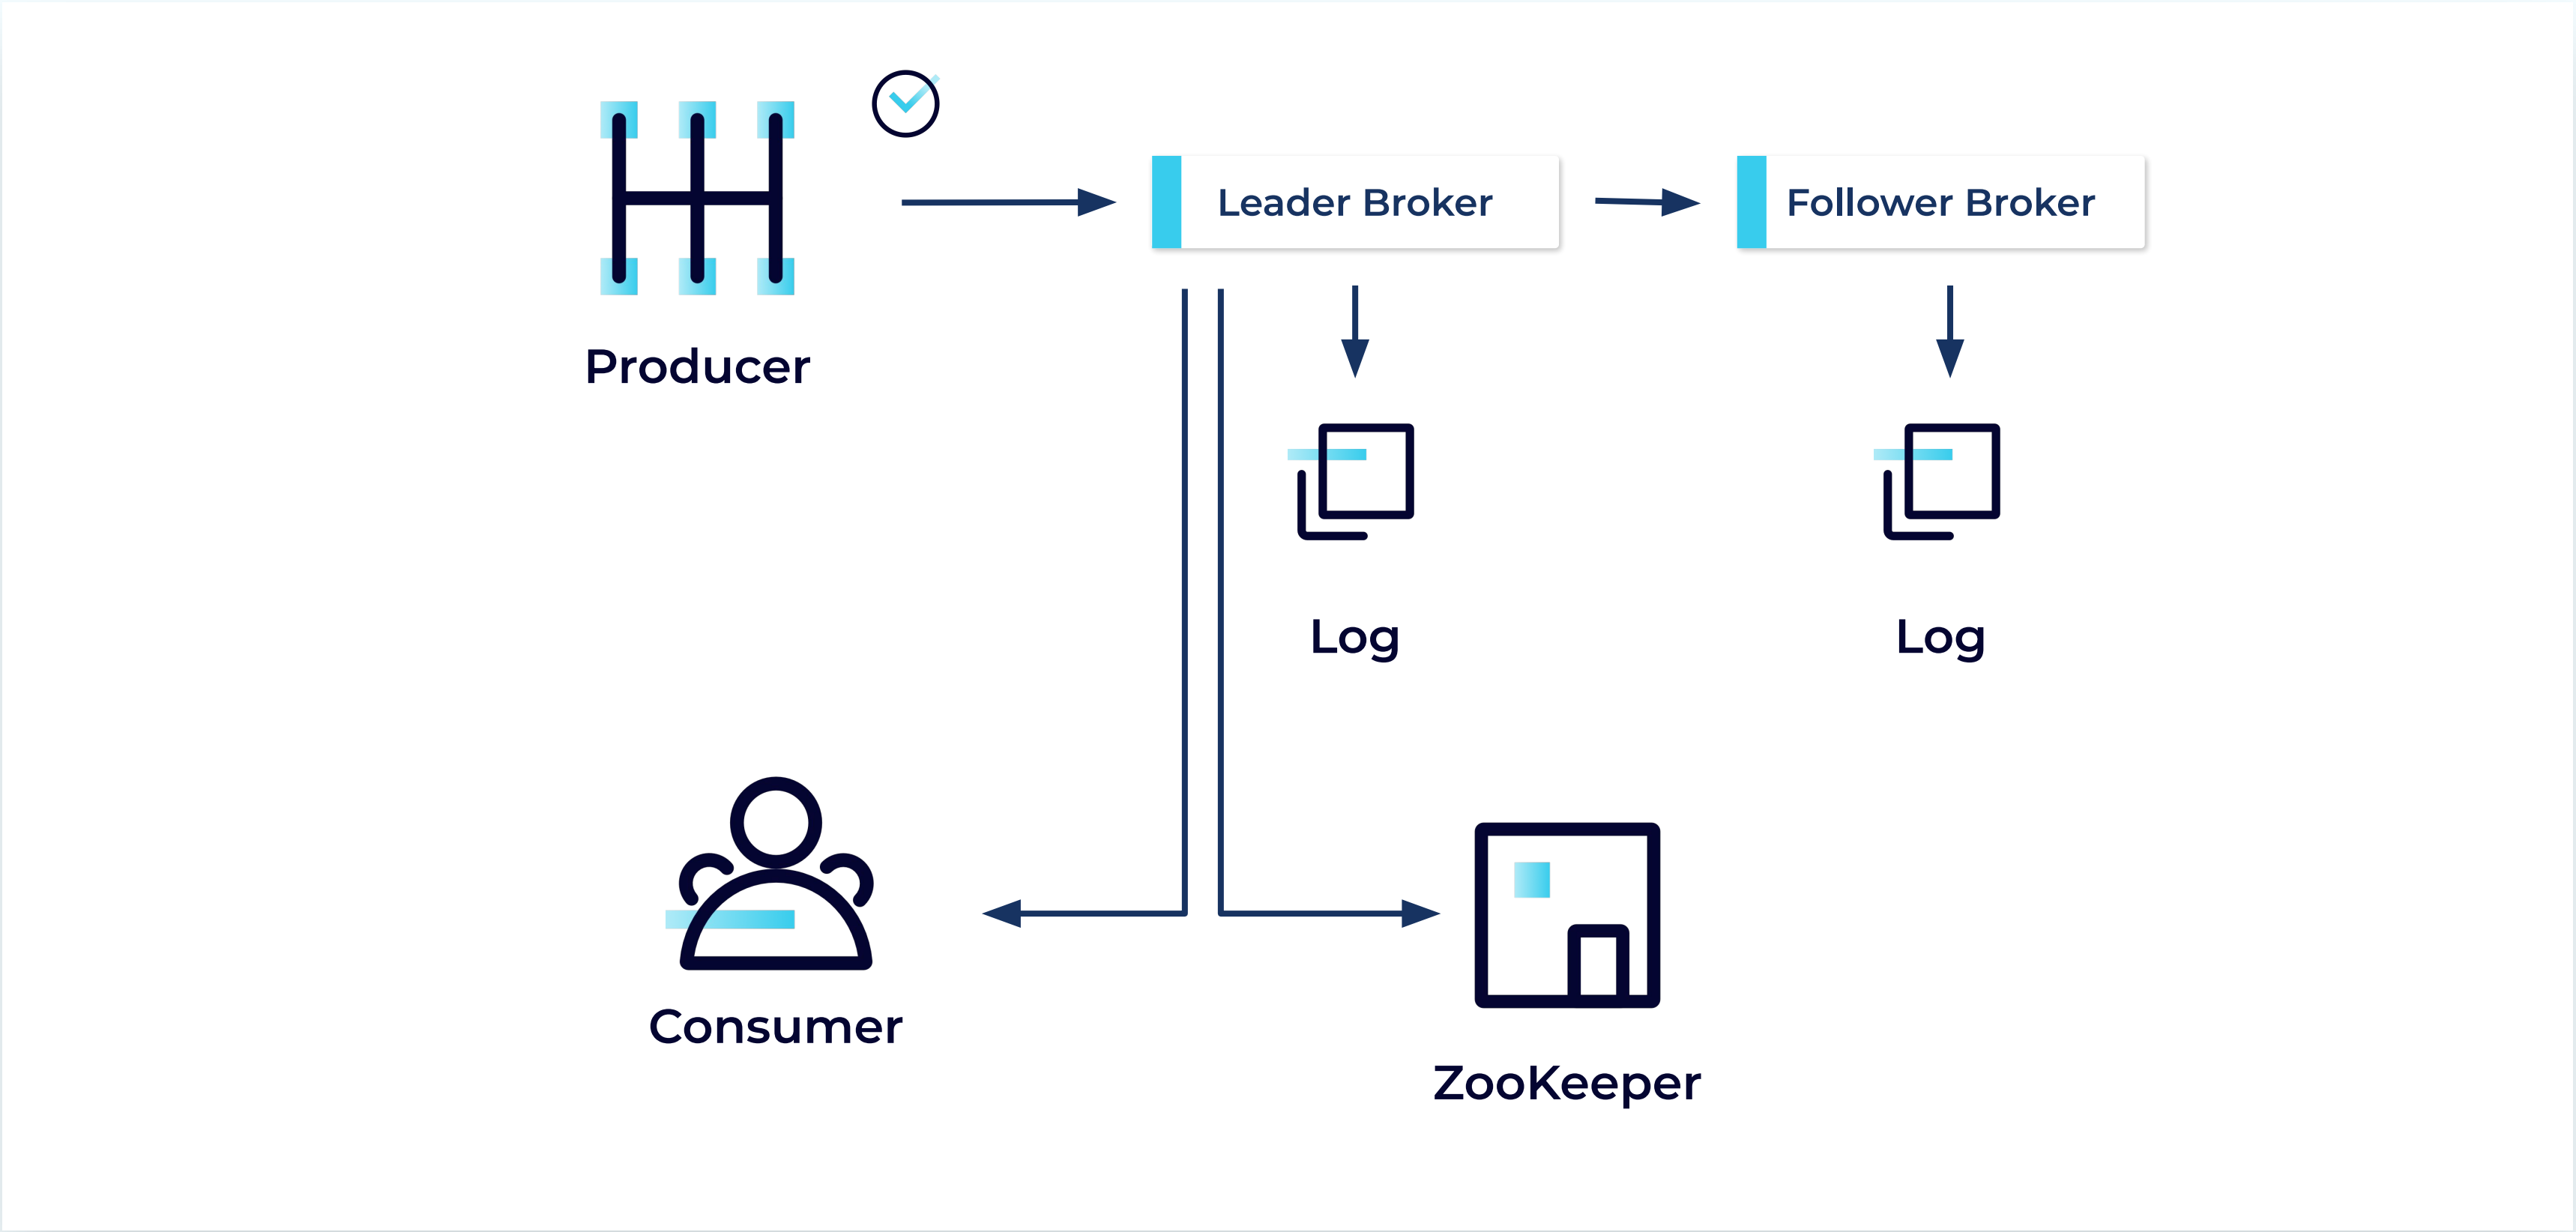
Data Flow in Kafka



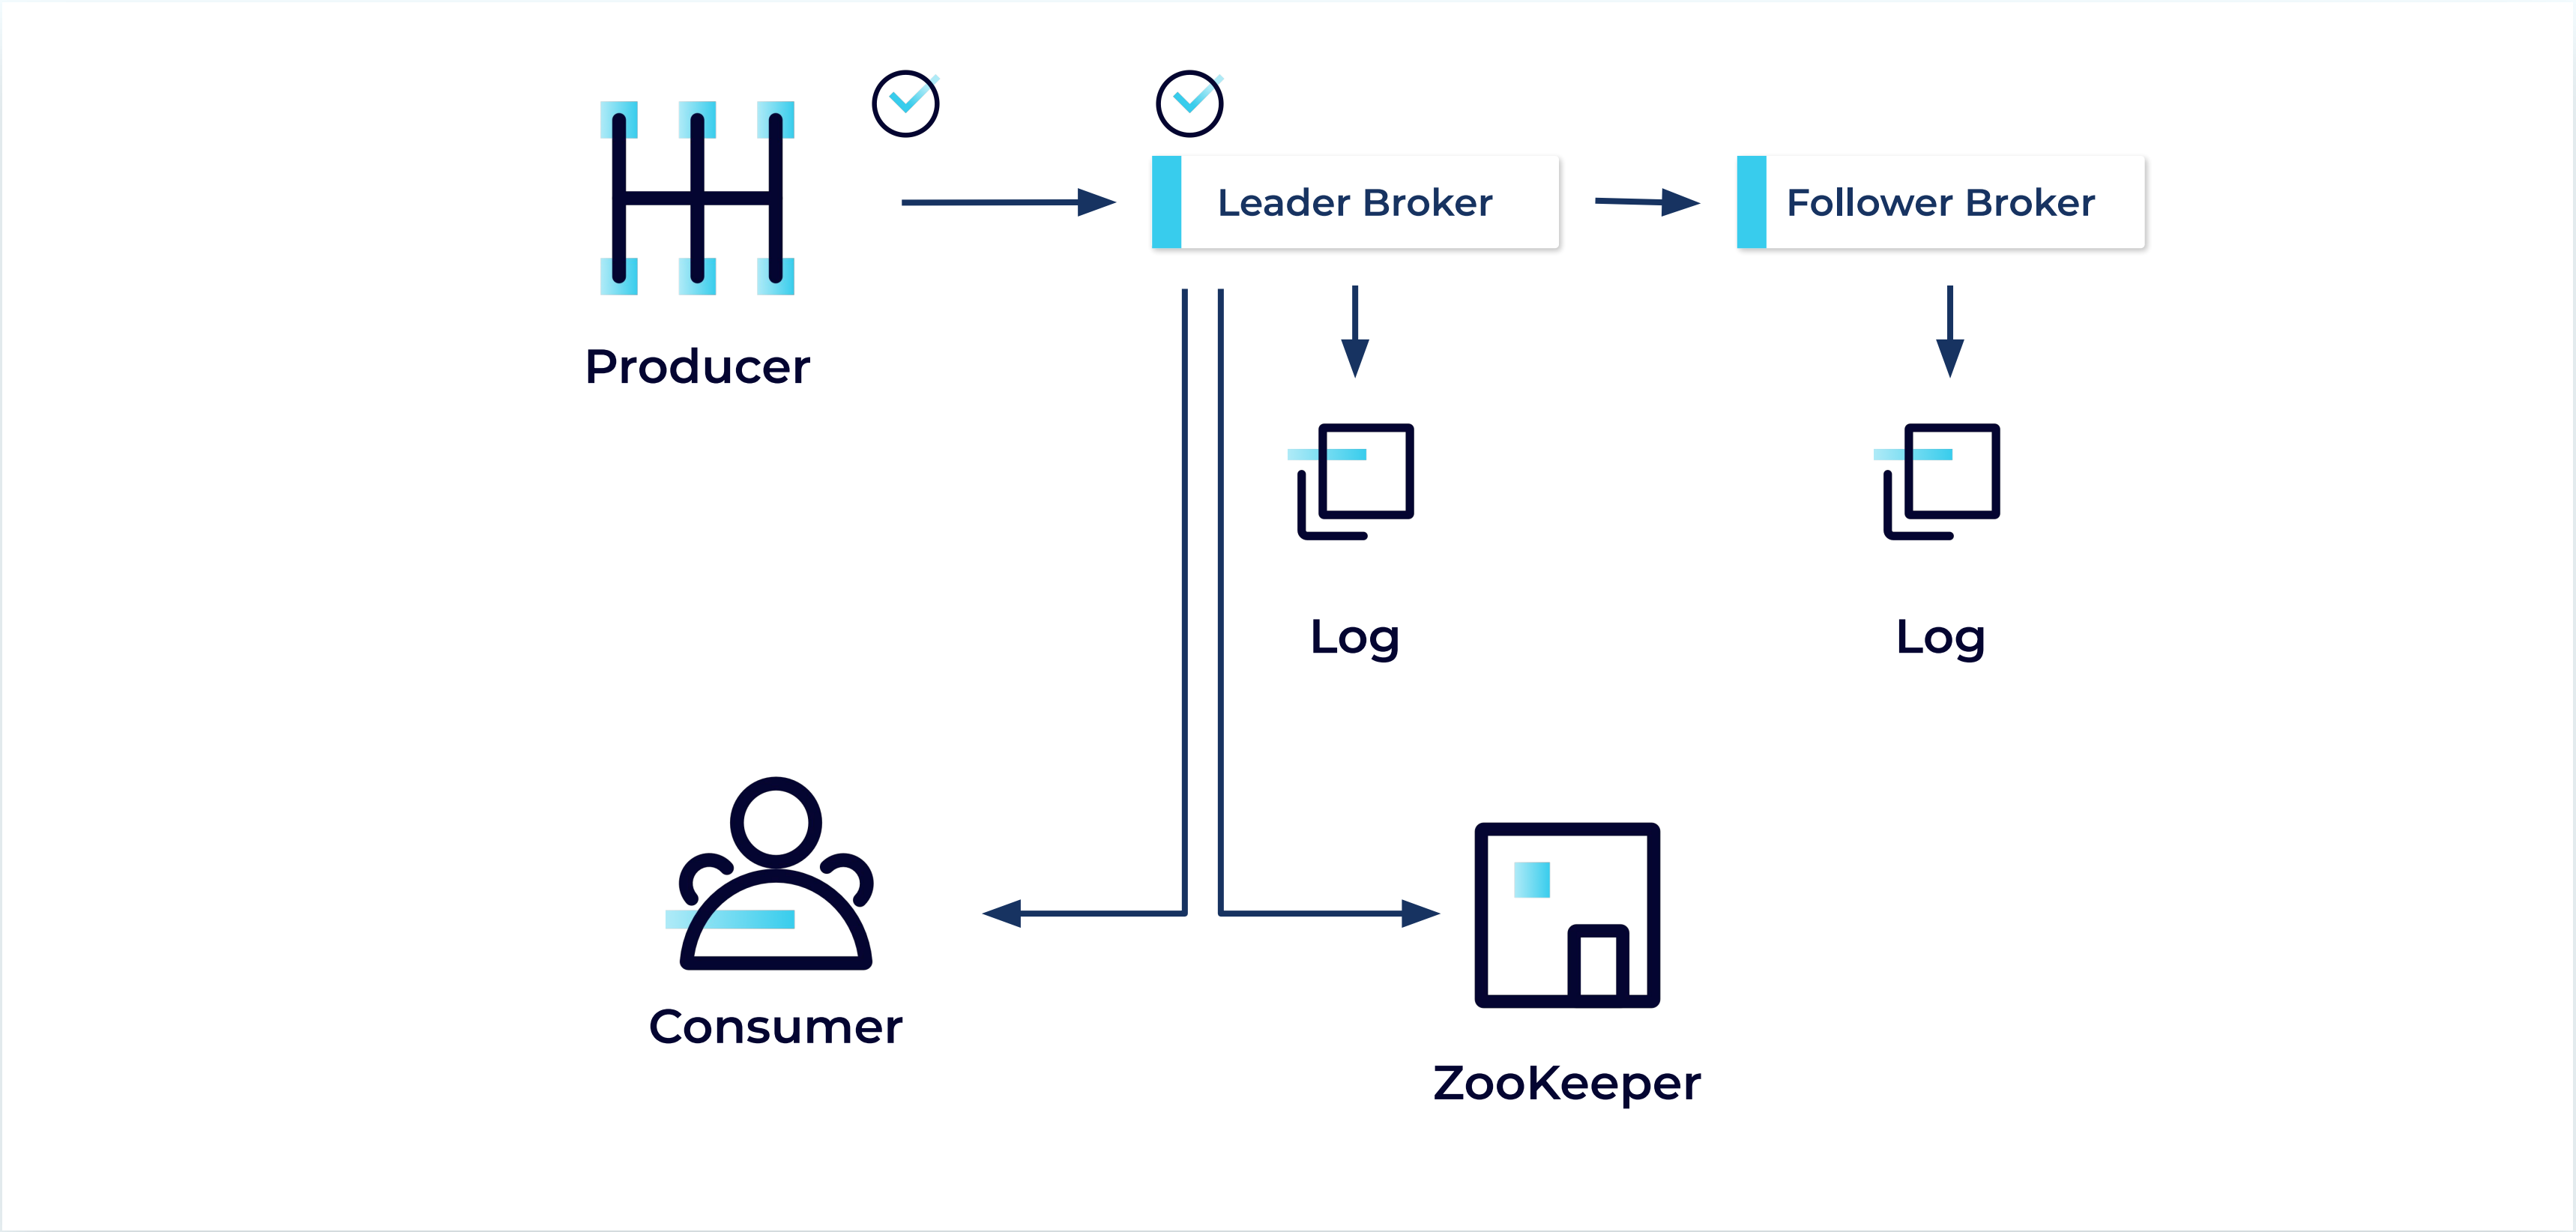
Data Flow in Kafka



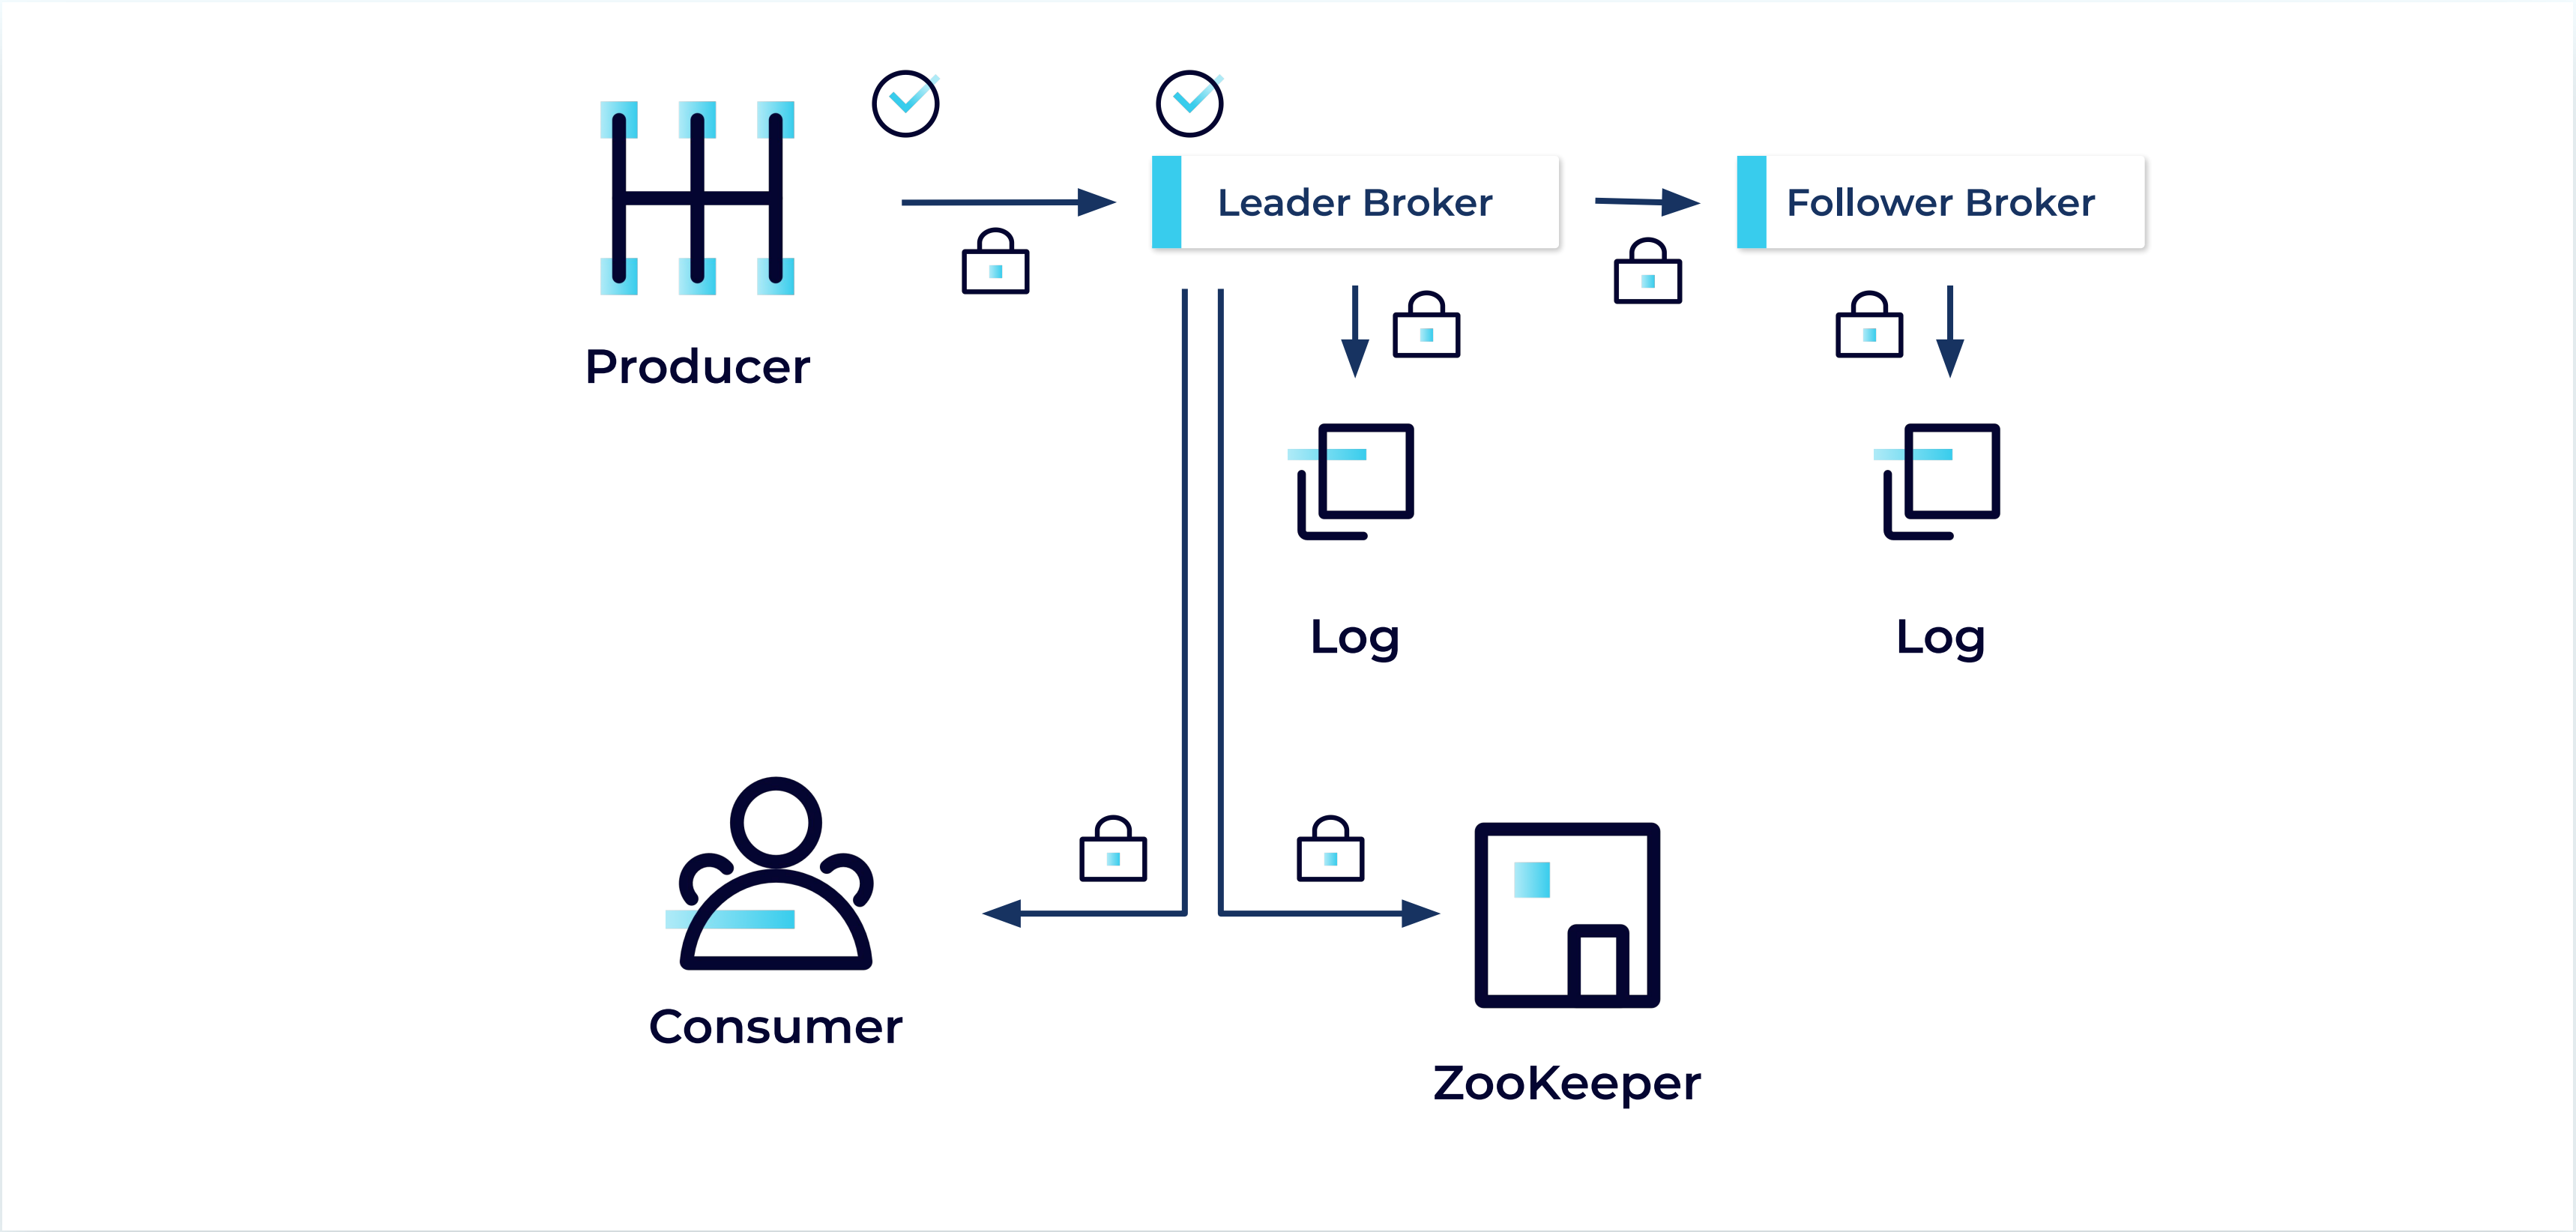
Data Flow in Kafka



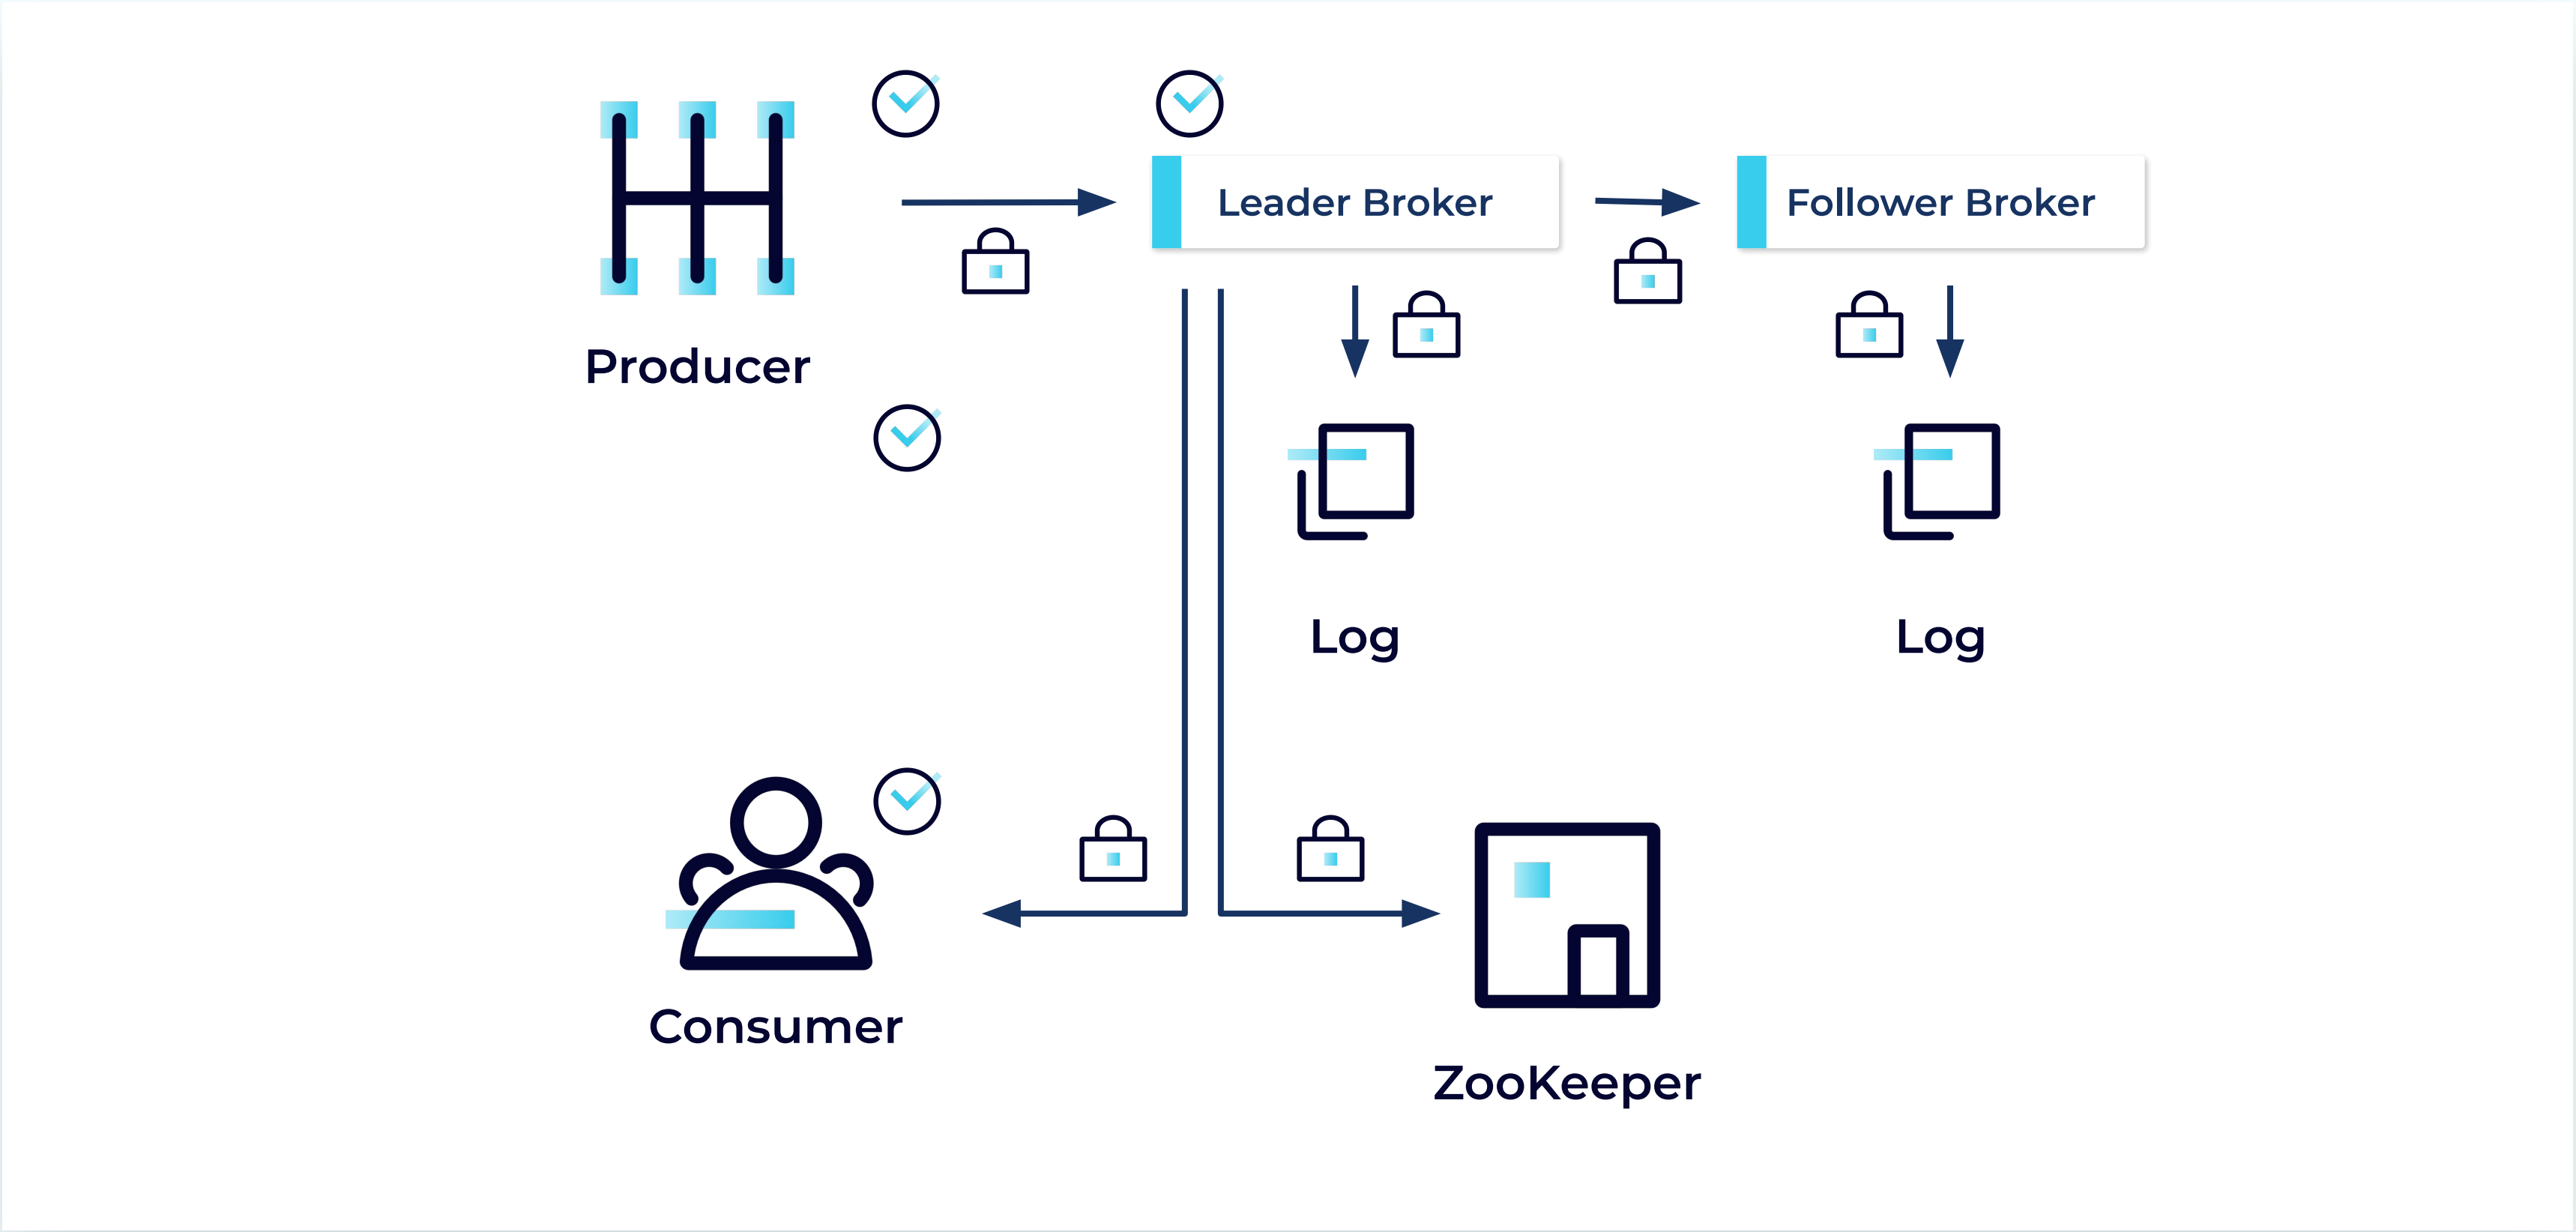
Data Flow in Kafka



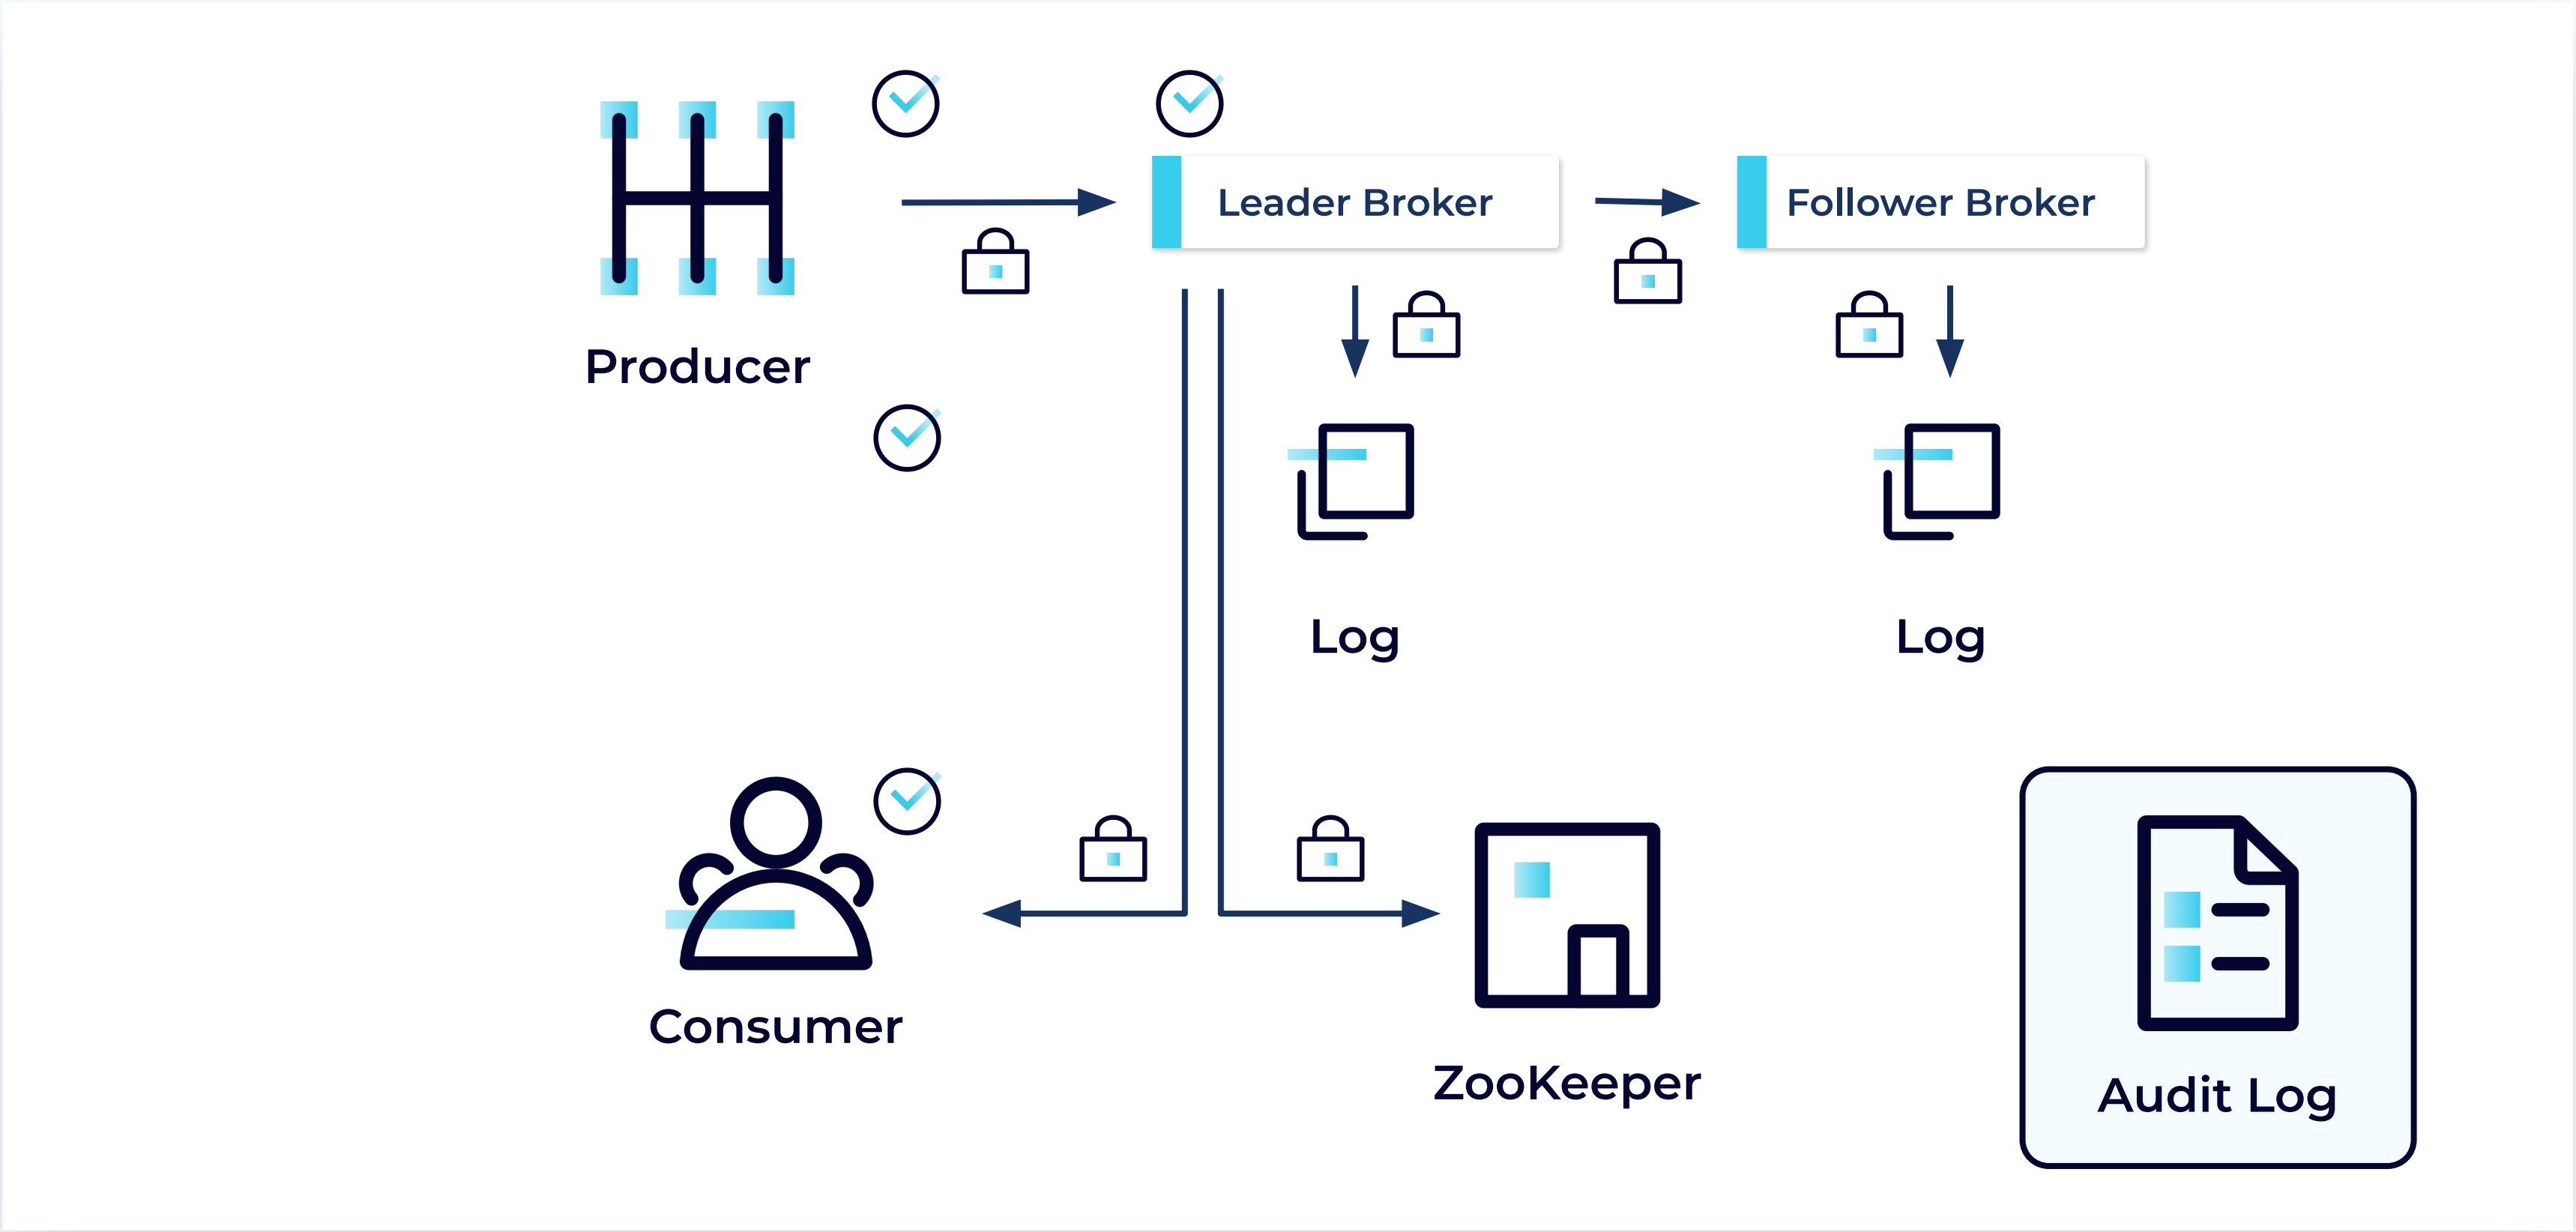
Data Flow in Kafka



Data Flow in Kafka




Data Flow in Kafka





Kafka Authentication Basics

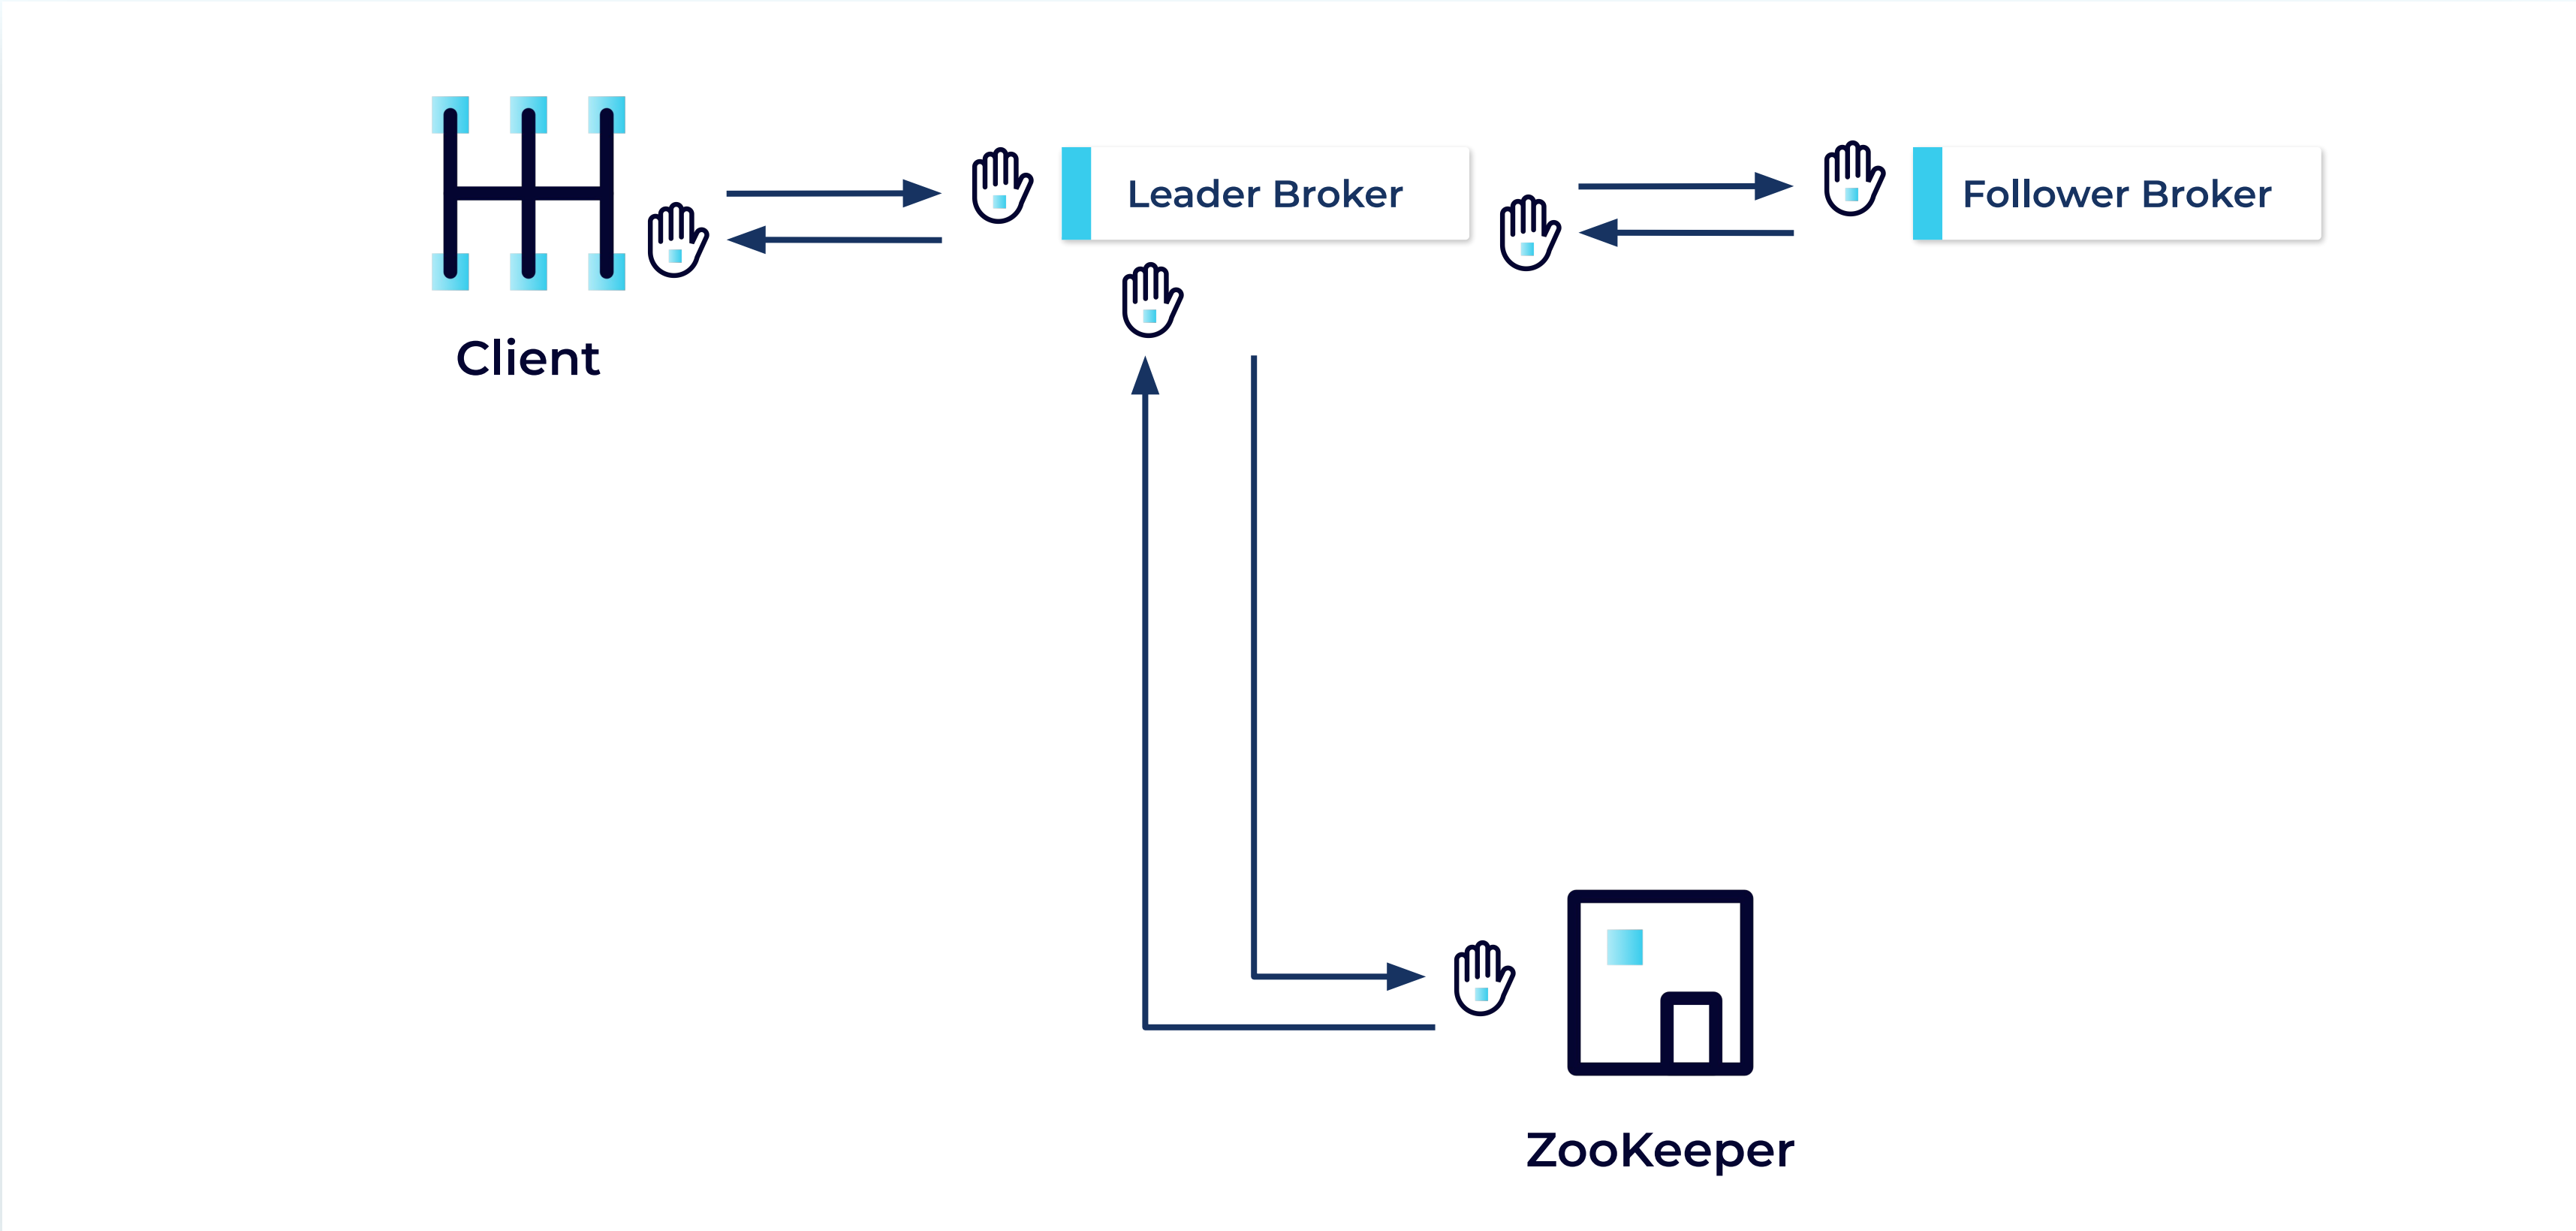
Data Flow in Kafka



Data Flow in Kafka



Data Flow in Kafka



Data Flow in Kafka

Username: sally



KafkaPrincipal

Interaction Types


Clients

- Users
- Applications
- Services

Brokers

- Broker-to-Broker communication

Listeners and Security Protocols

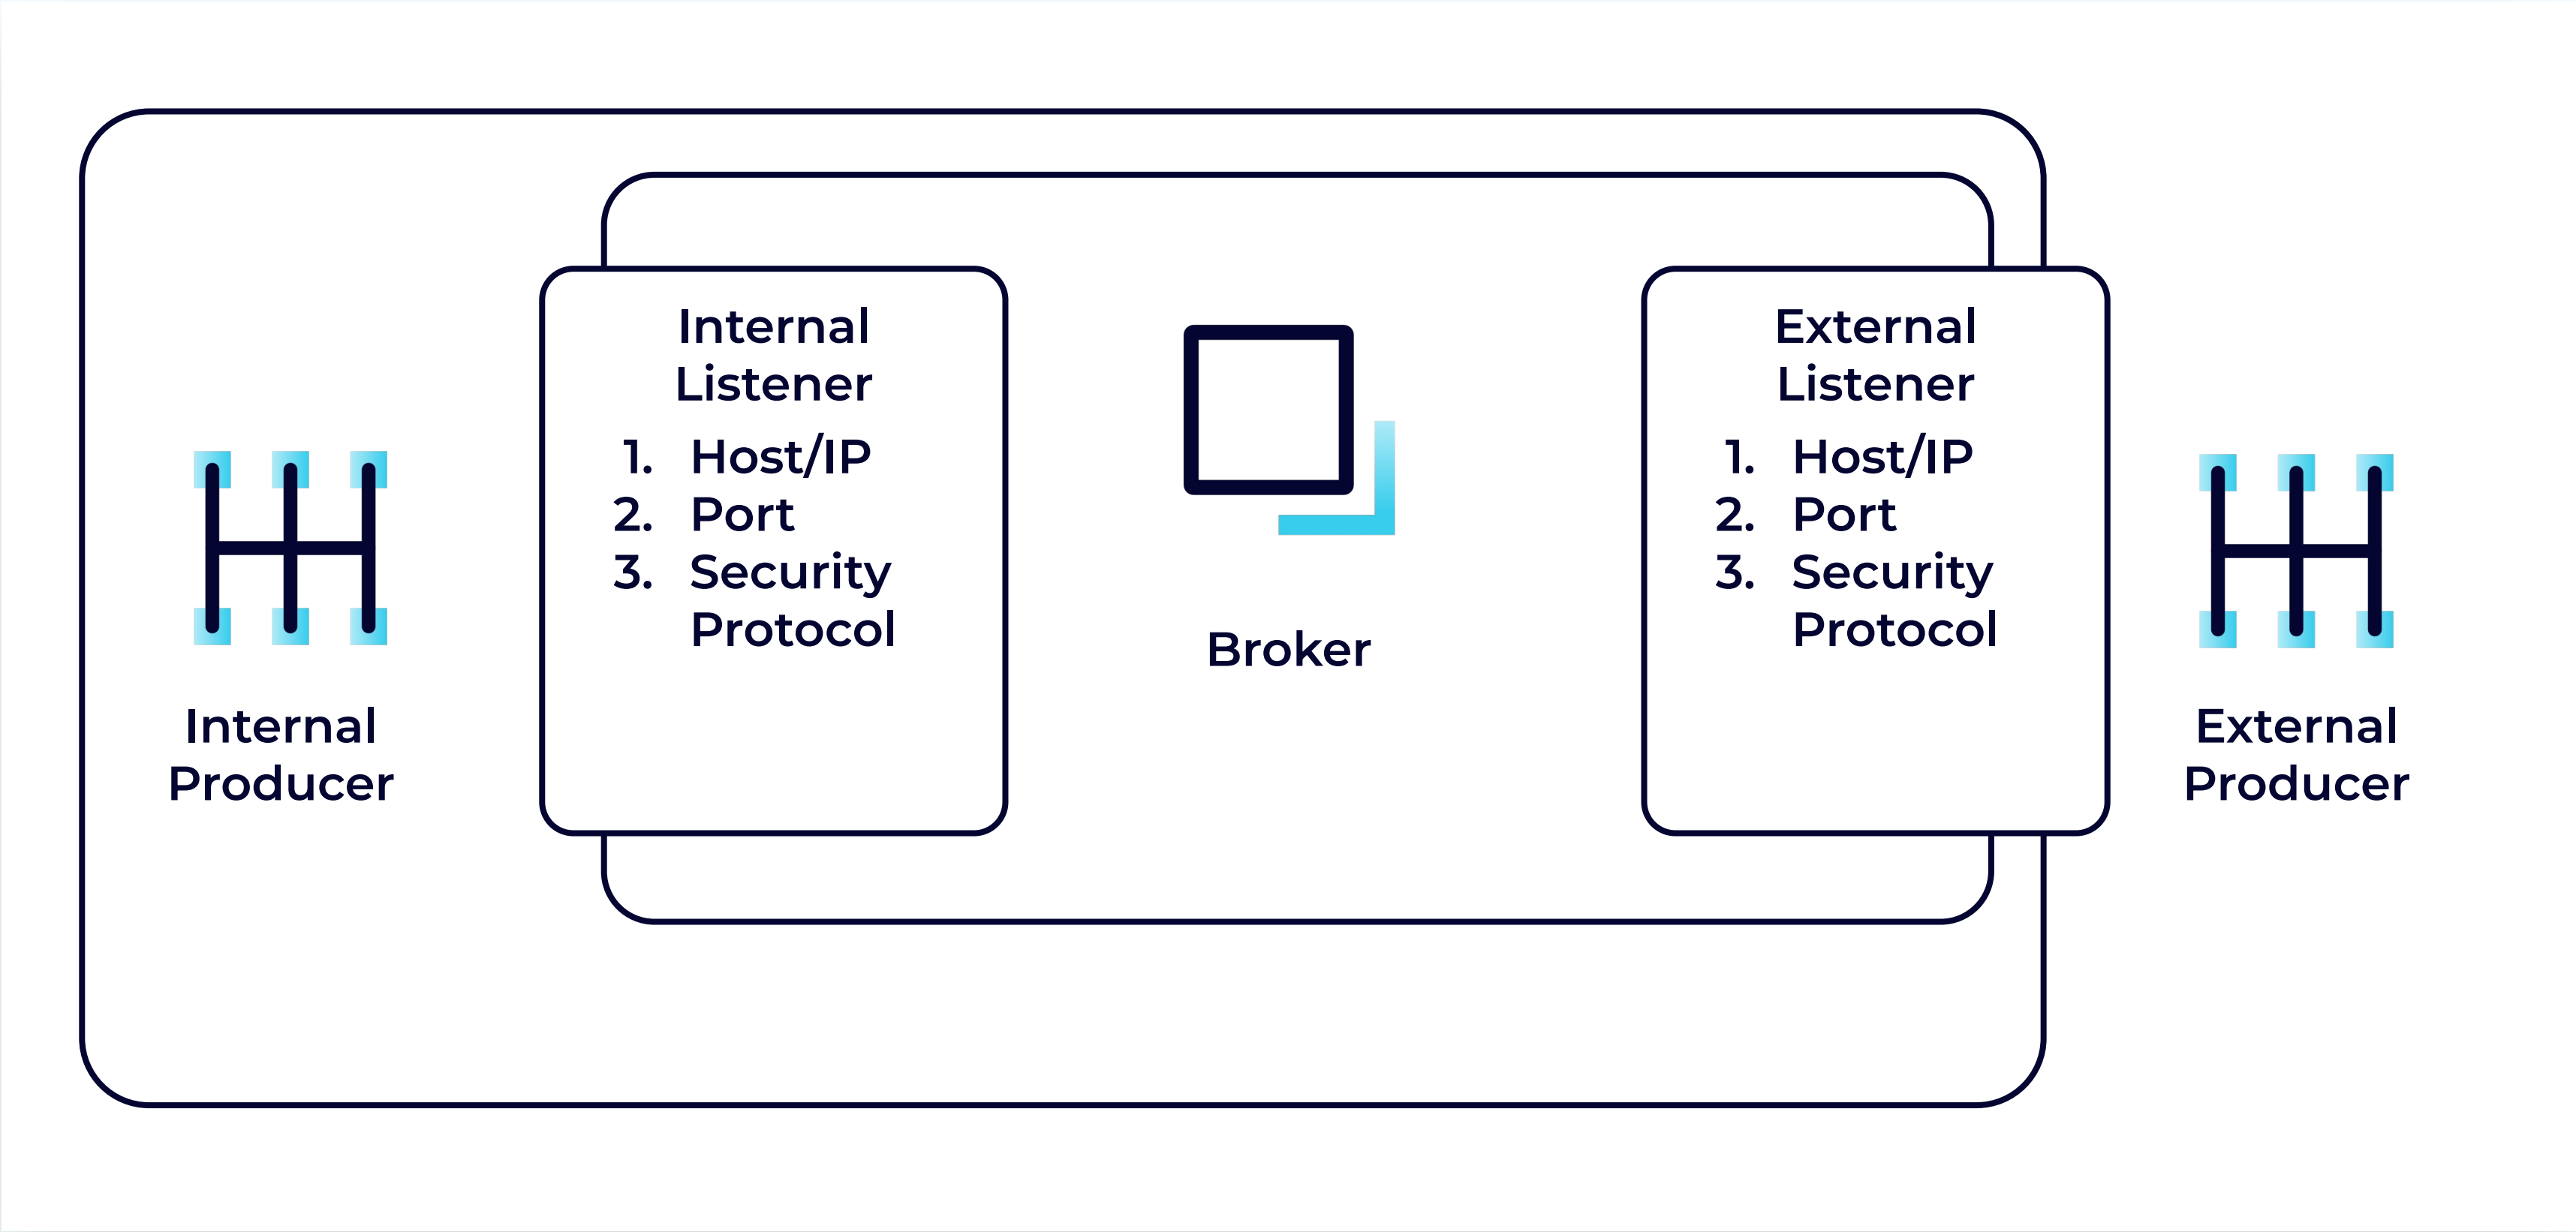


Securing Kafka

Security Protocols

- **PLAINTEXT - not secure**
- **SASL_PLAINTEXT - not secure**
- **SSL - secure**
- **SASL_SSL - secure**

Listeners and Security Protocols



Configuring listeners on the broker



Broker configuration

listeners=EXTERNAL://:9092,INTERNAL://10.0.0.2:9093,BROKER://10.0.0.2:9094
advertised.listeners=EXTERNAL://broker1.example.com:9092,INTERNAL://
broker1.local:9093,BROKER://broker1.local:9094

listener.security.protocol.map=EXTERNAL:SASL_SSL,INTERNAL:SSL,BROKER:SSL

inter.broker.listener.name=BROKER

Configuring the client




Client configuration


security.protocol=SASL_SSL


bootstrap.servers=broker1.example.com:9092,broker2.example.com:9092



Kafka Authentication with SSL and SASL_SSL








Broker configuration

ssl.client.auth=required



Inter-broker Authentication

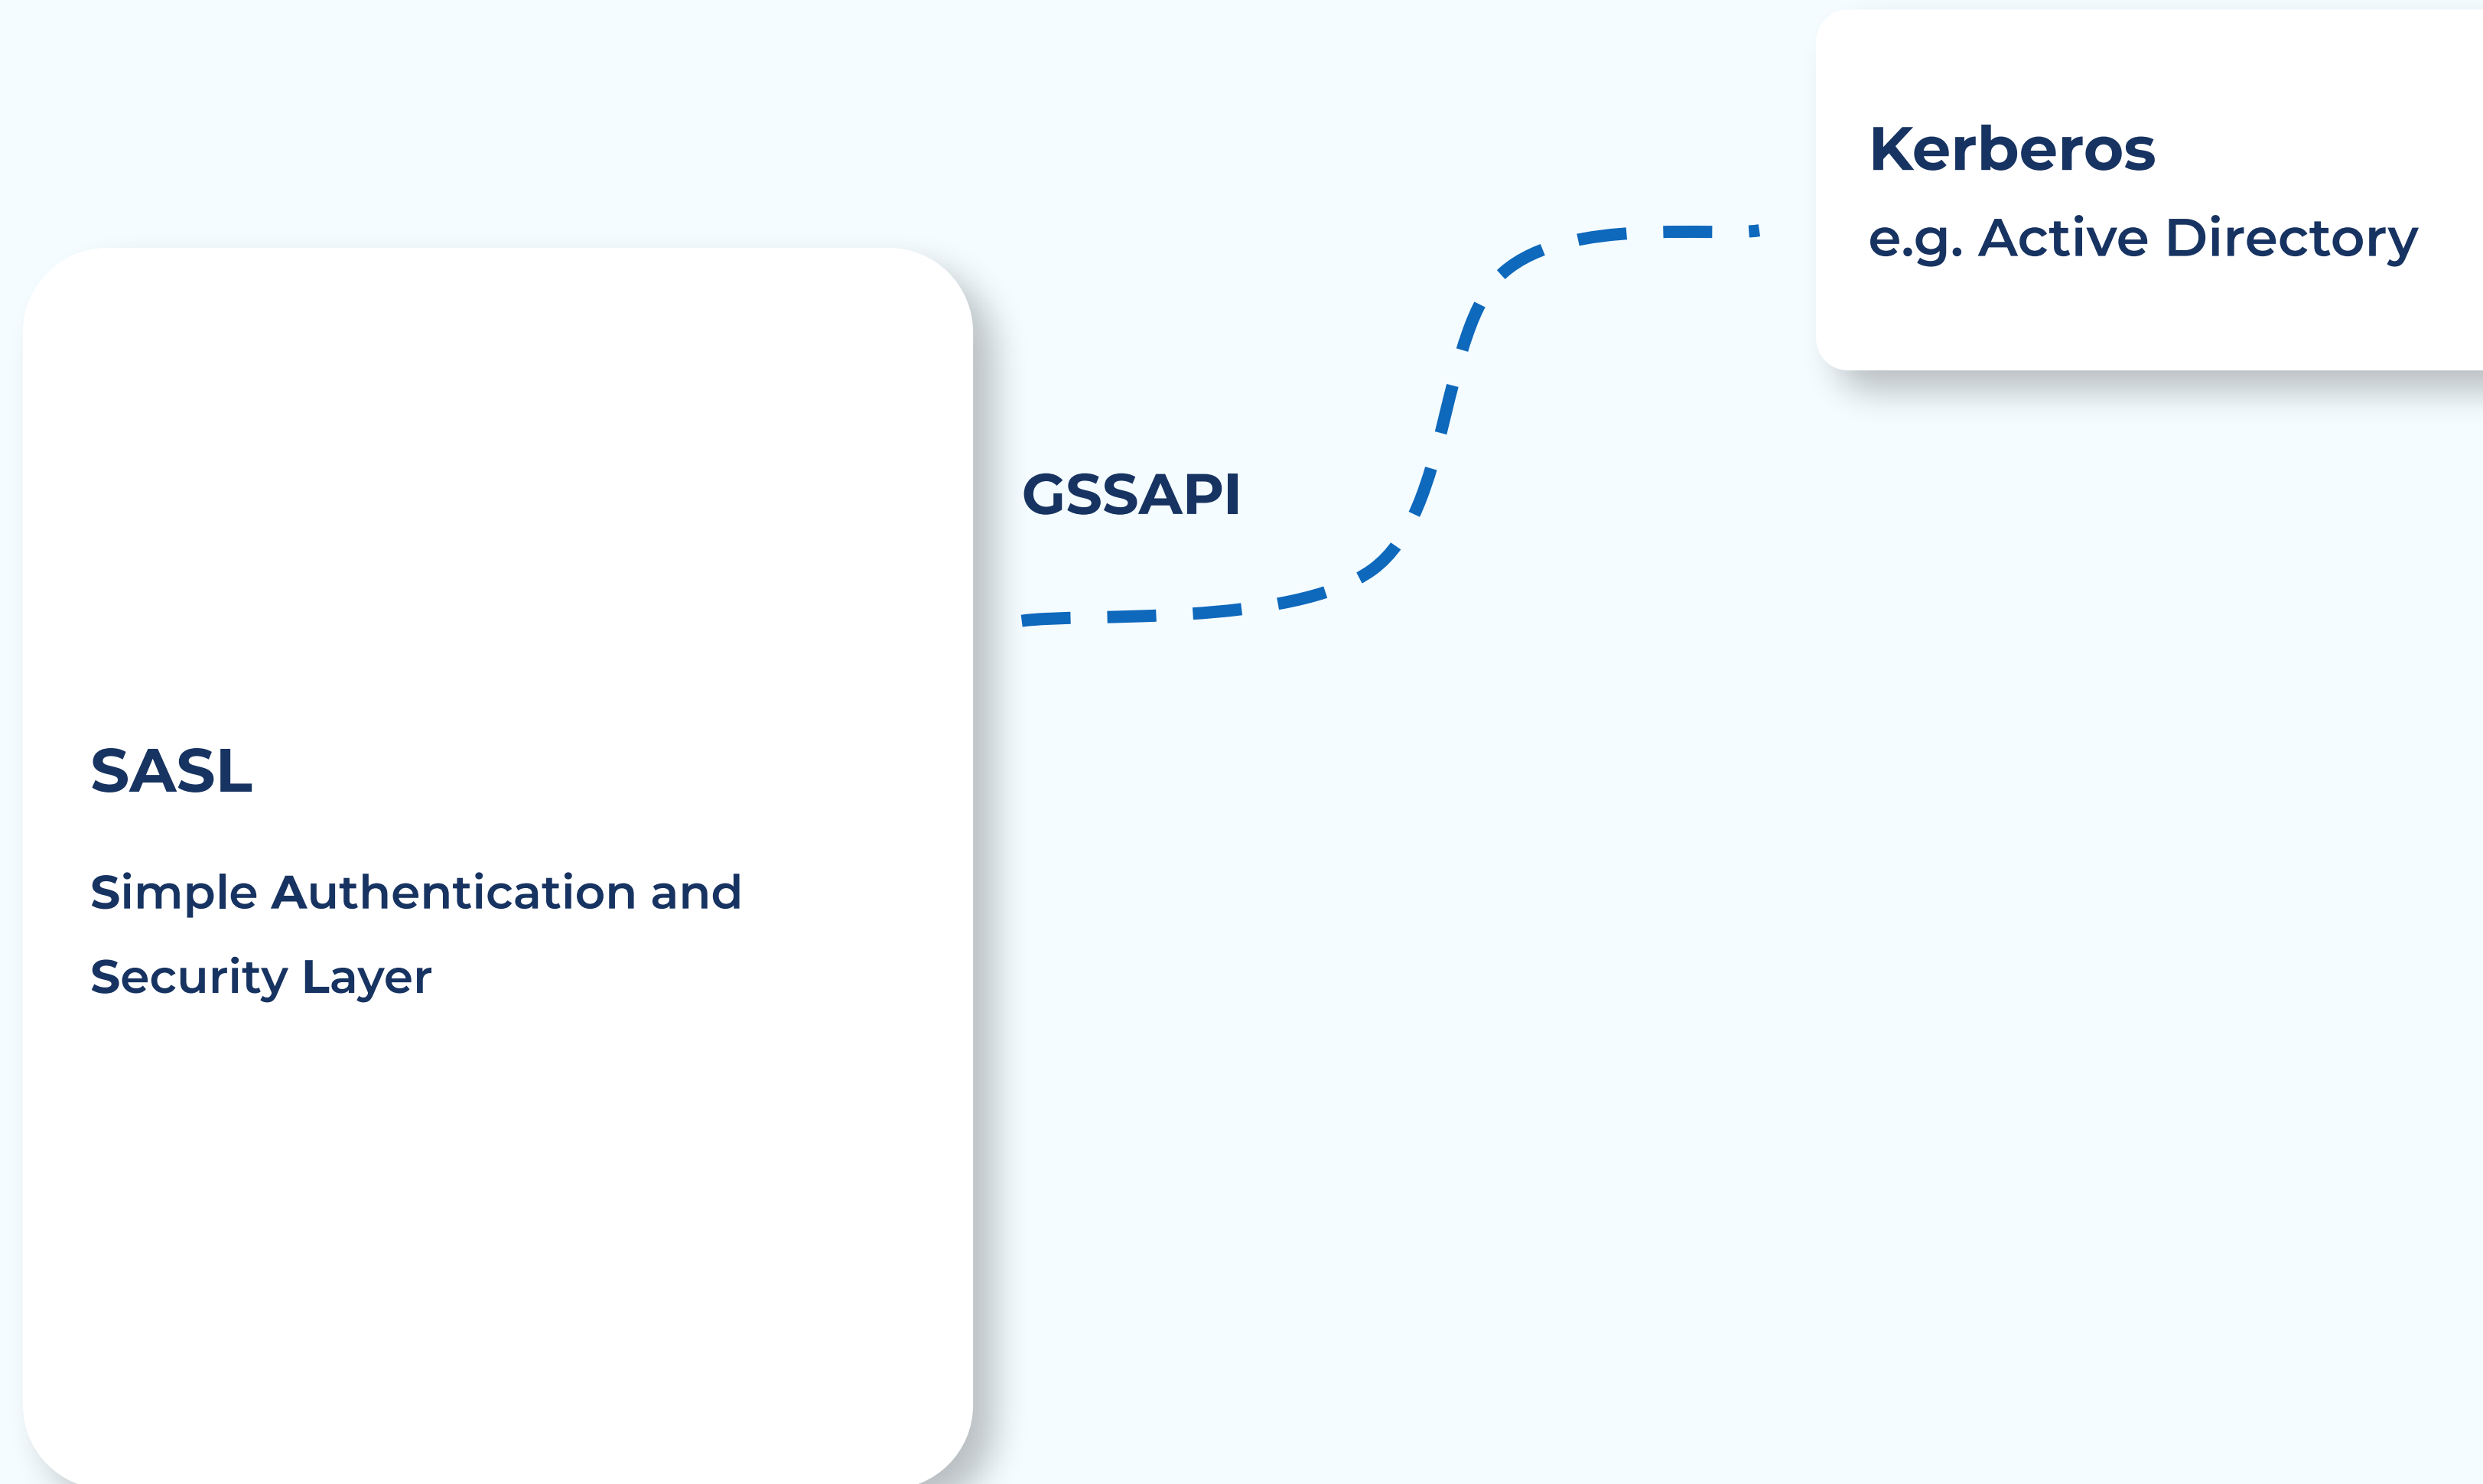


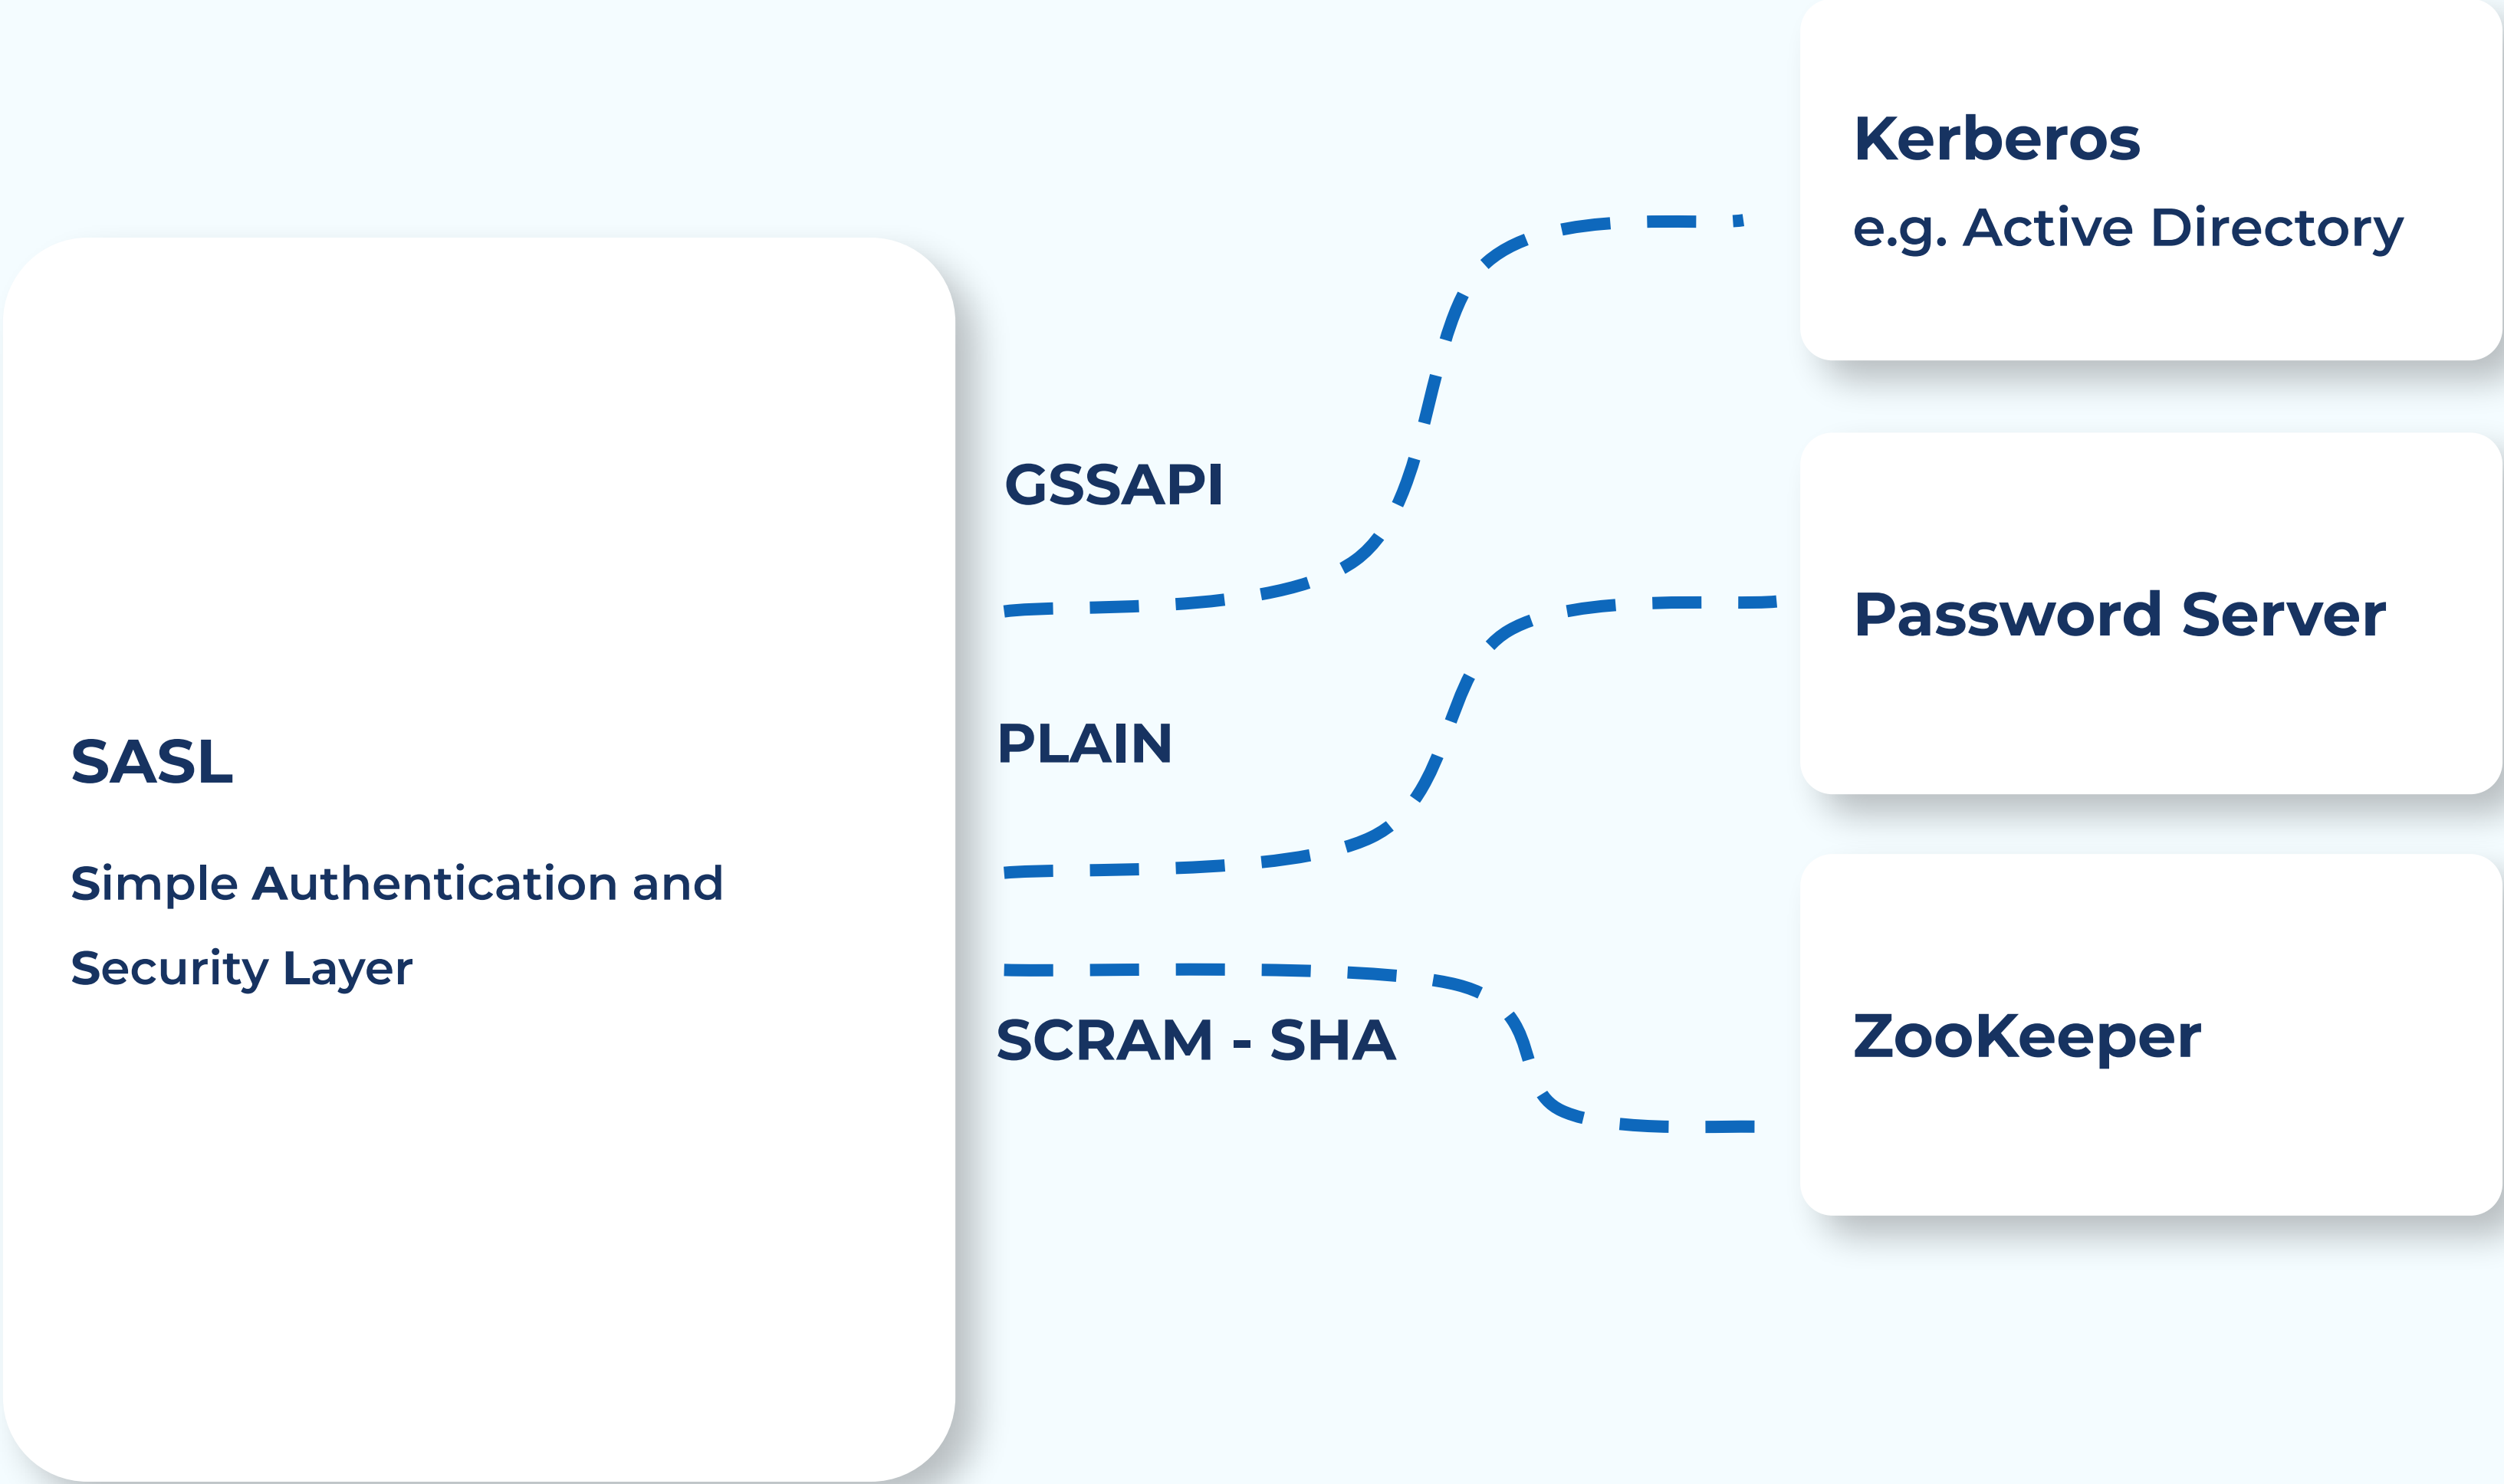
inter.broker.listener.name

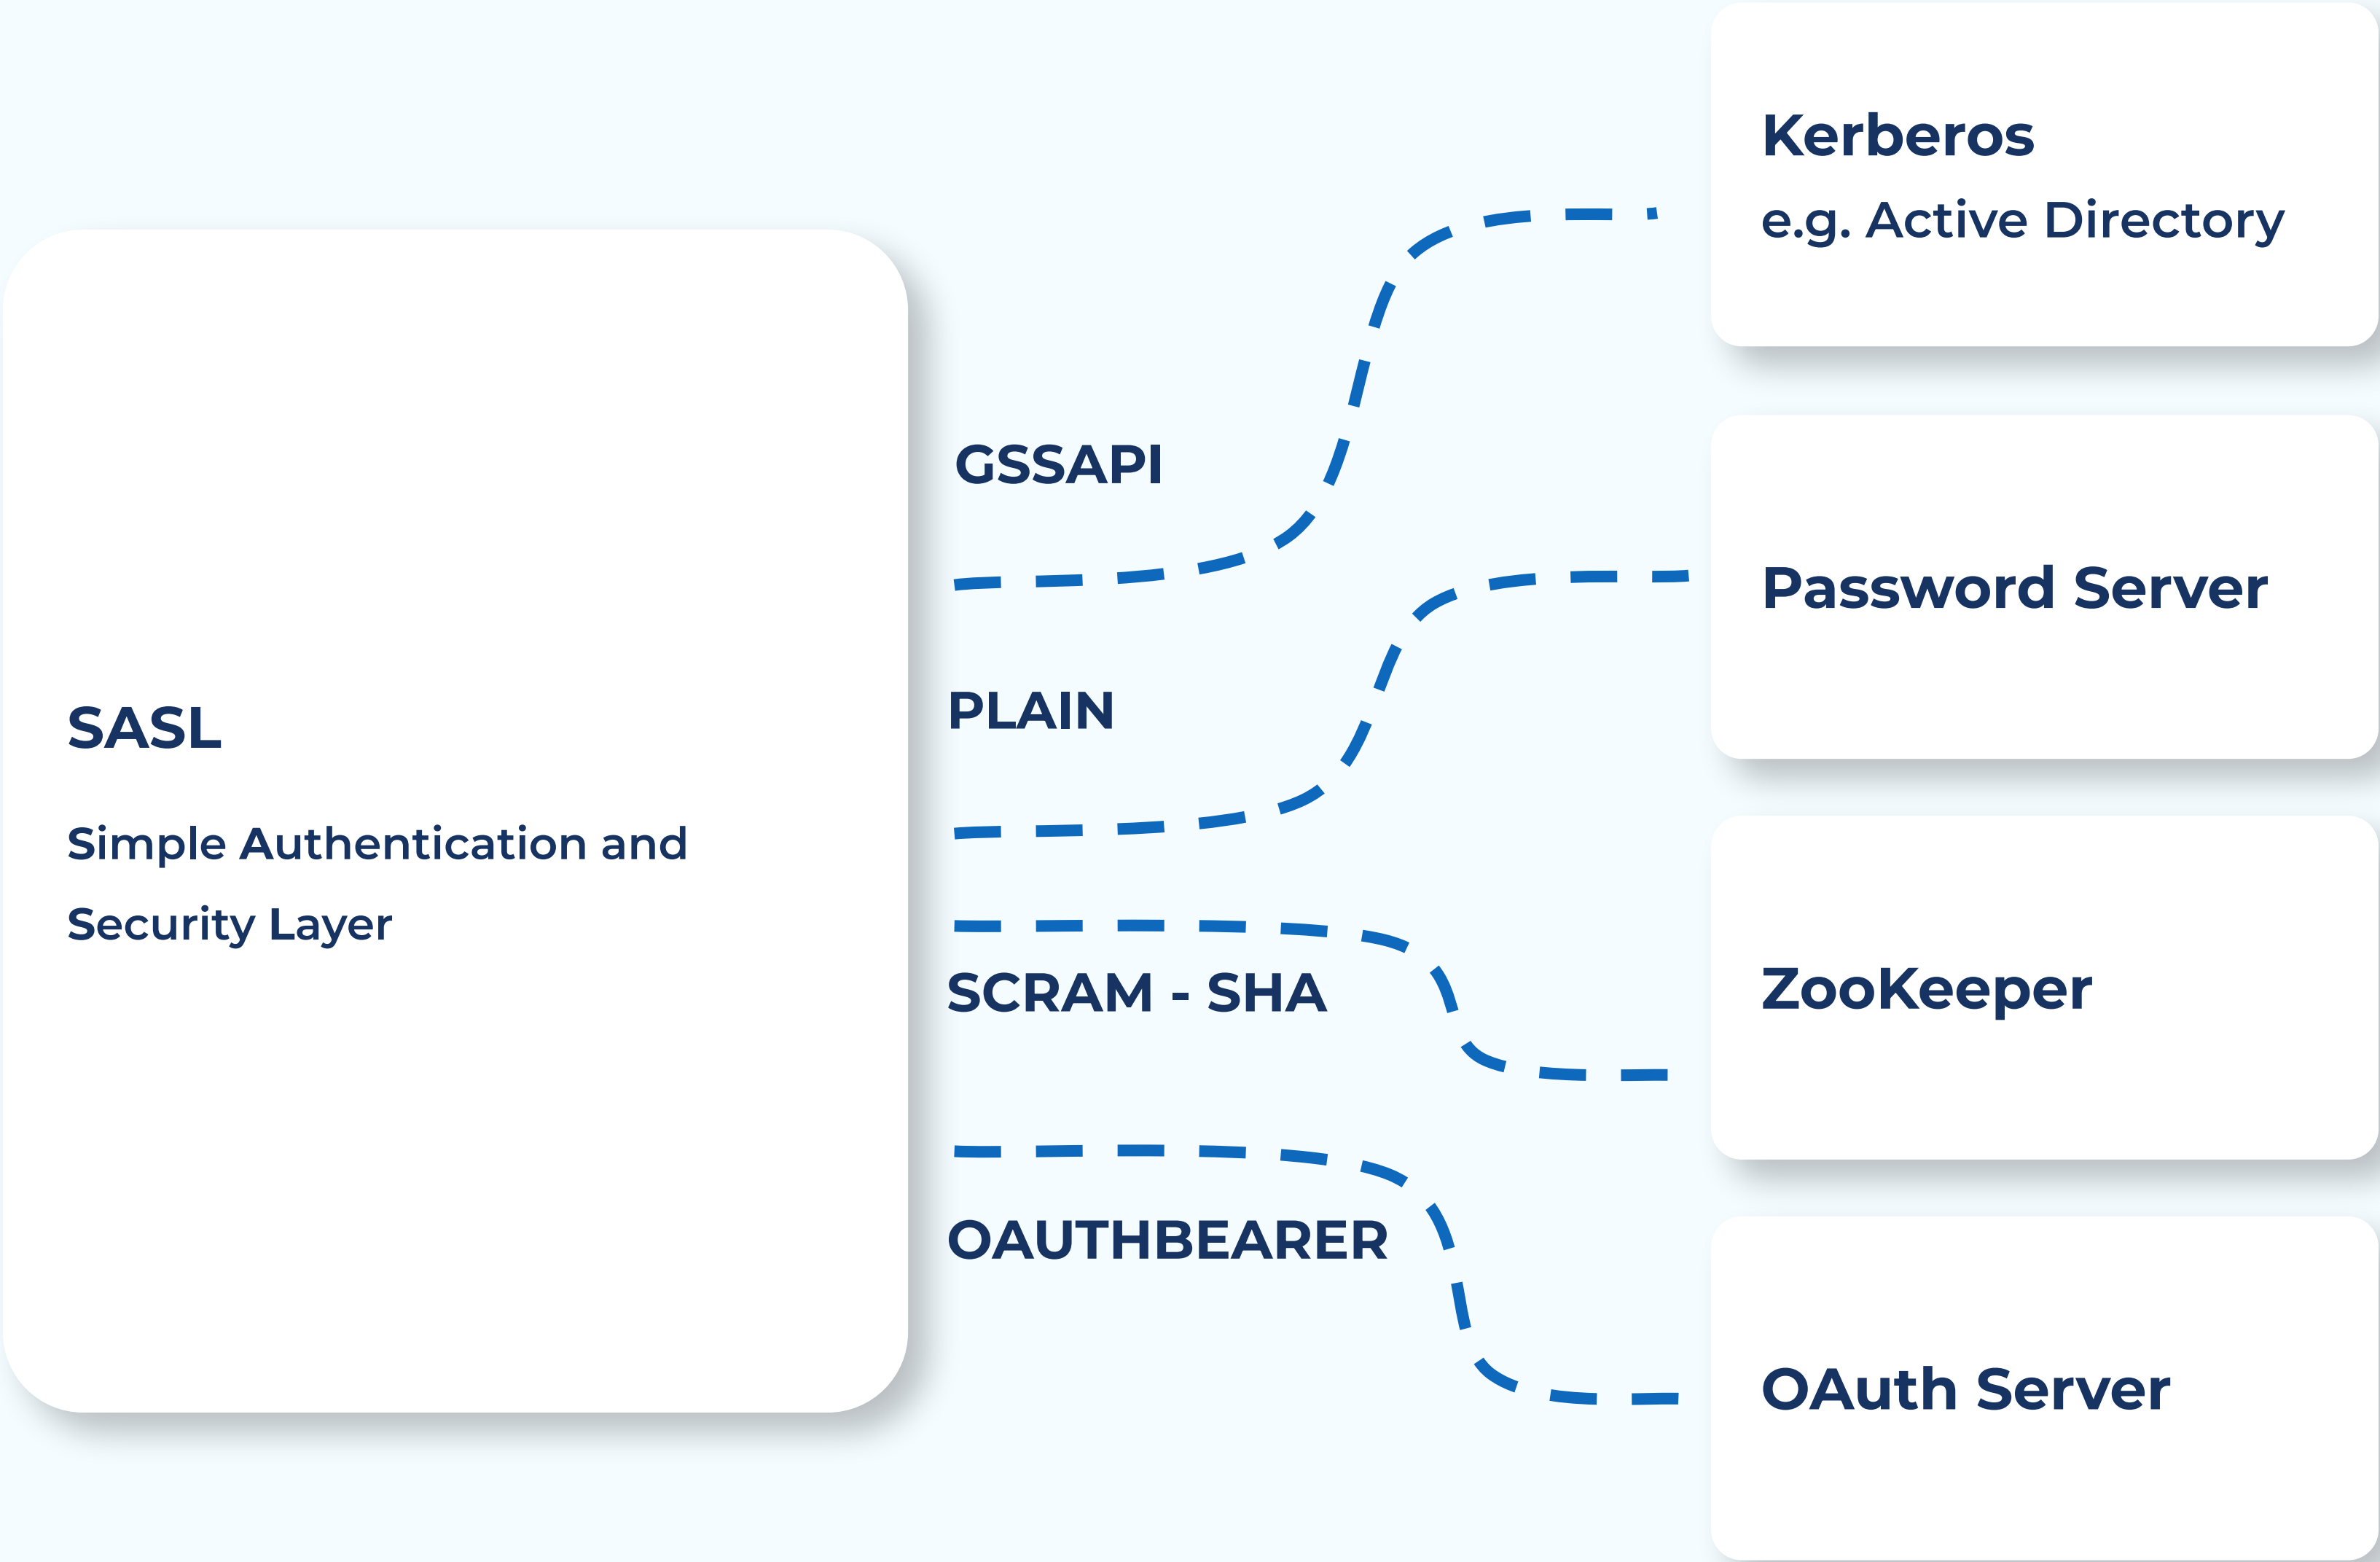
security.inter.broker.protocol

SASL

**Simple Authentication and
Security Layer**







Considerations



- Use filesystem permissions to restrict access to files containing security information

Considerations



- Use filesystem permissions to restrict access to files containing security information
- Avoid storing passwords in plaintext anywhere on the system

Considerations

- Use filesystem permissions to restrict access to files containing security information
- Avoid storing passwords in plaintext anywhere on the system
- Use disk encryption or a secure credential store

Considerations

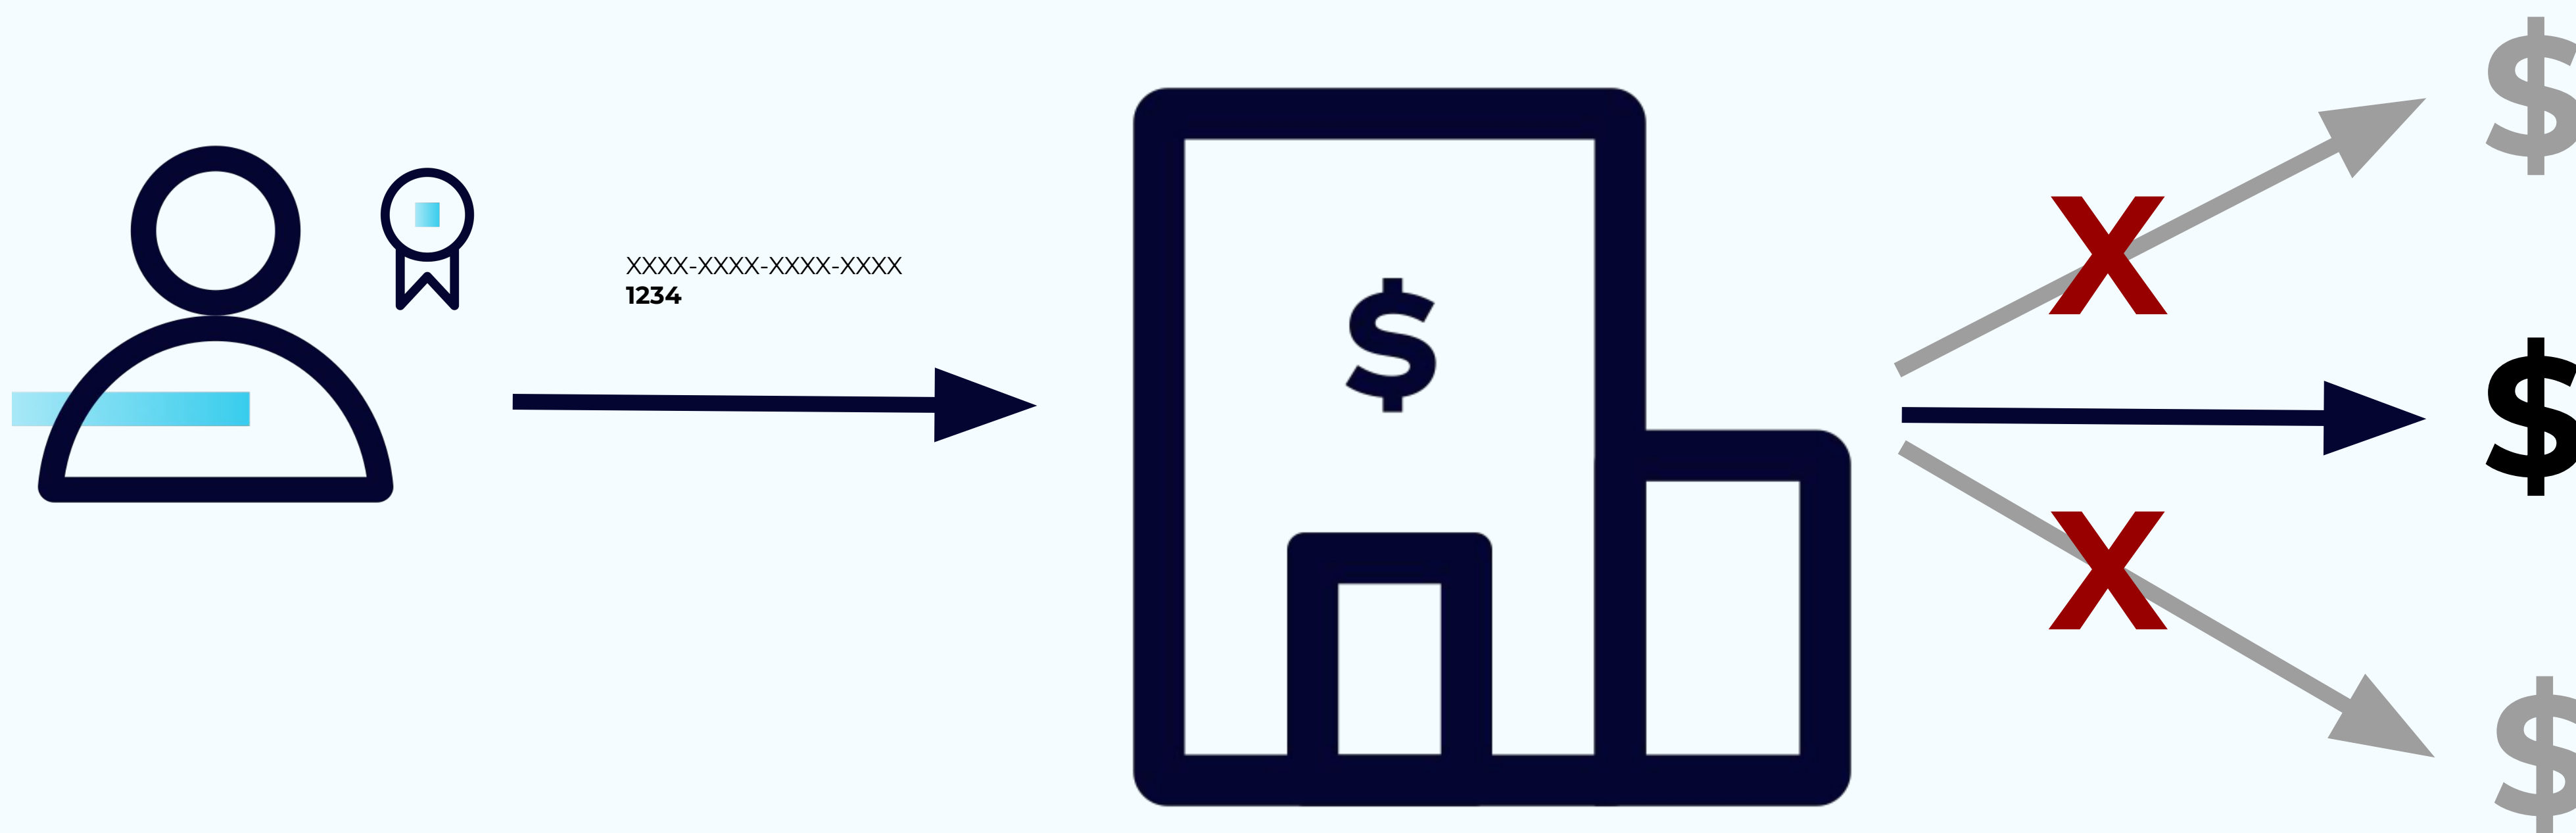
- Use filesystem permissions to restrict access to files containing security information
- Avoid storing passwords in plaintext anywhere on the system
- Use disk encryption or a secure credential store
- Use quotas to limit the impact of malicious client behaviours

Considerations

- Use filesystem permissions to restrict access to files containing security information
- Avoid storing passwords in plaintext anywhere on the system
- Use disk encryption or a secure credential store
- Use quotas to limit the impact of malicious client behaviours
- Apply change control to configuration



Authorization



Access control lists (ACLs)

Principal

Resource type and name

User:Alice has **Allow** permission for **Write** to **Topic:customer**

Permission type

Operation

Access control lists (ACLs)



Pattern type: Literal or Prefixed

Resource name: Supports *

Principal: Supports User:*

Host: Source IP or *

Creating ACLs



```
kafka-acls --bootstrap-server localhost:9092 \  
  --command-config adminclient-configs.conf \  
  --add \  
  --allow-principal User:alice \  
  --allow-principal User:fred \  
  --operation read \  
  --operation write \  
  --topic finance
```

Real-world principal names



SSL certificate subject name

User:CN=tracking.example.org,OU=TEST,O=Sales,L=PaloAlto,ST=Ca,C=US

Kerberos

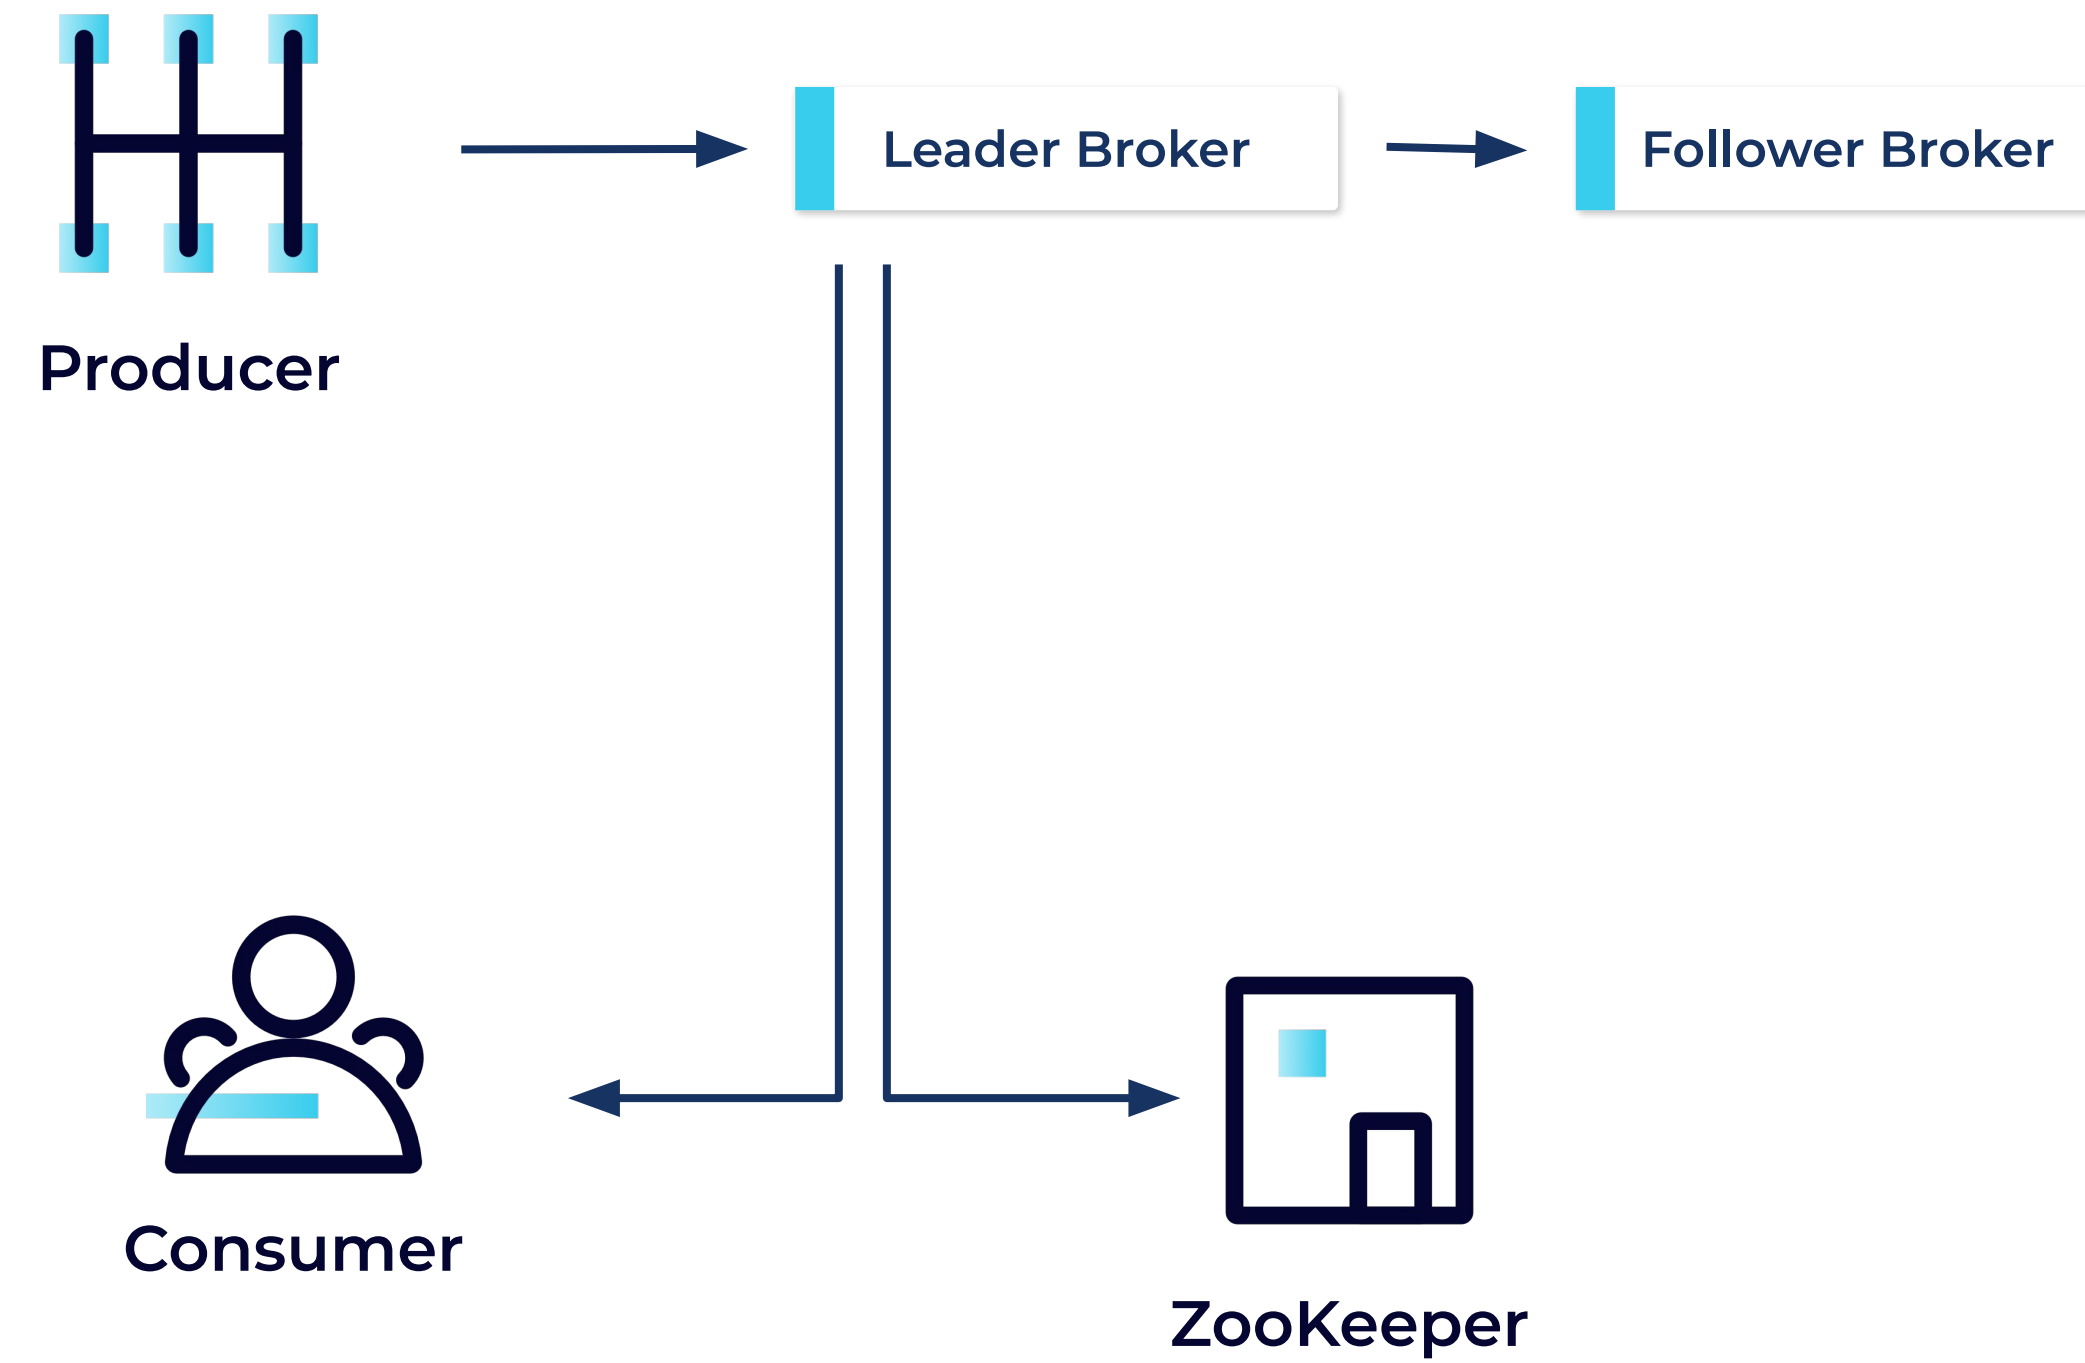
User:tracking

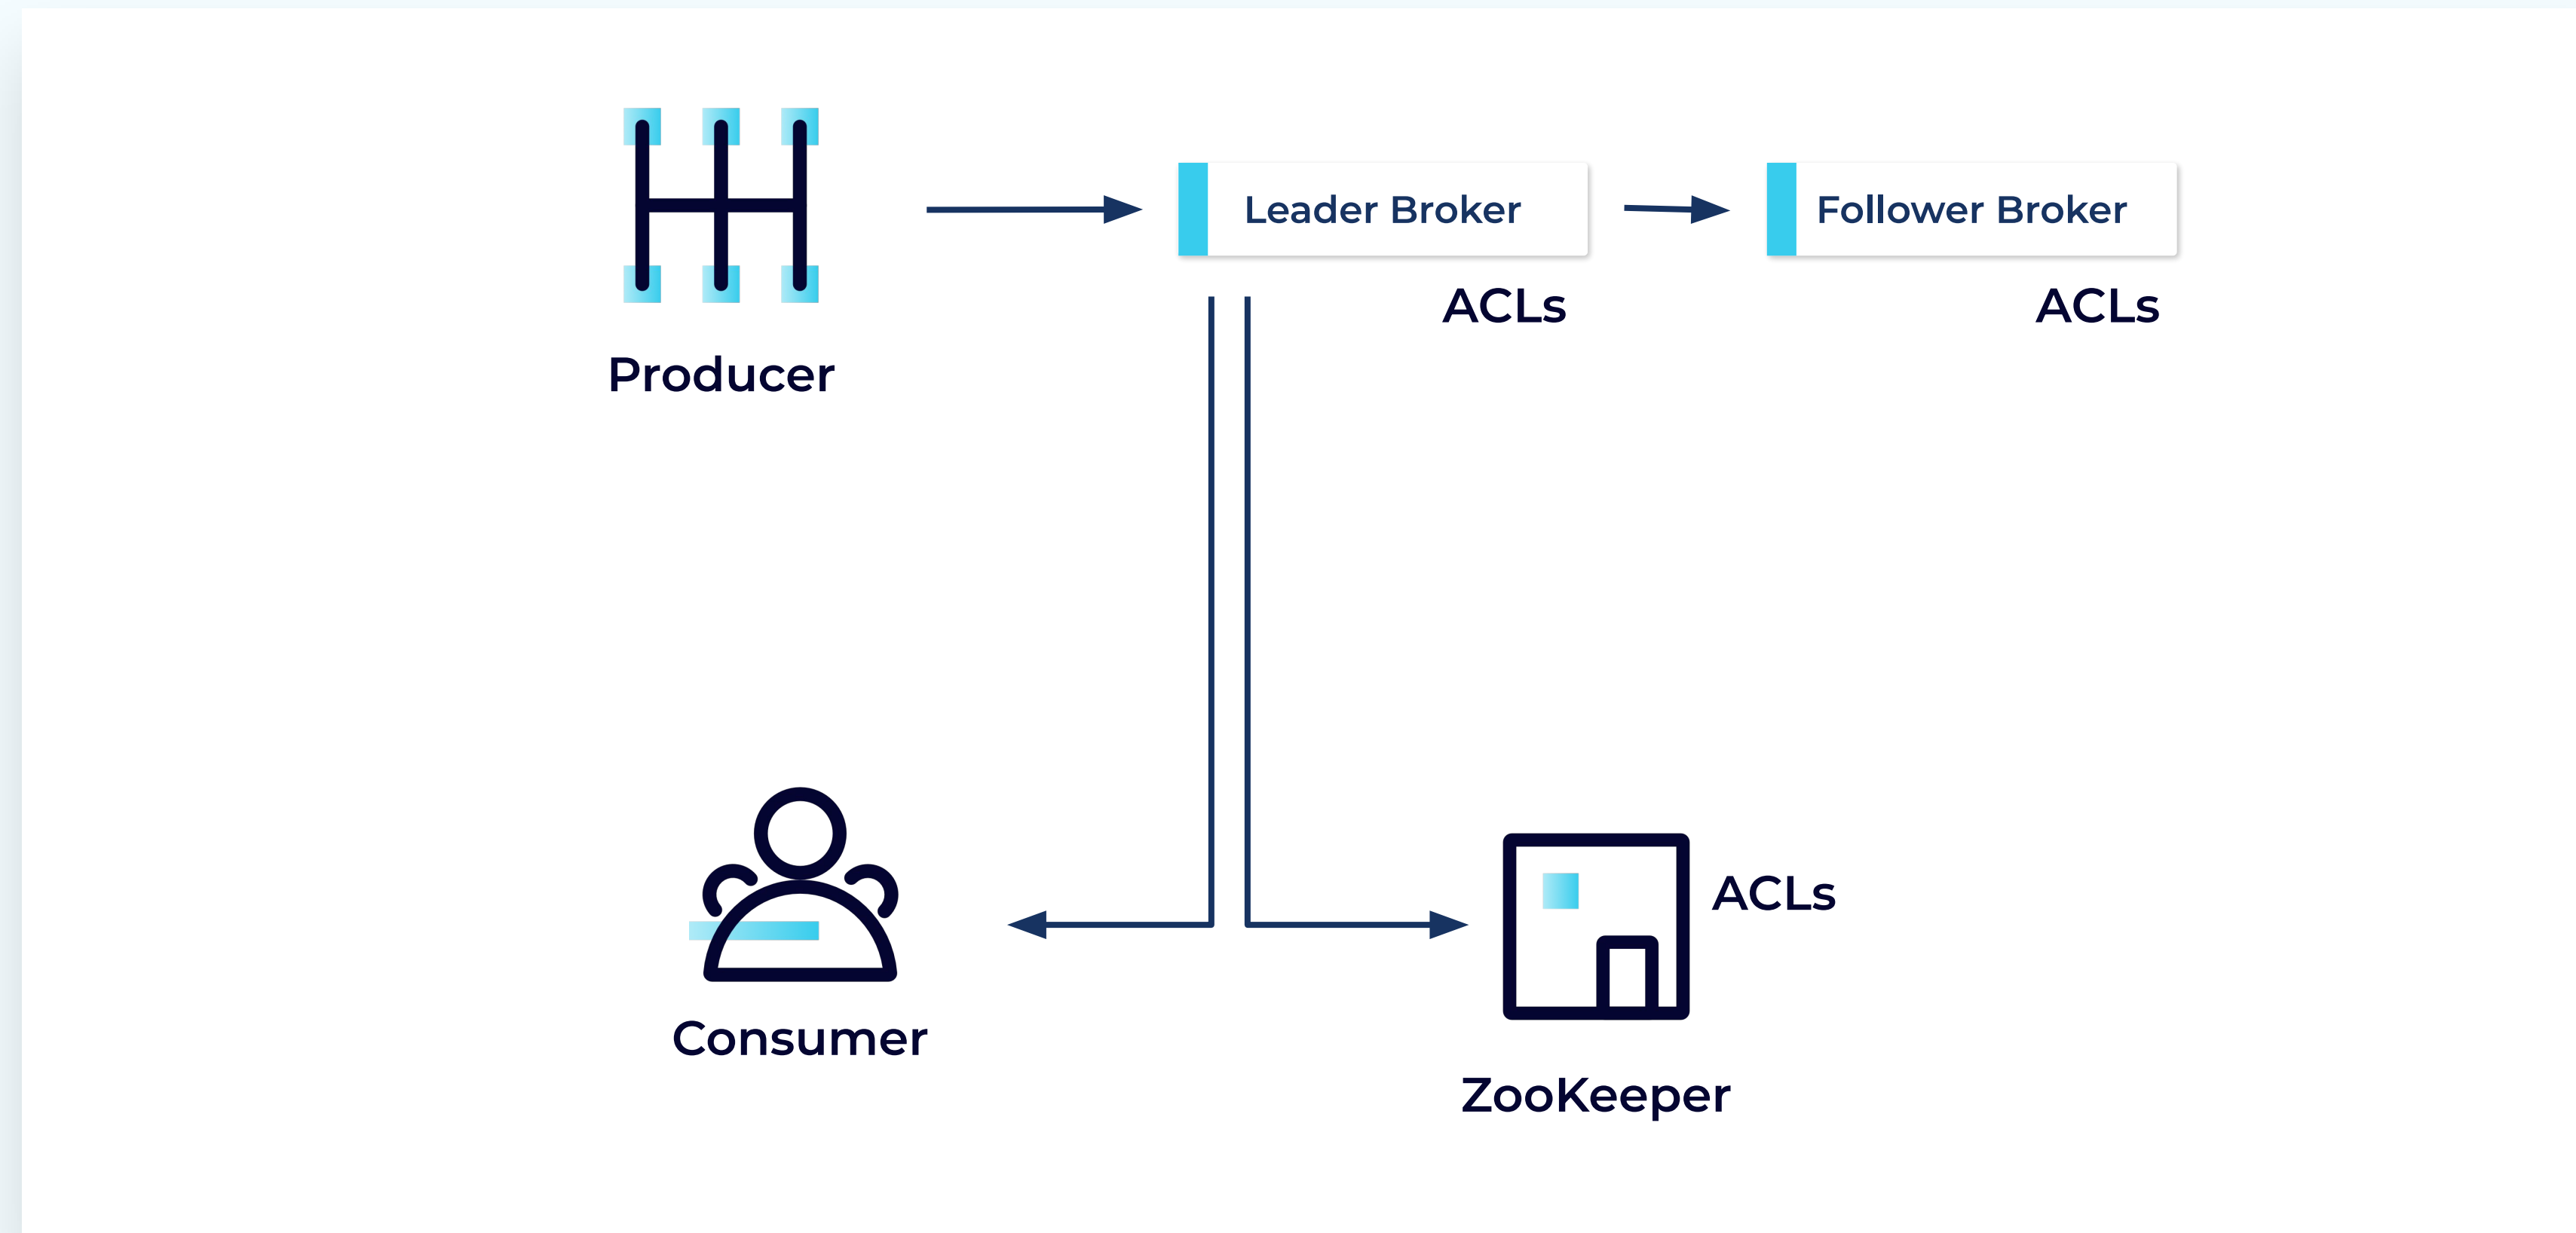
SSL subject names to principals

ssl.principal.mapping.rules

Kerberos identities to preferred principal format

sasl.kerberos.principal.to.local.rules





How Kafka applies ACLs

Broker configuration (with ZooKeeper)

authorizer.class.name=kafka.security.authorizer.AclAuthorizer

How Kafka applies ACLs

Broker configuration (with ZooKeeper)

authorizer.class.name=kafka.security.authorizer.AclAuthorizer

KRaft (no ZooKeeper)

Uses StandardAuthorizer by default

ACLs stored in __cluster_metadata topic

How Kafka applies ACLs

Broker configuration (with ZooKeeper)

authorizer.class.name=kafka.security.authorizer.AclAuthorizer

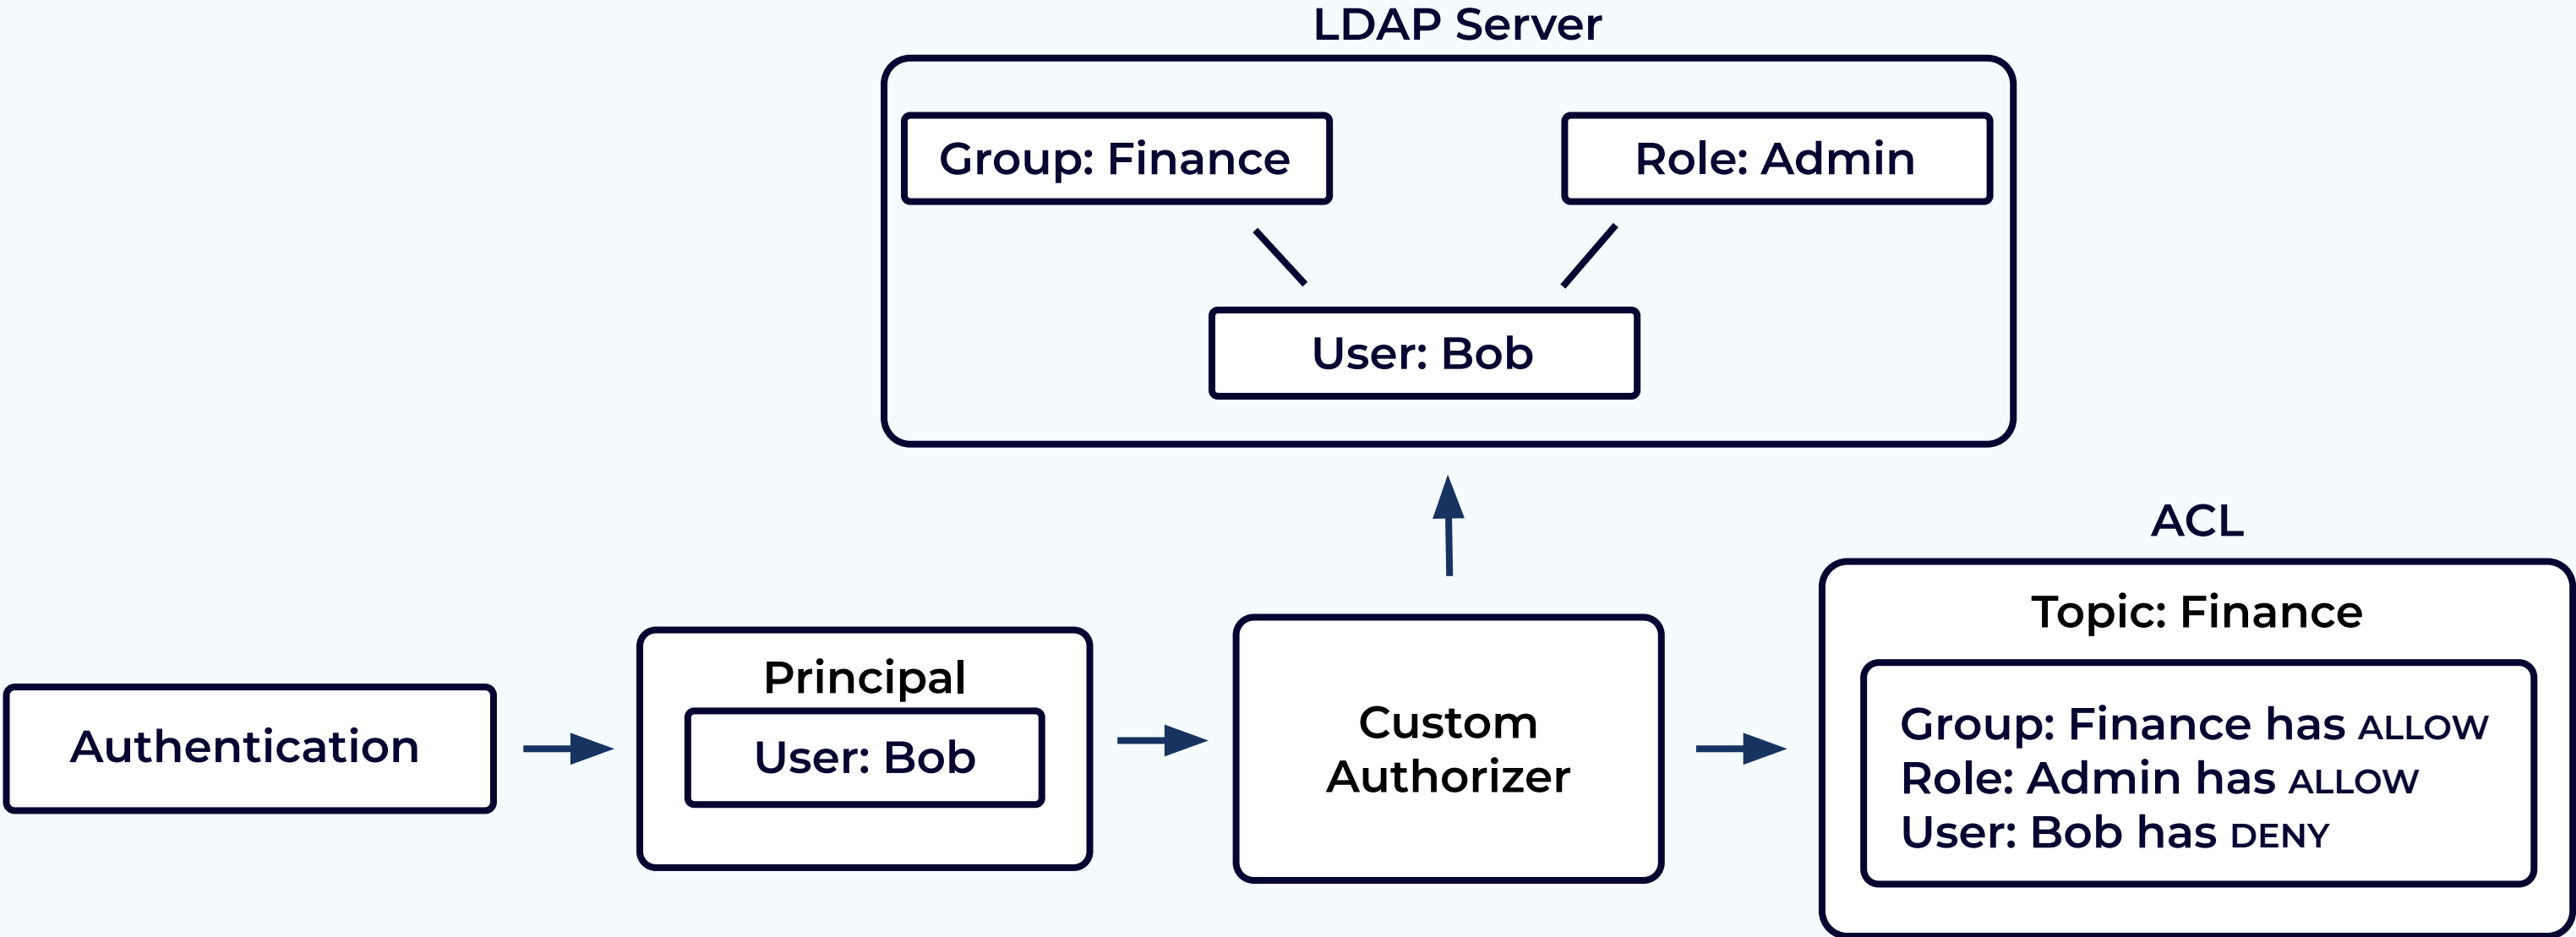
KRaft (no ZooKeeper)

Uses StandardAuthorizer by default

ACLs stored in __cluster_metadata topic

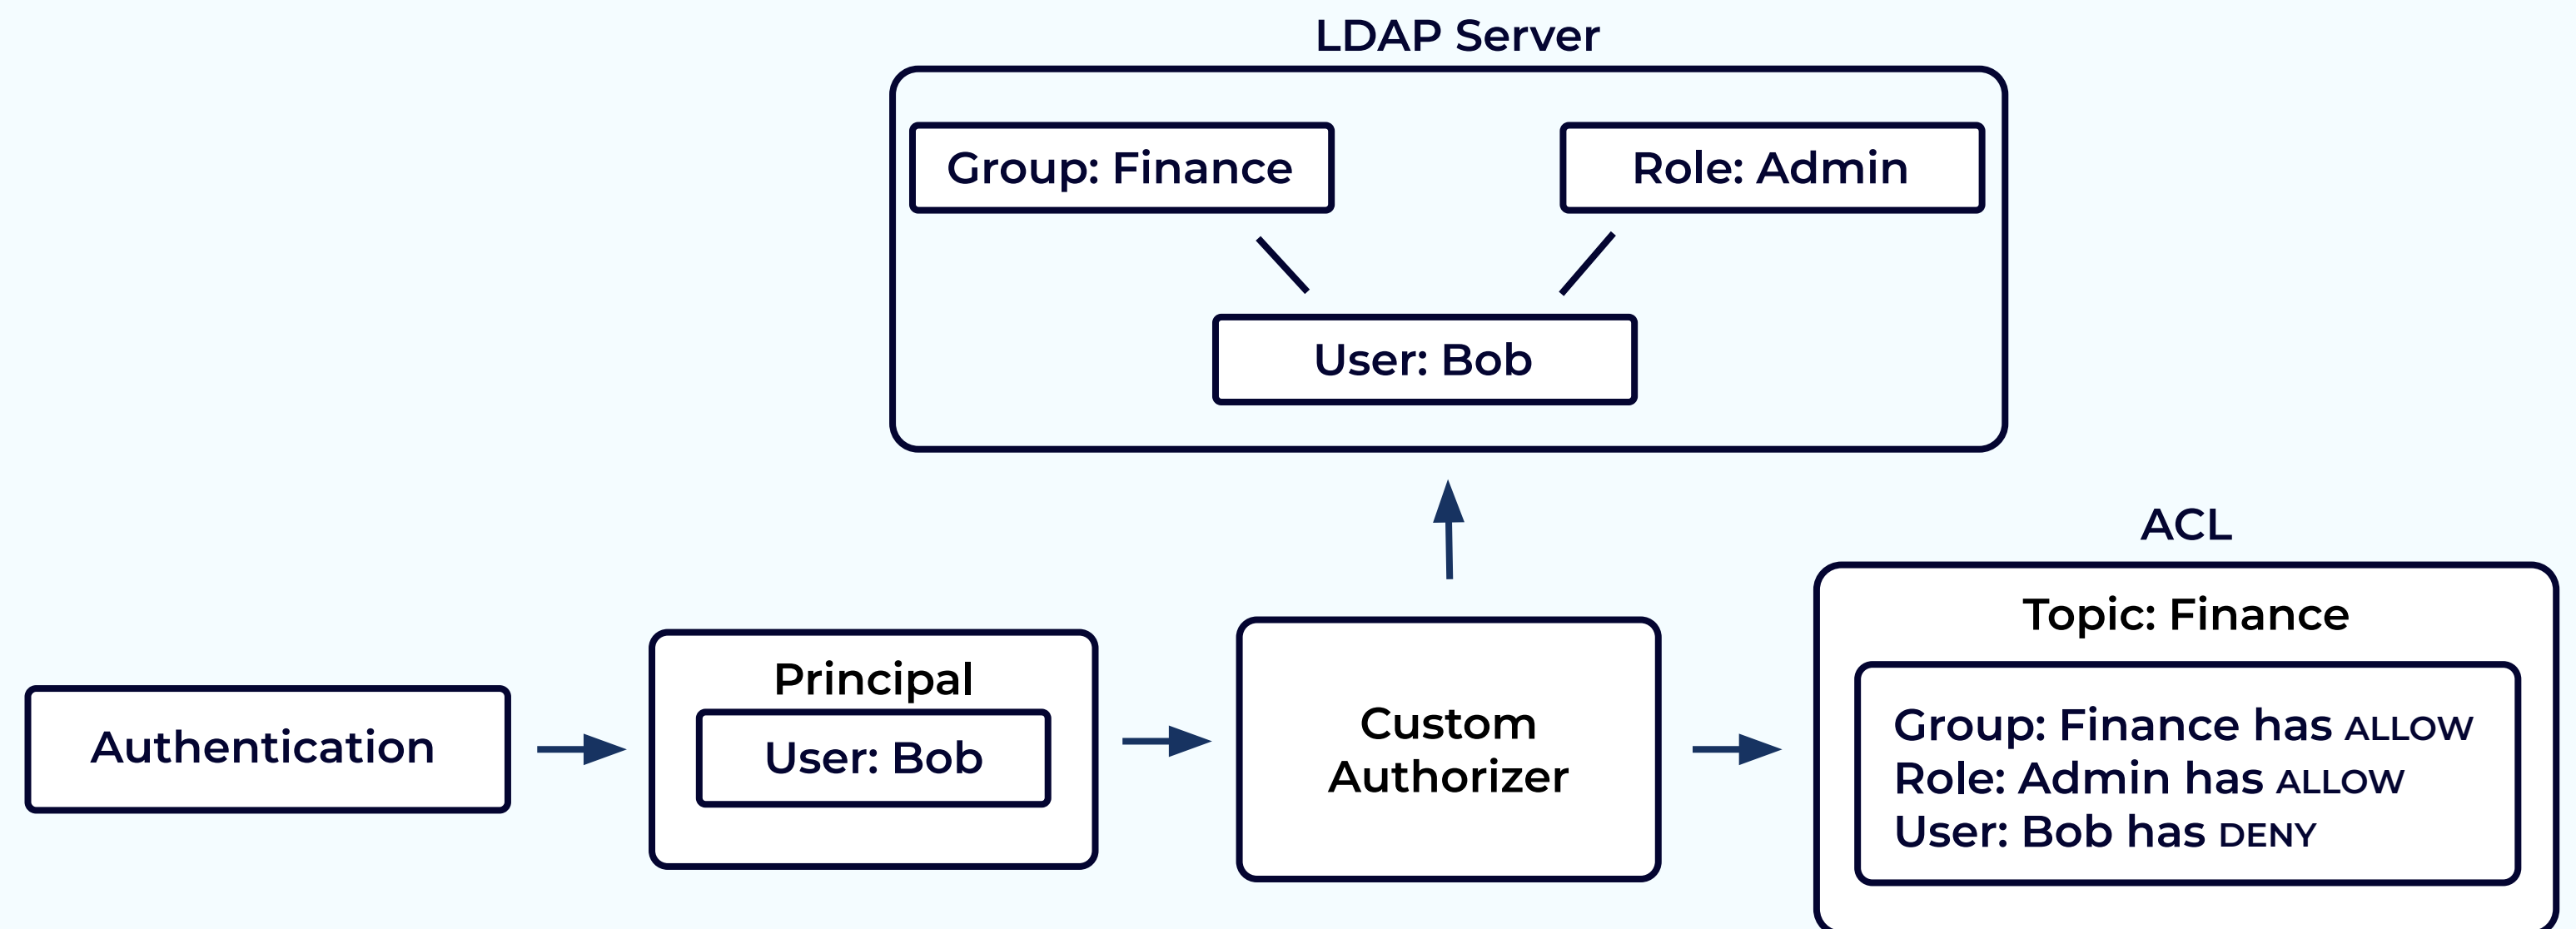
Deny > Allow

Customizing access control



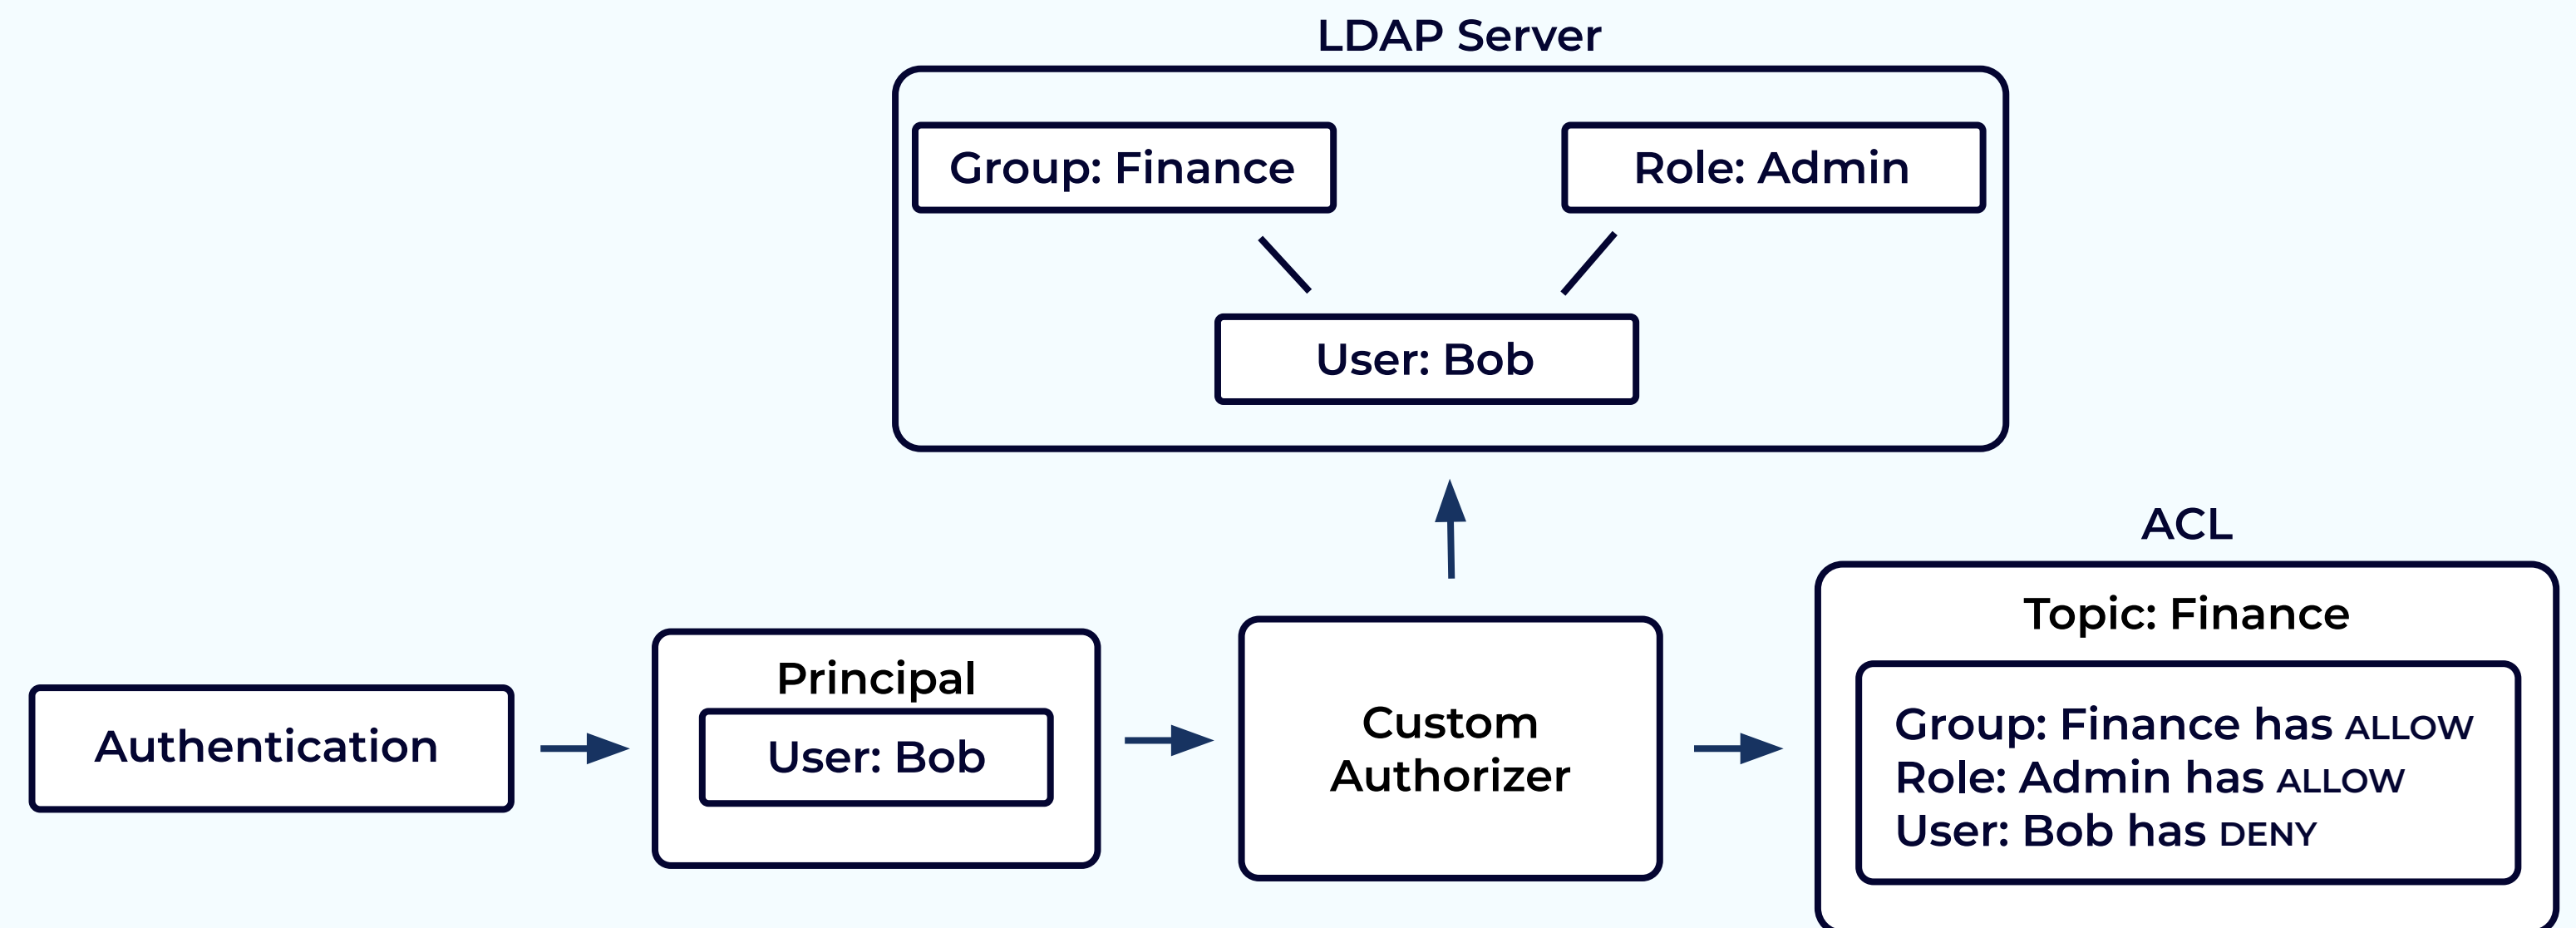
Customizing access control

- Store user, group and role associations in e.g. an external LDAP server



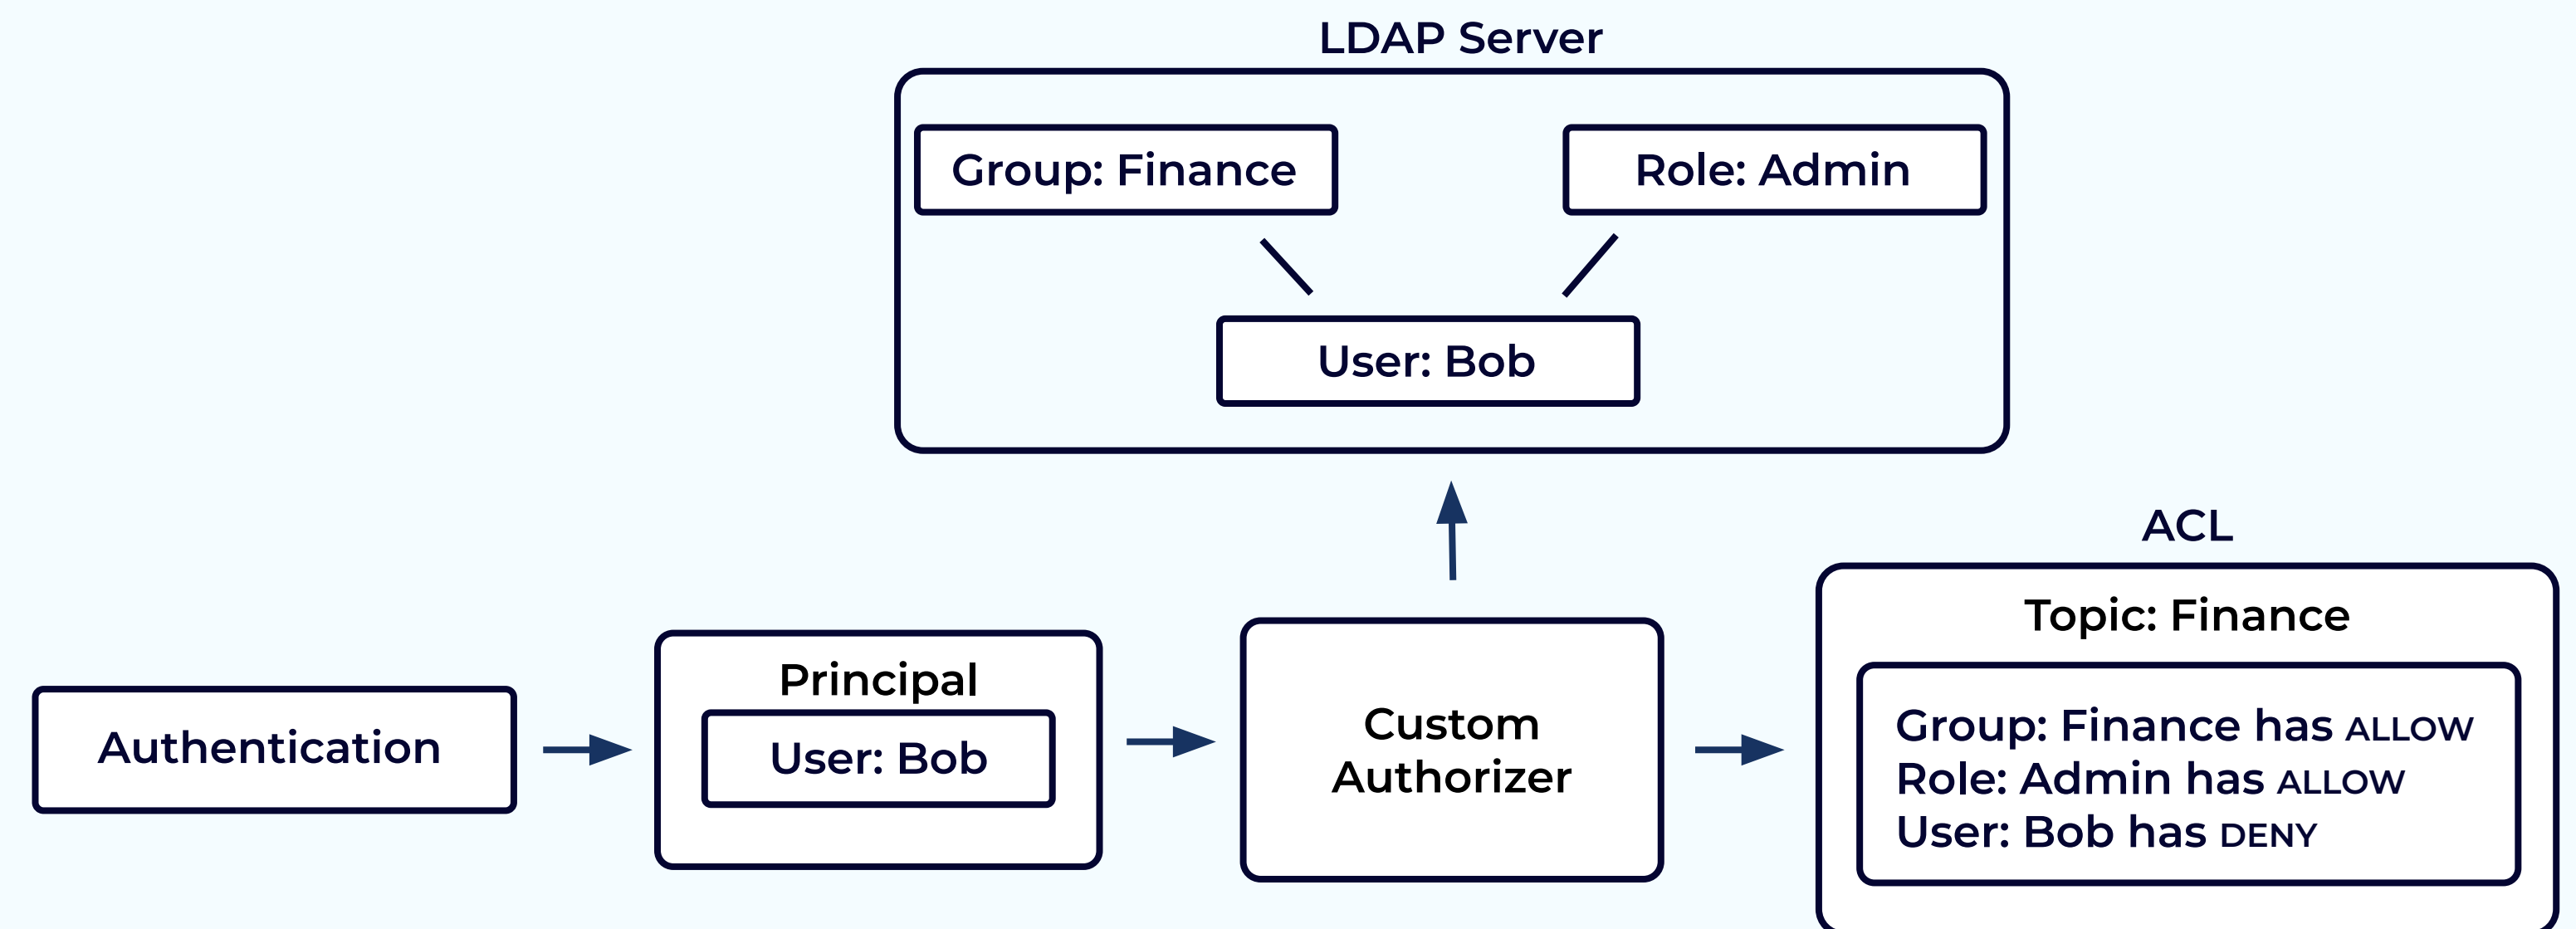
Customizing access control

- Store user, group and role associations in e.g. an external LDAP server
- Add user, group and role permissions to the ACLs stored in ZooKeeper



Customizing access control

- Store user, group and role associations in e.g. an external LDAP server
- Add user, group and role permissions to the ACLs stored in ZooKeeper
- Implement custom Authorizer that fetches a user's groups and roles from LDAP, and resolves against a resource's ACLs.



Considerations



- Use super users and `allow.everyone.if.no.acl.found` with caution – or not at all

Considerations



- Use super users and `allow.everyone.if.no.acl.found` with caution – or not at all
- Don't grant access to ANONYMOUS principal

Considerations



- Use super users and `allow.everyone.if.no.acl.found` with caution – or not at all
- Don't grant access to ANONYMOUS principal
- Automate the process of creating user credentials and assigning ACLs for all environments

Considerations



- Use super users and `allow.everyone.if.no.acl.found` with caution – or not at all
- Don't grant access to ANONYMOUS principal
- Automate the process of creating user credentials and assigning ACLs for all environments
- Adjust `connections.max.reauth.ms` to force connection to reauthenticate at intervals

Considerations

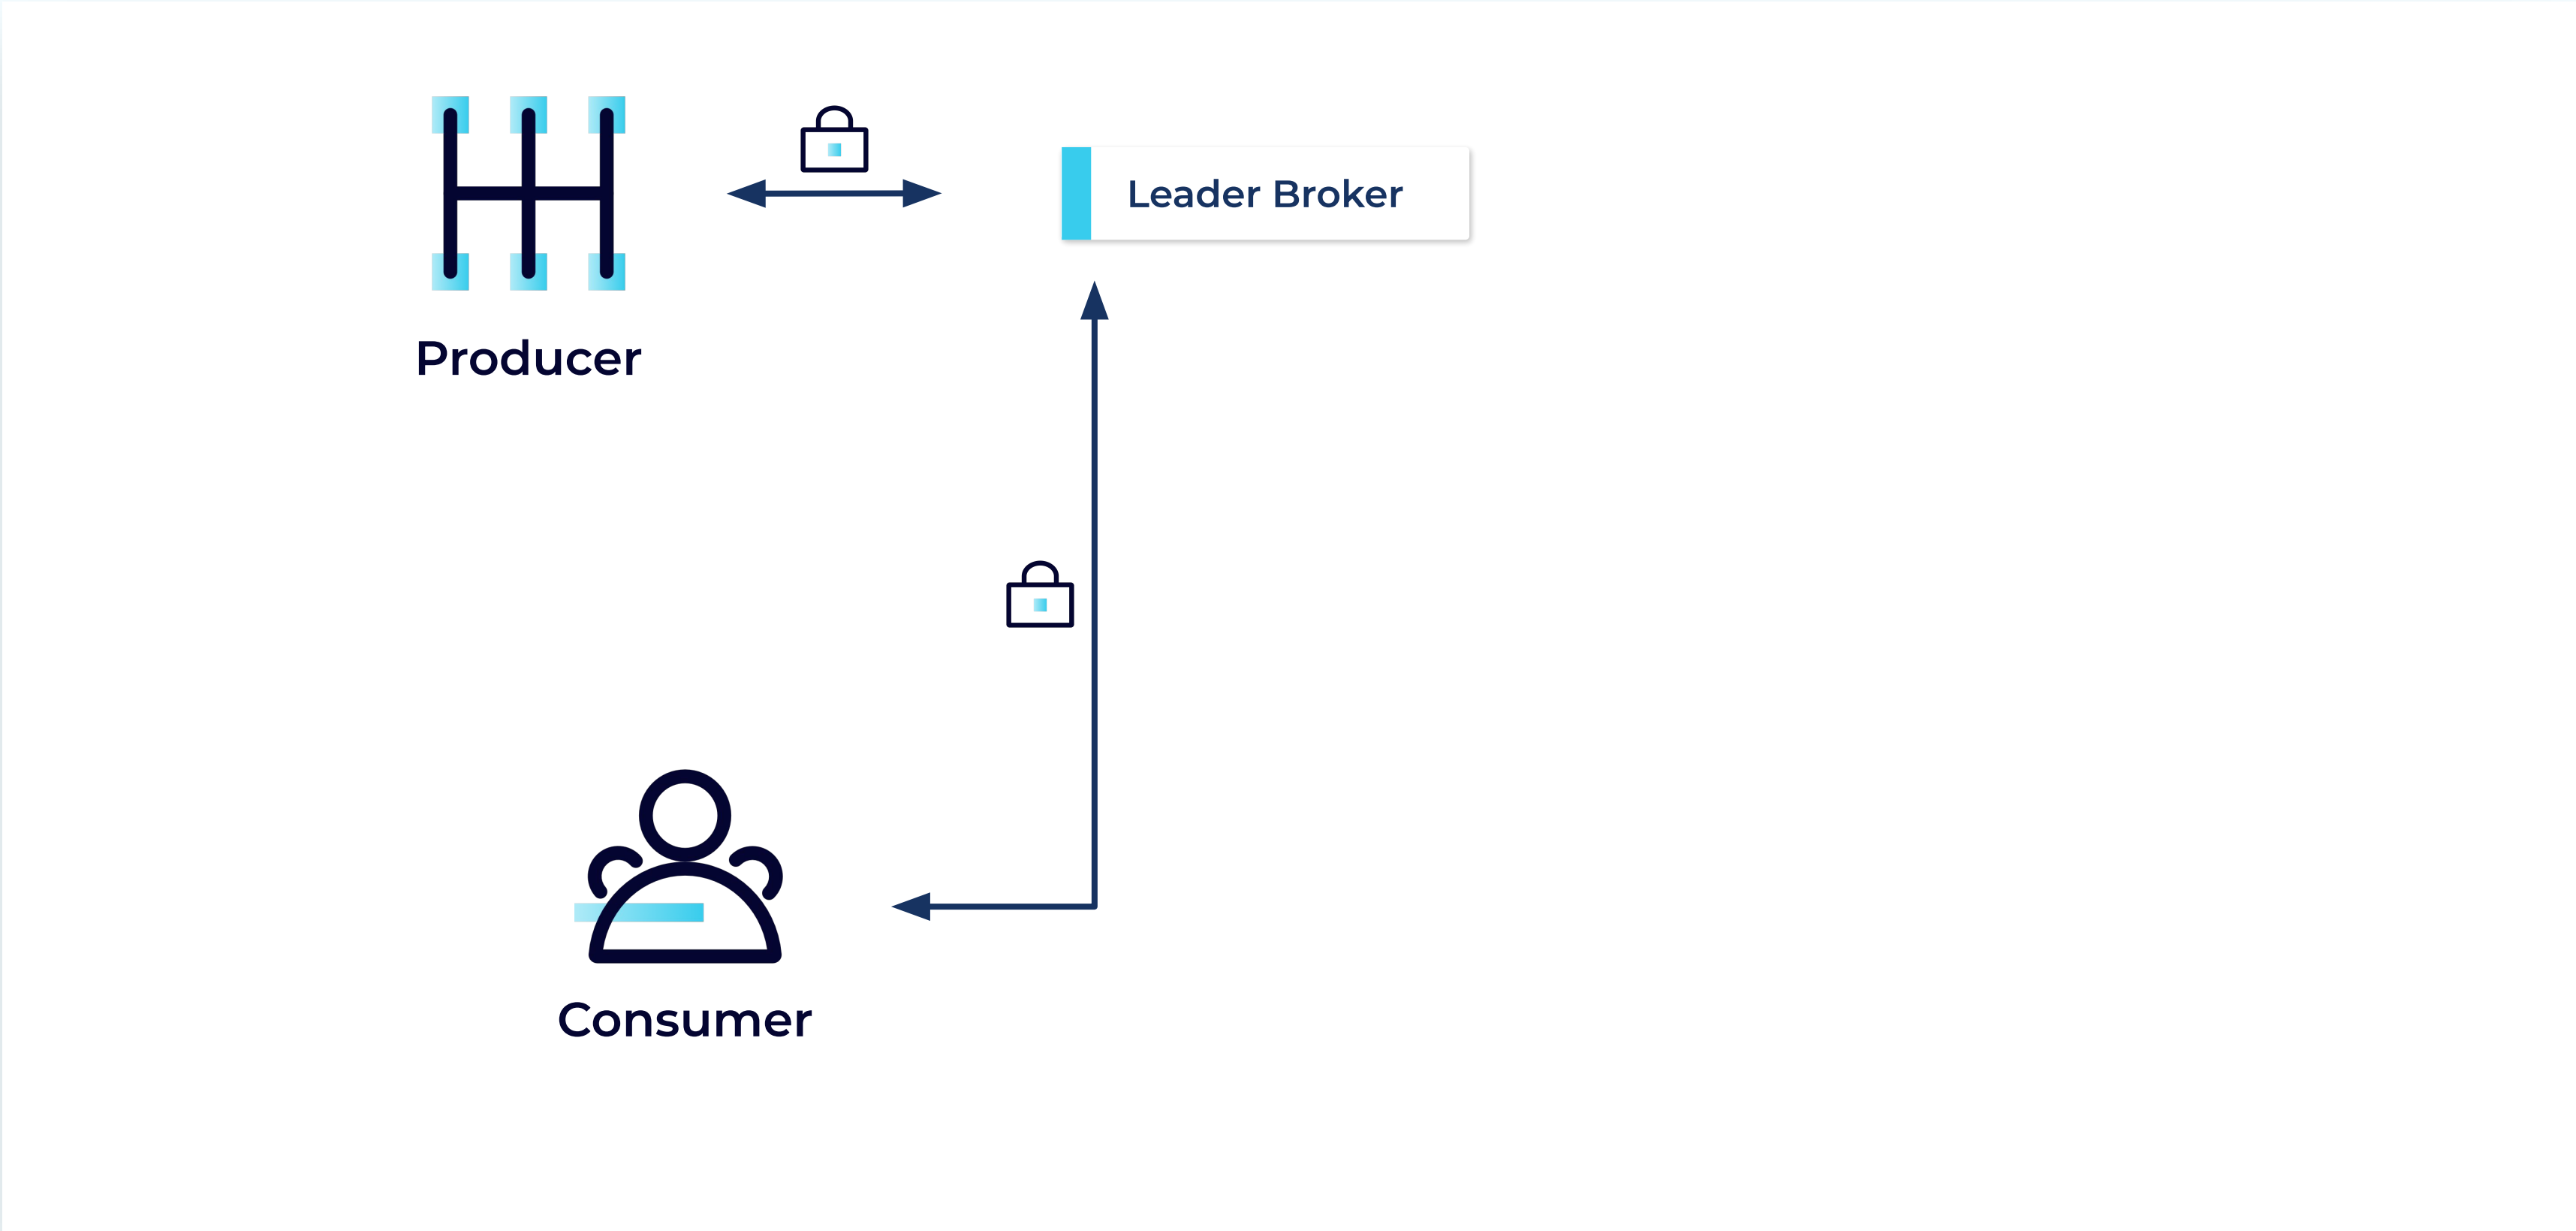


- Use super users and `allow.everyone.if.no.acl.found` with caution – or not at all
- Don't grant access to ANONYMOUS principal
- Automate the process of creating user credentials and assigning ACLs for all environments
- Adjust `connections.max.reauth.ms` to force connection to reauthenticate at intervals
- Use Deny ACLs to prevent actions from compromised users

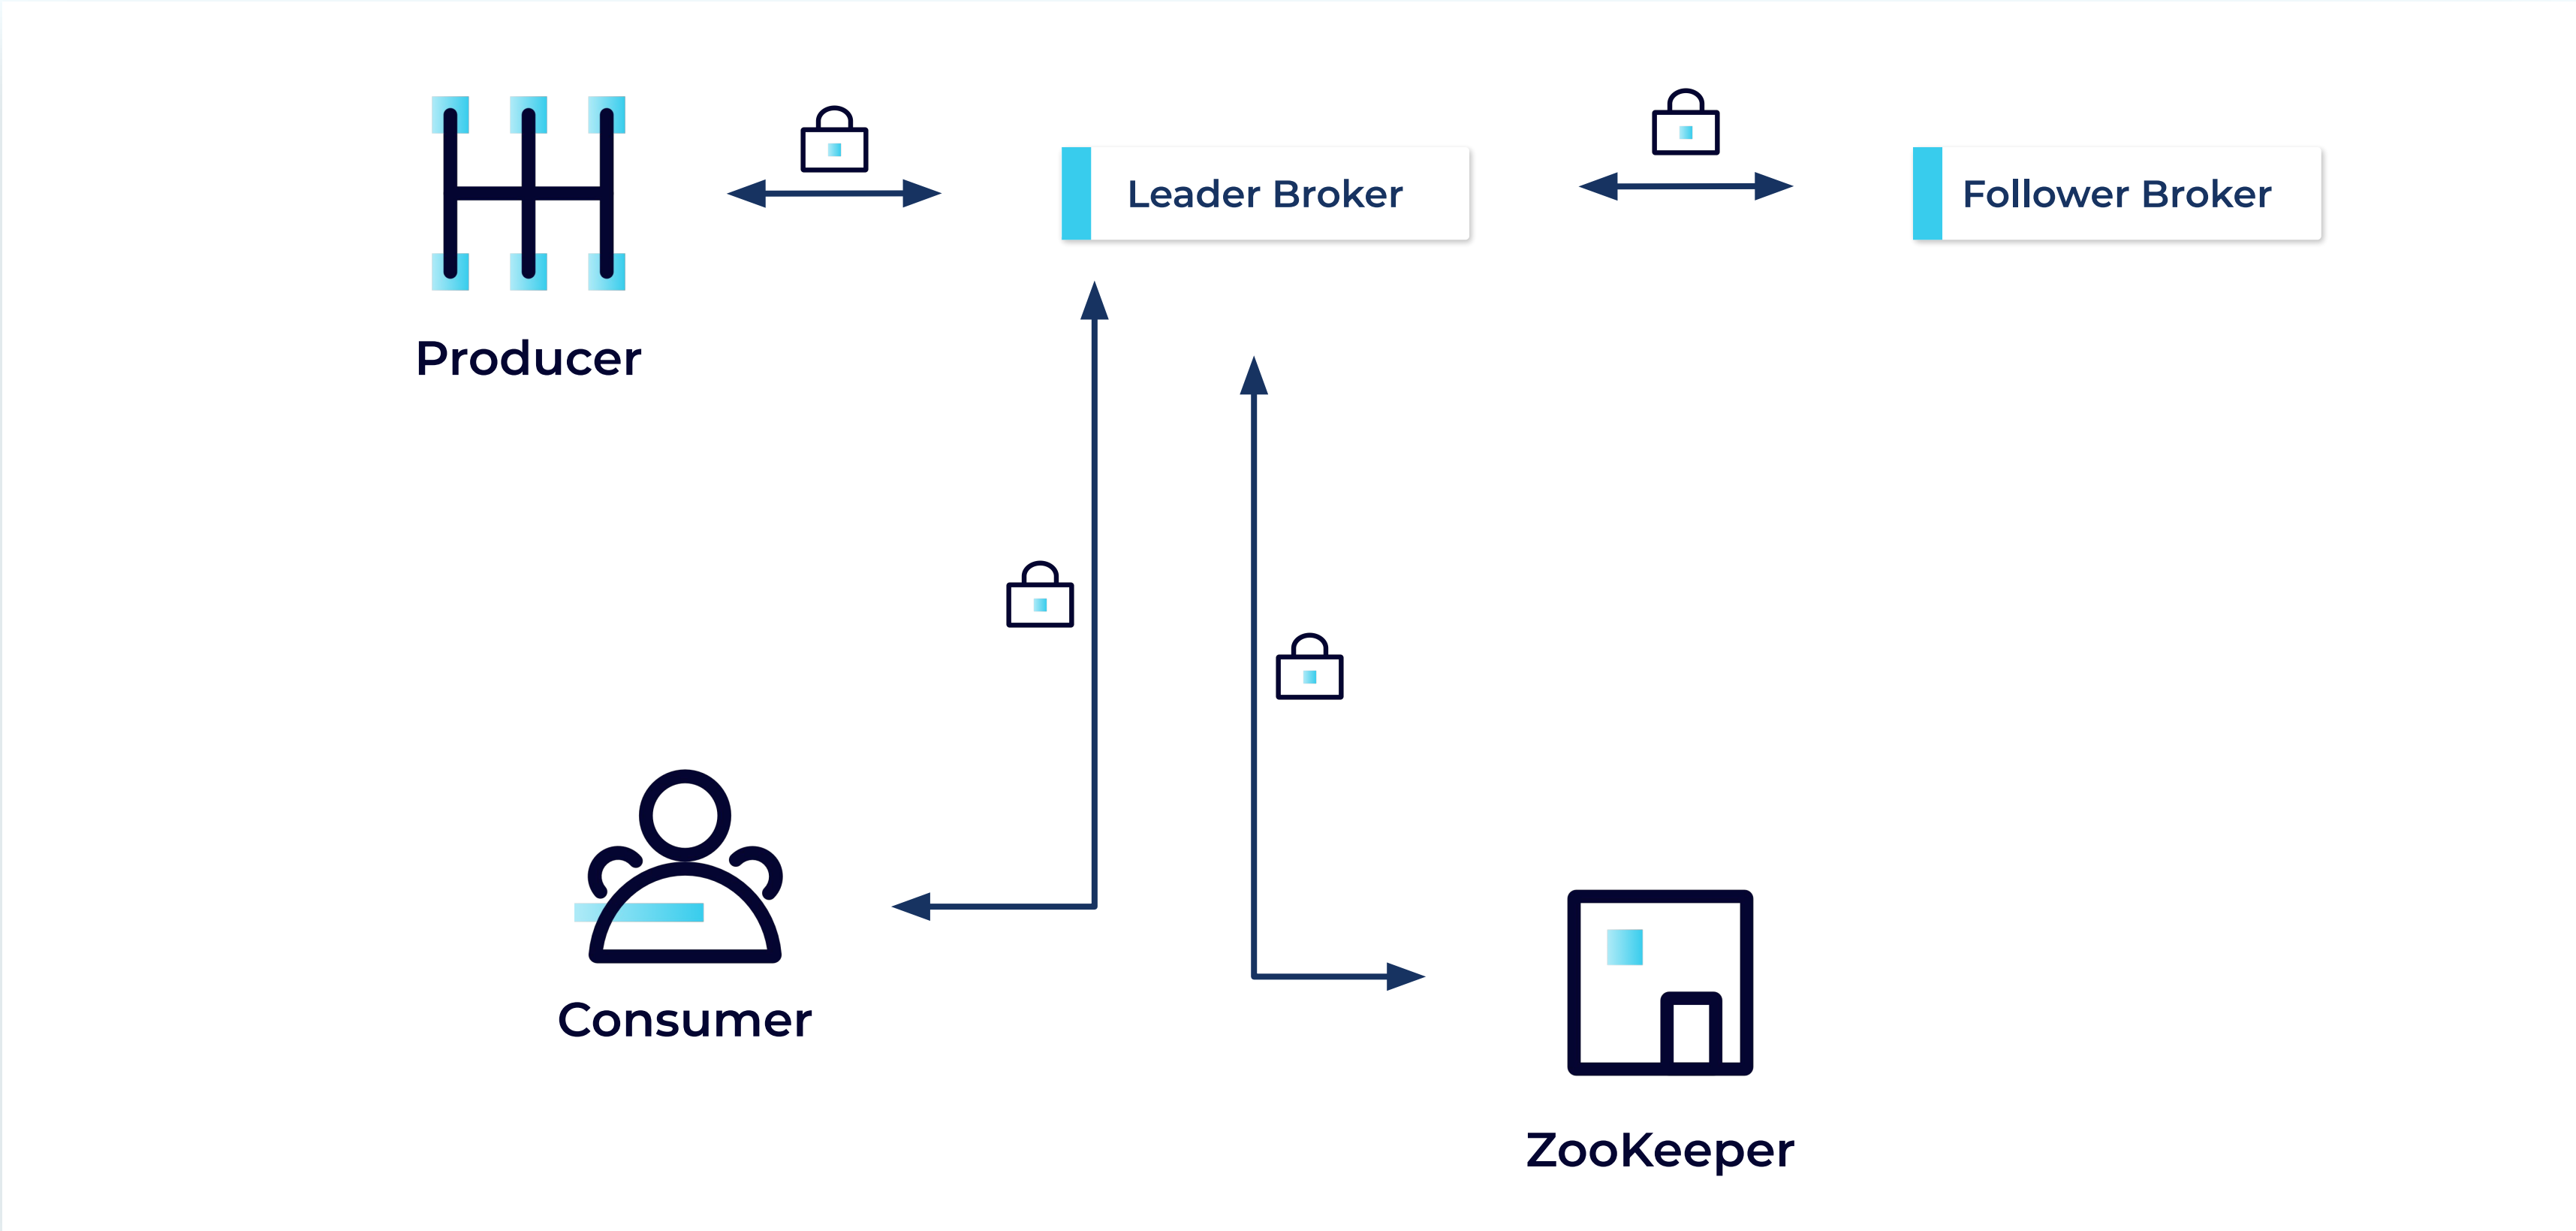


Encryption

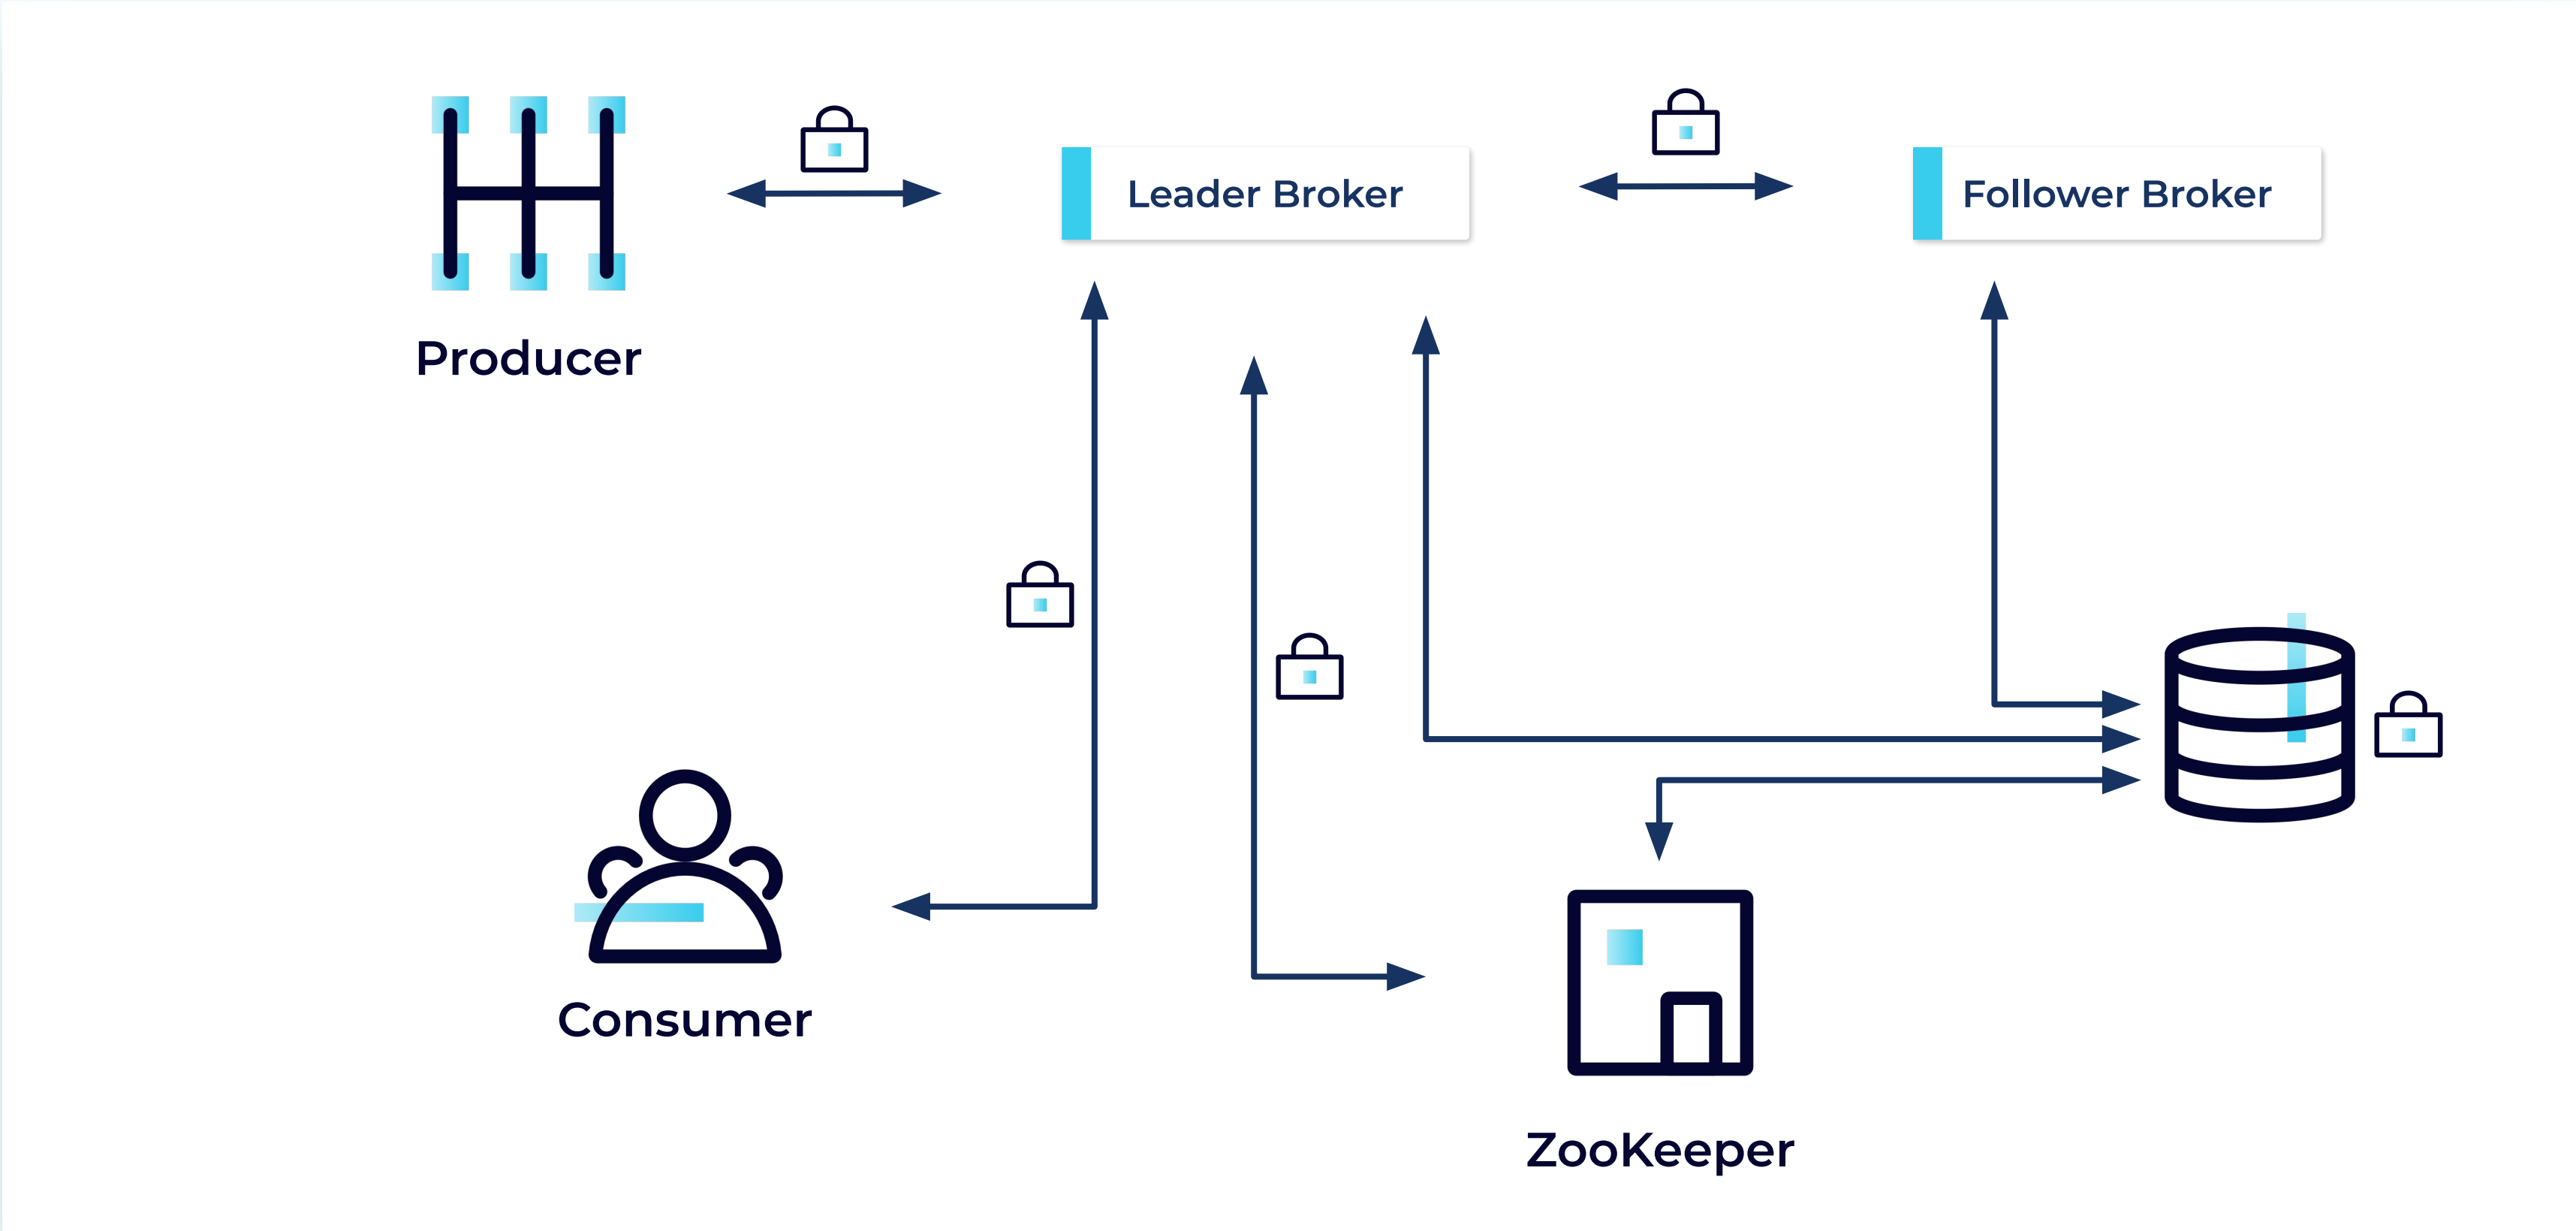
Encryption



Encryption




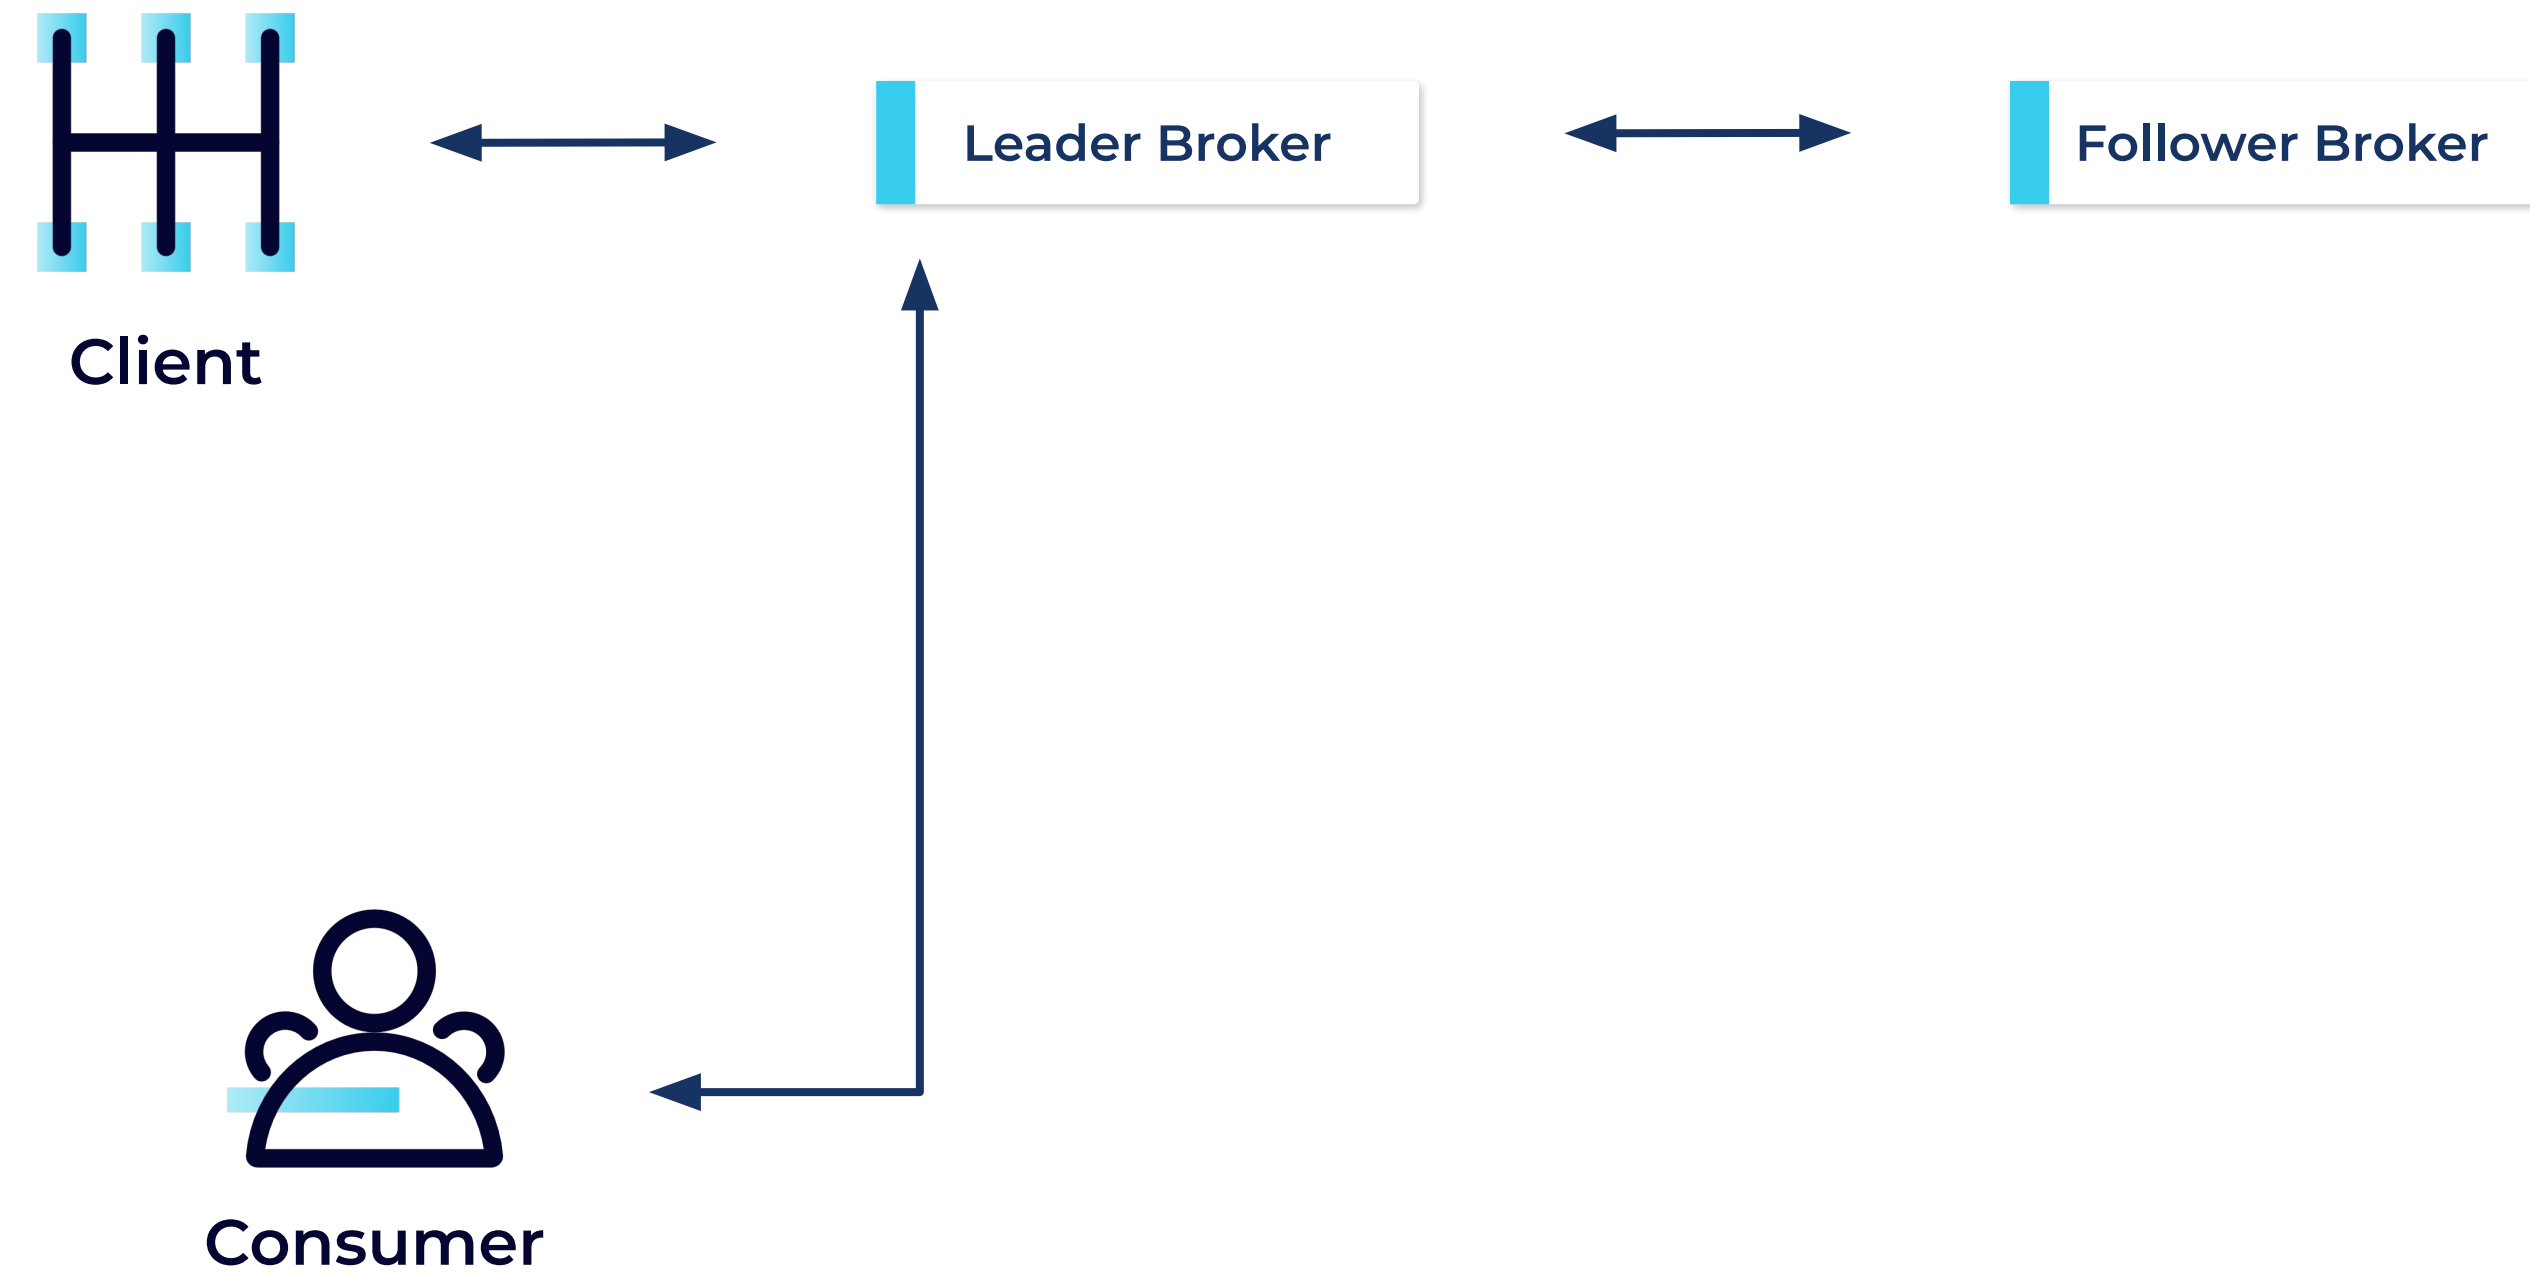
Encryption

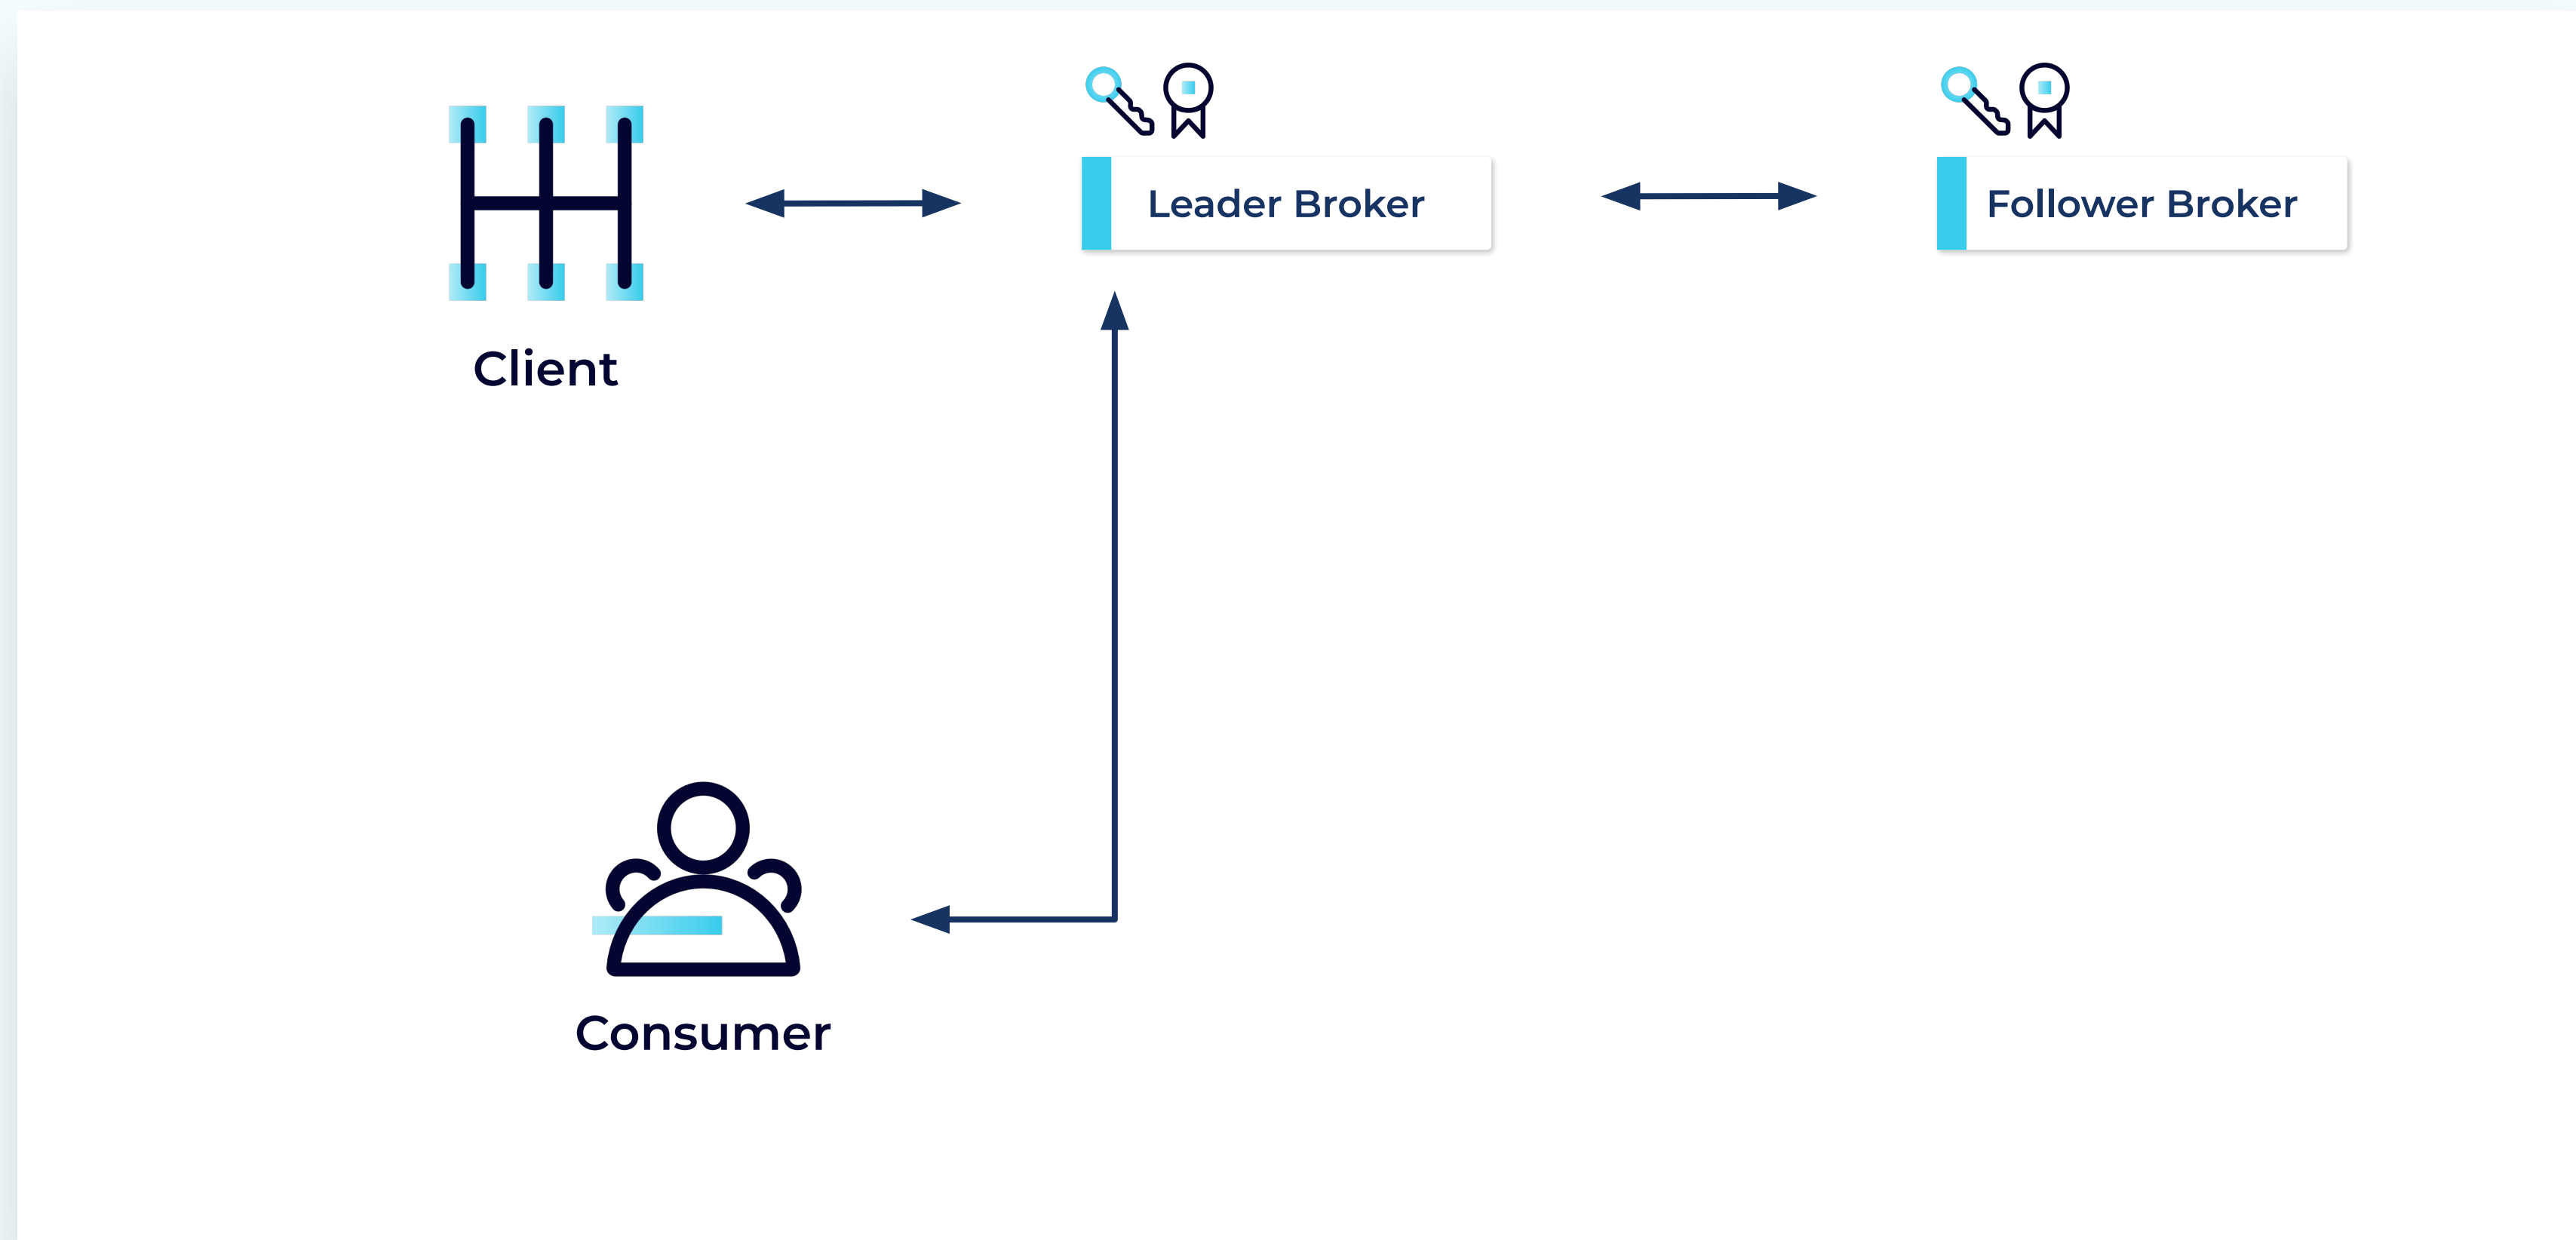


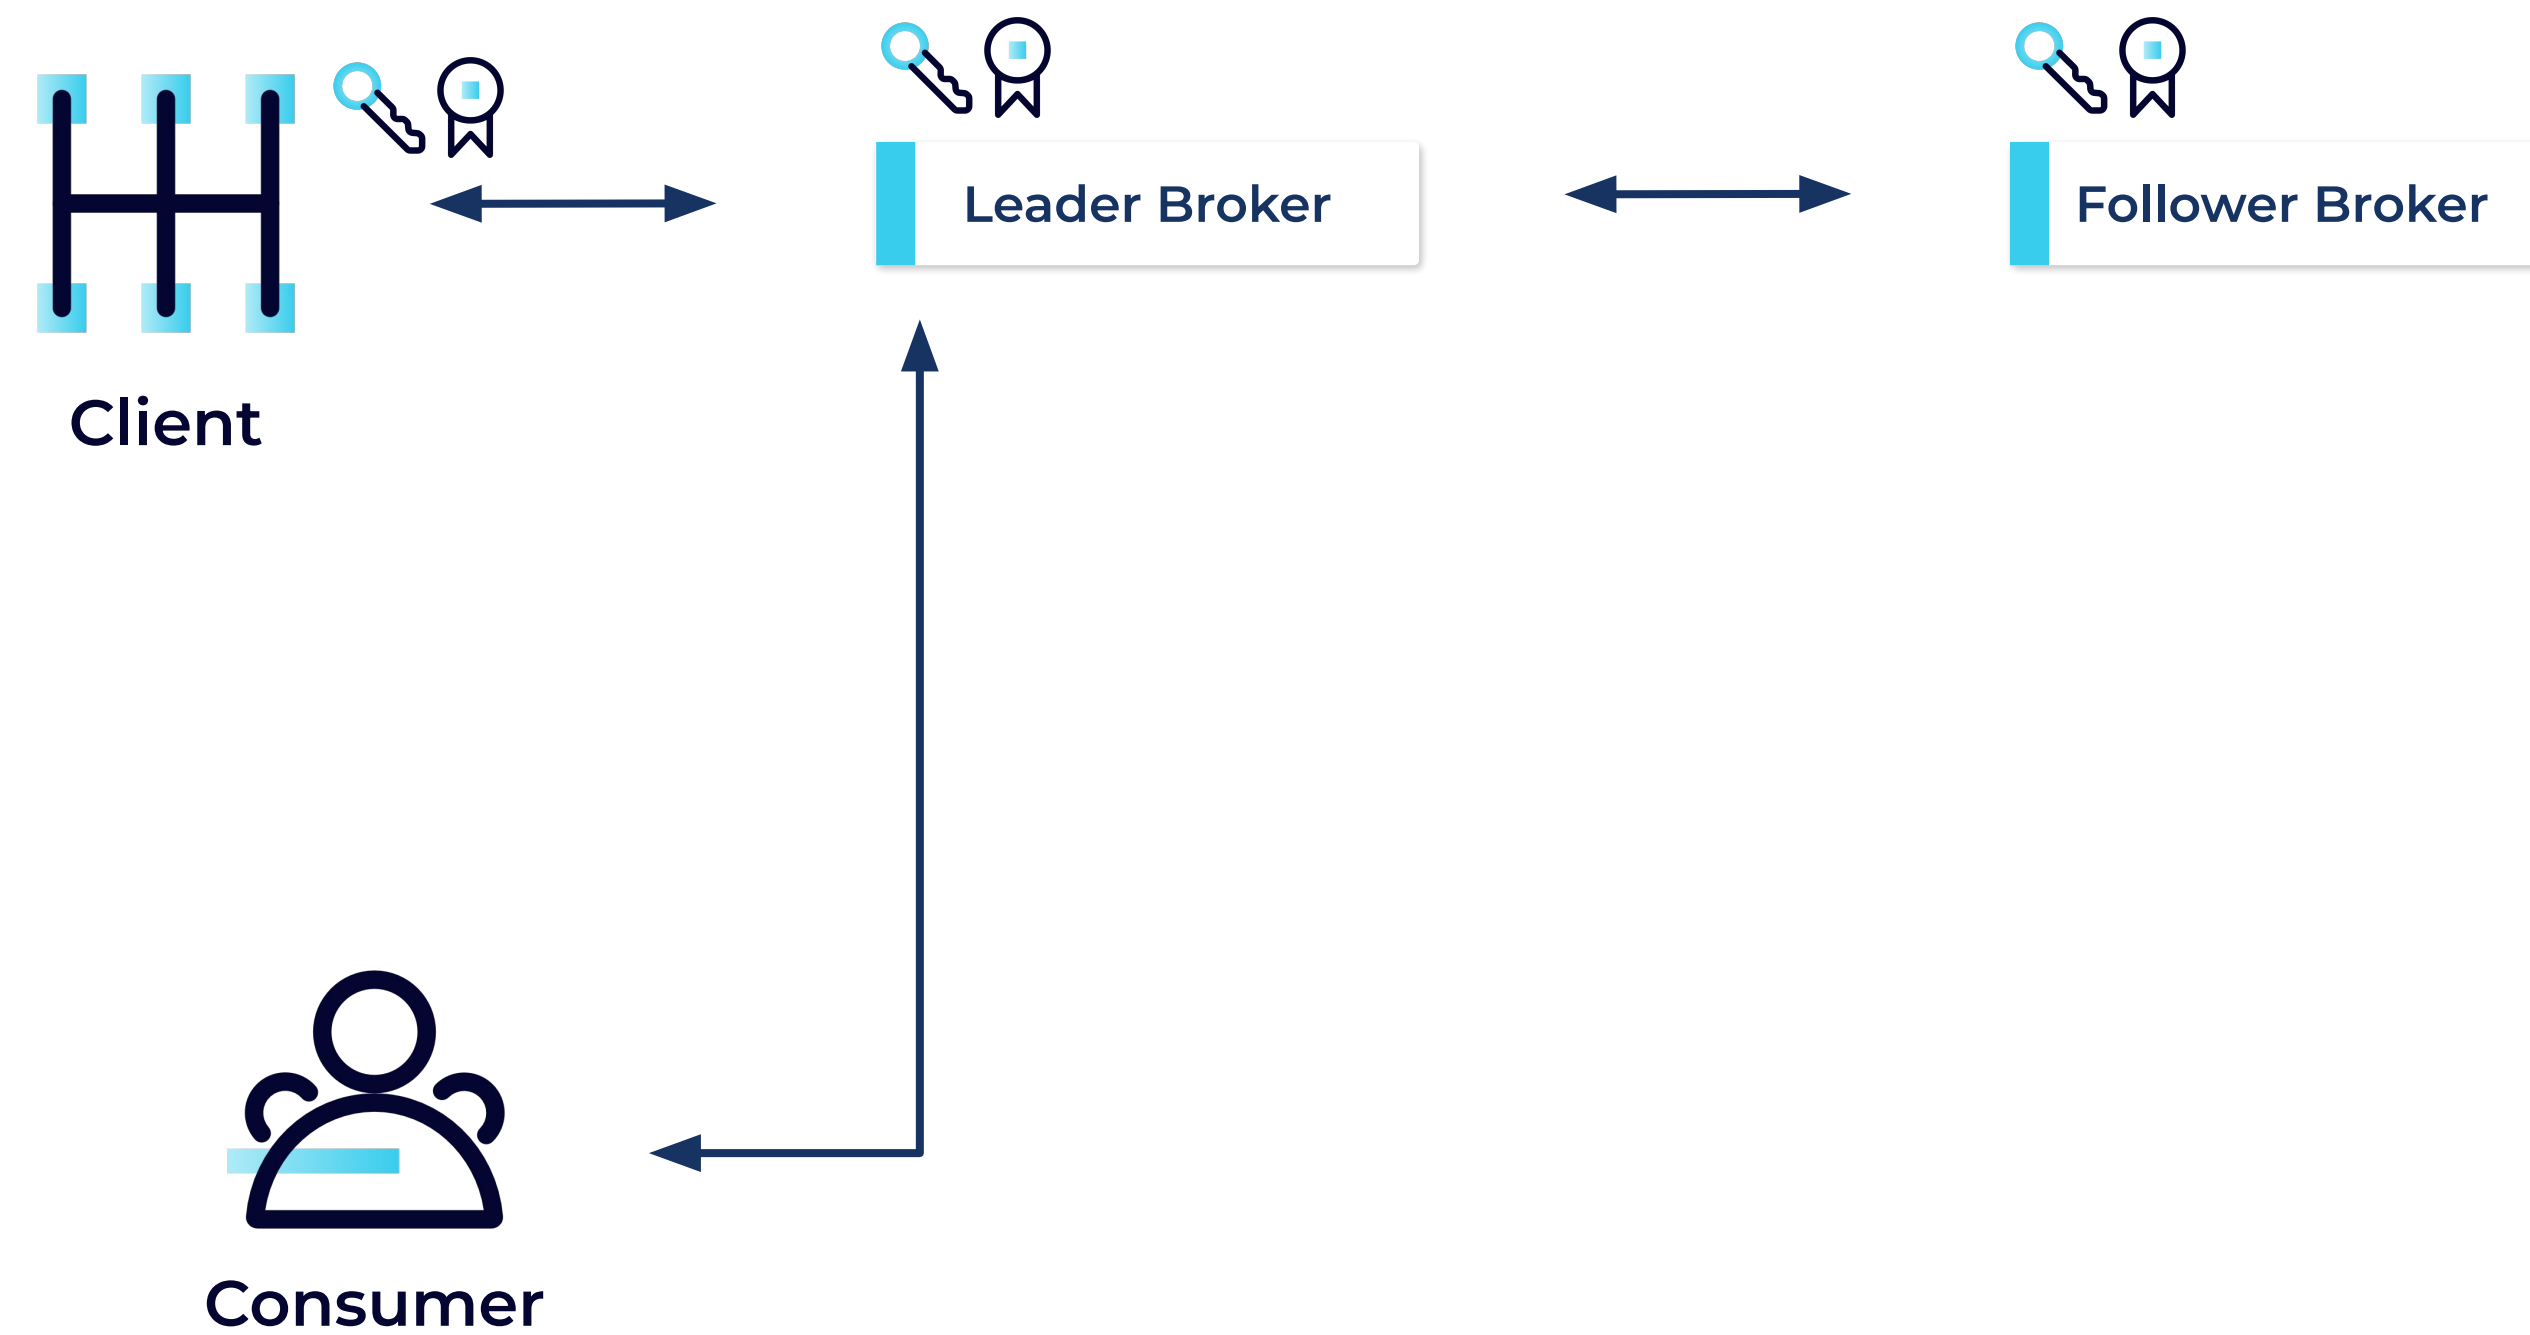
Broker configuration

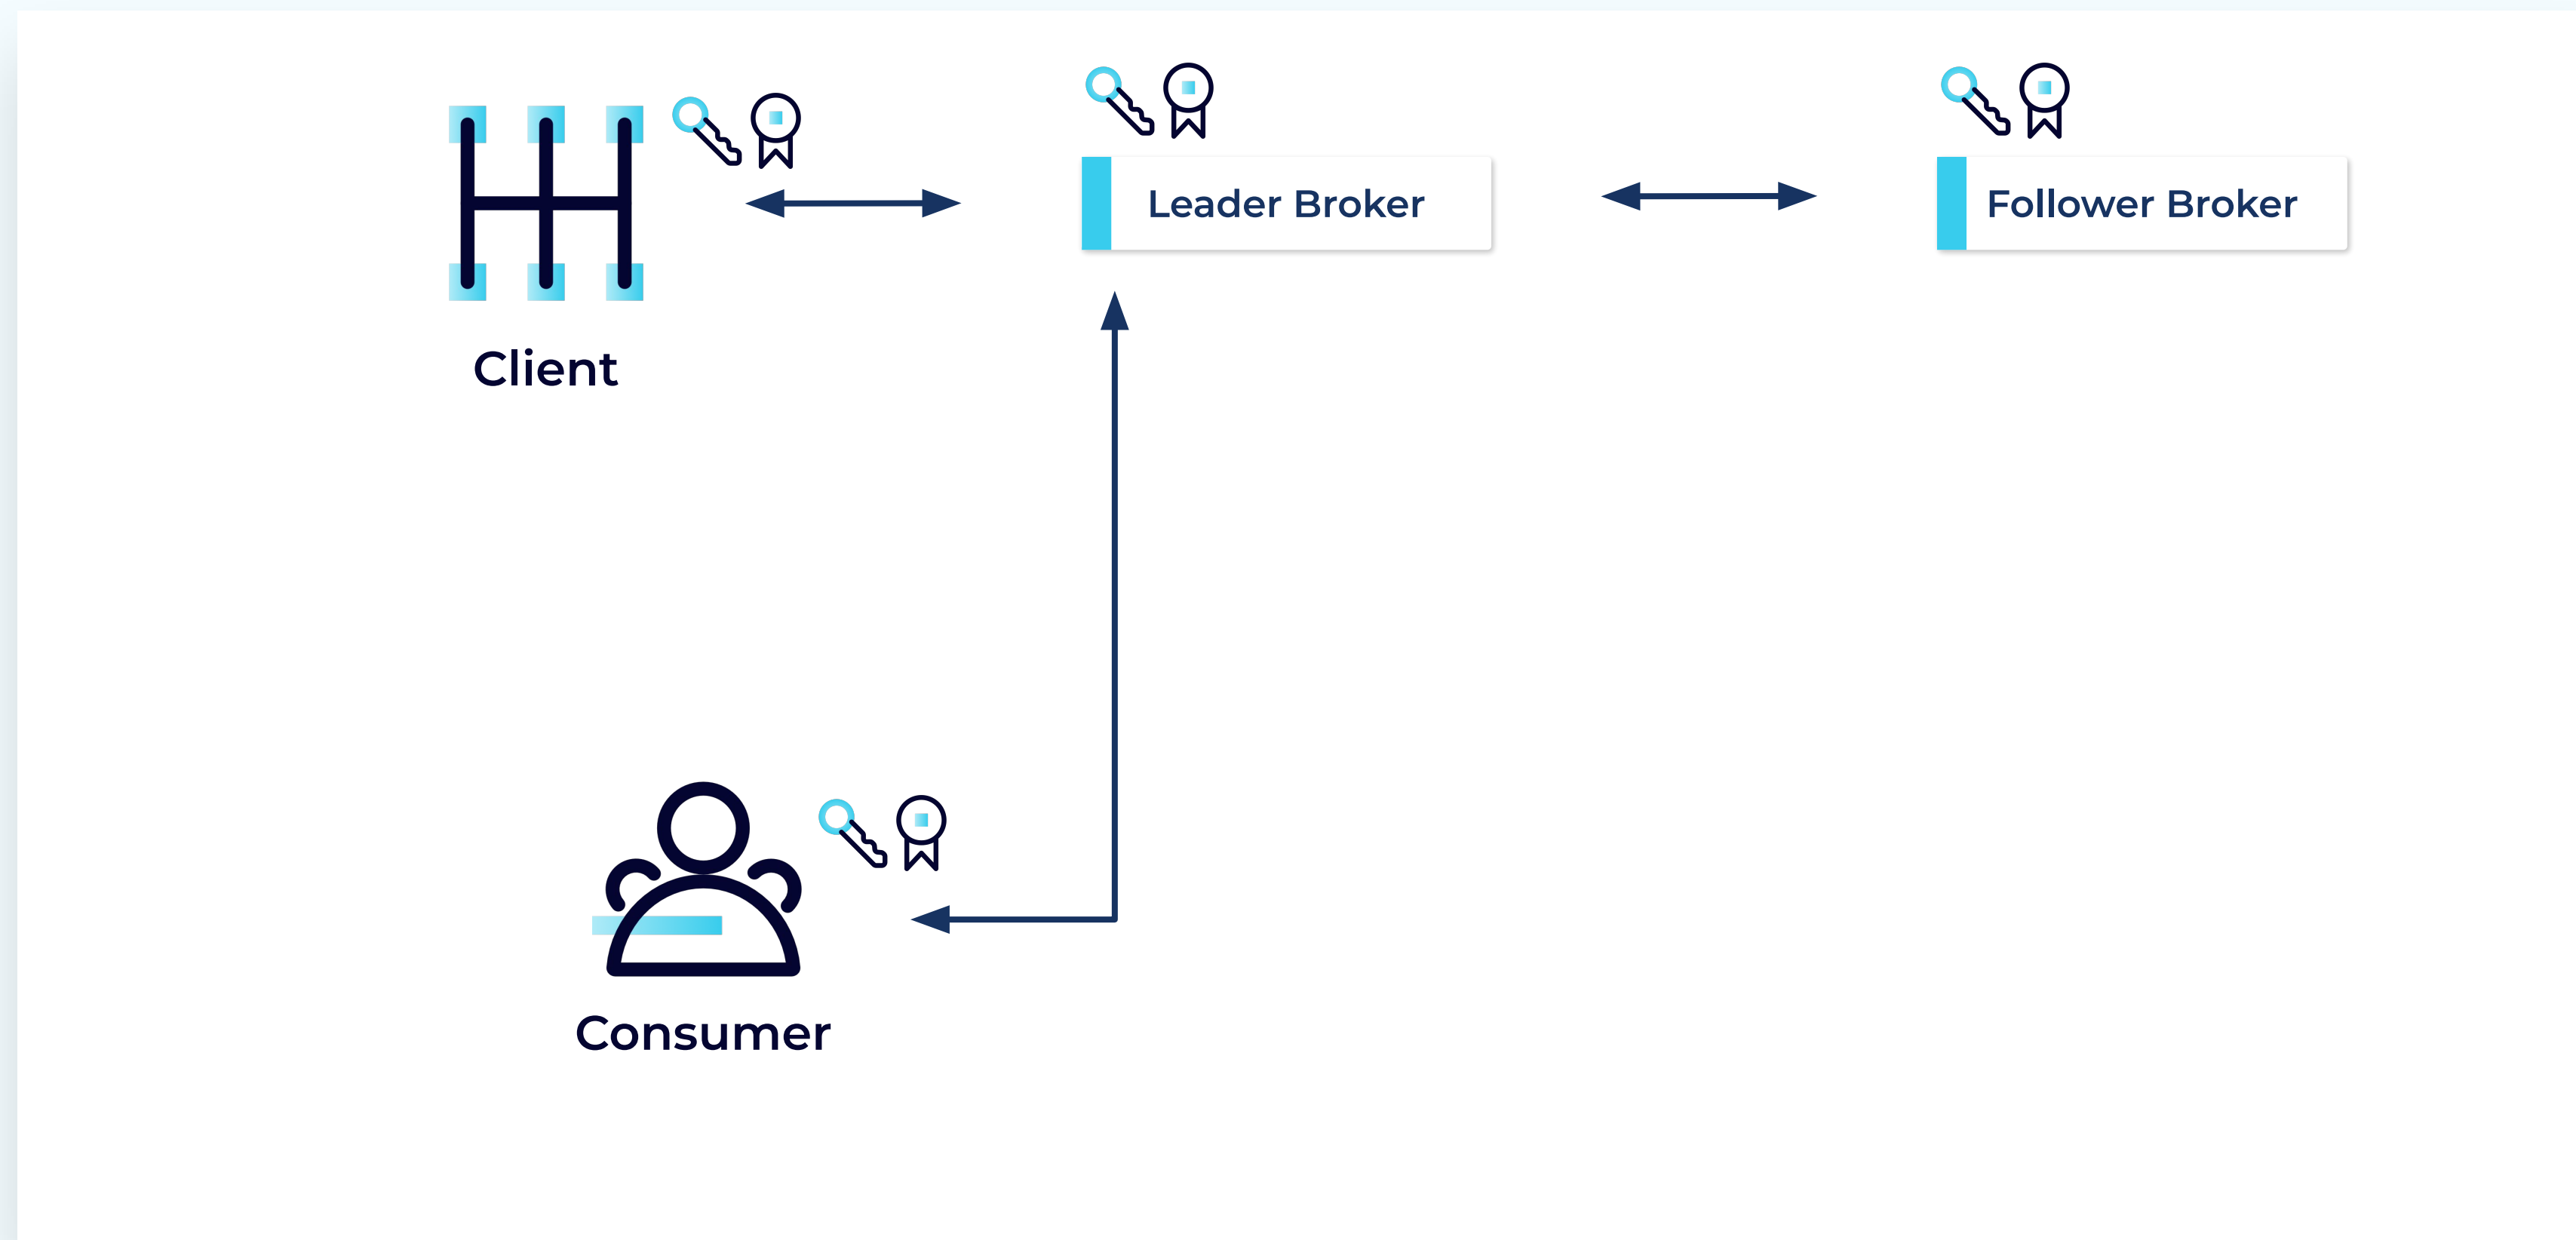
ssl.client.auth=required











Inter-broker TLS Communication


`security.inter.broker.protocol=SSL`

Encrypting data at rest


- Whole disk or volume encryption
 - Filesystem ACLs to restrict access
- Platform capabilities
 - E.g. Encrypted AWS EBS volumes
- Appliances
 - E.g. Vormetric, Gemalto



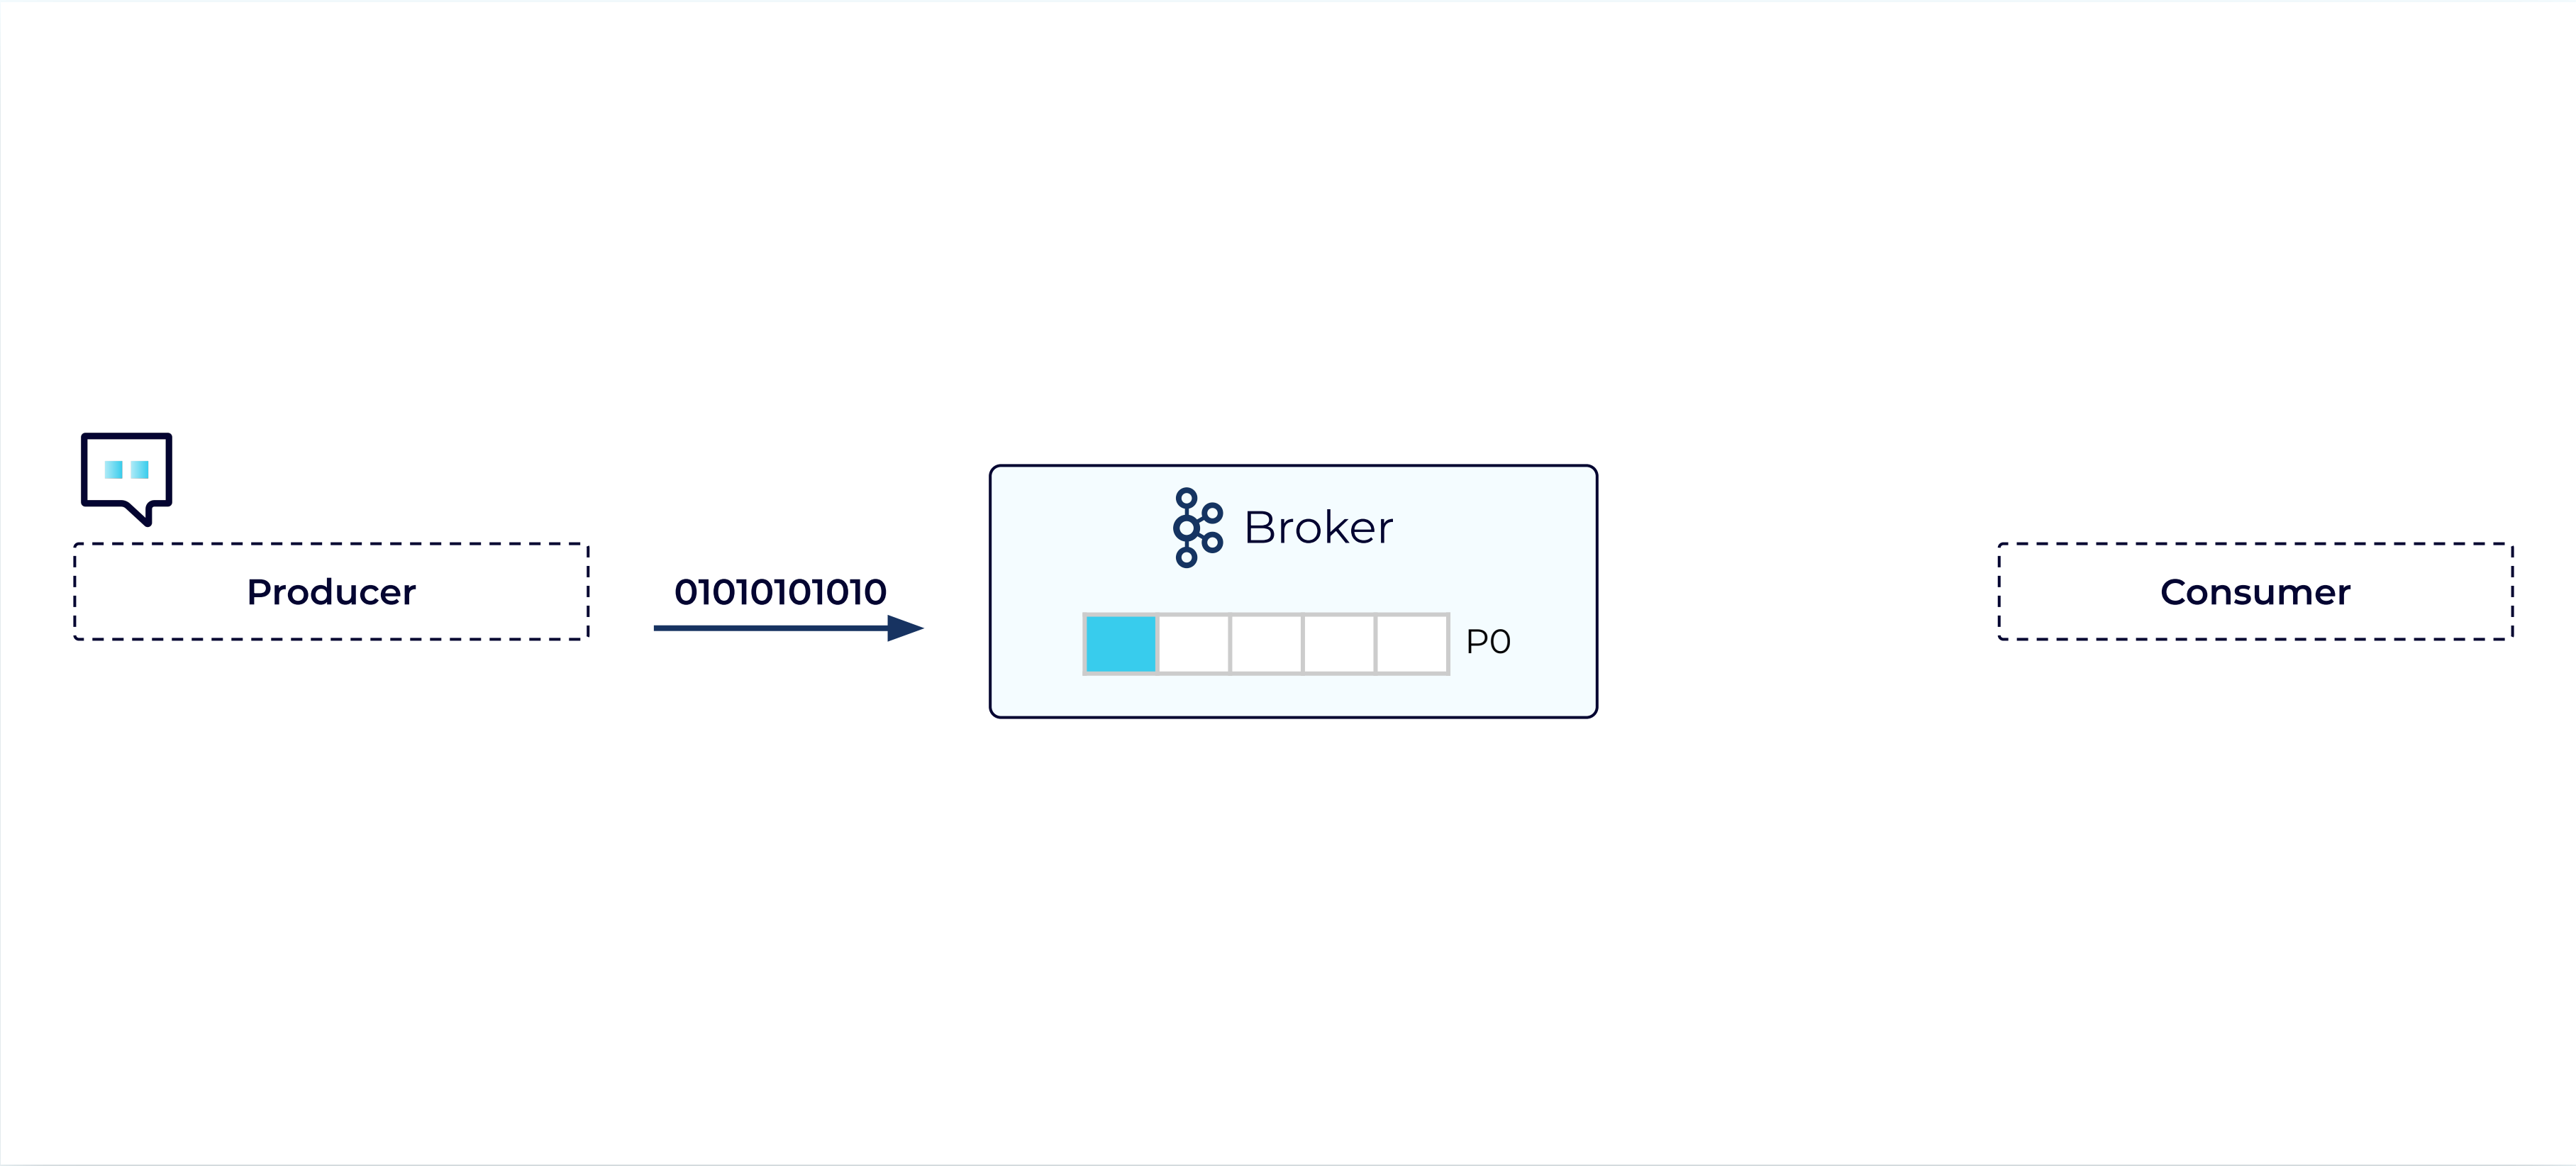
End-to-End Encryption



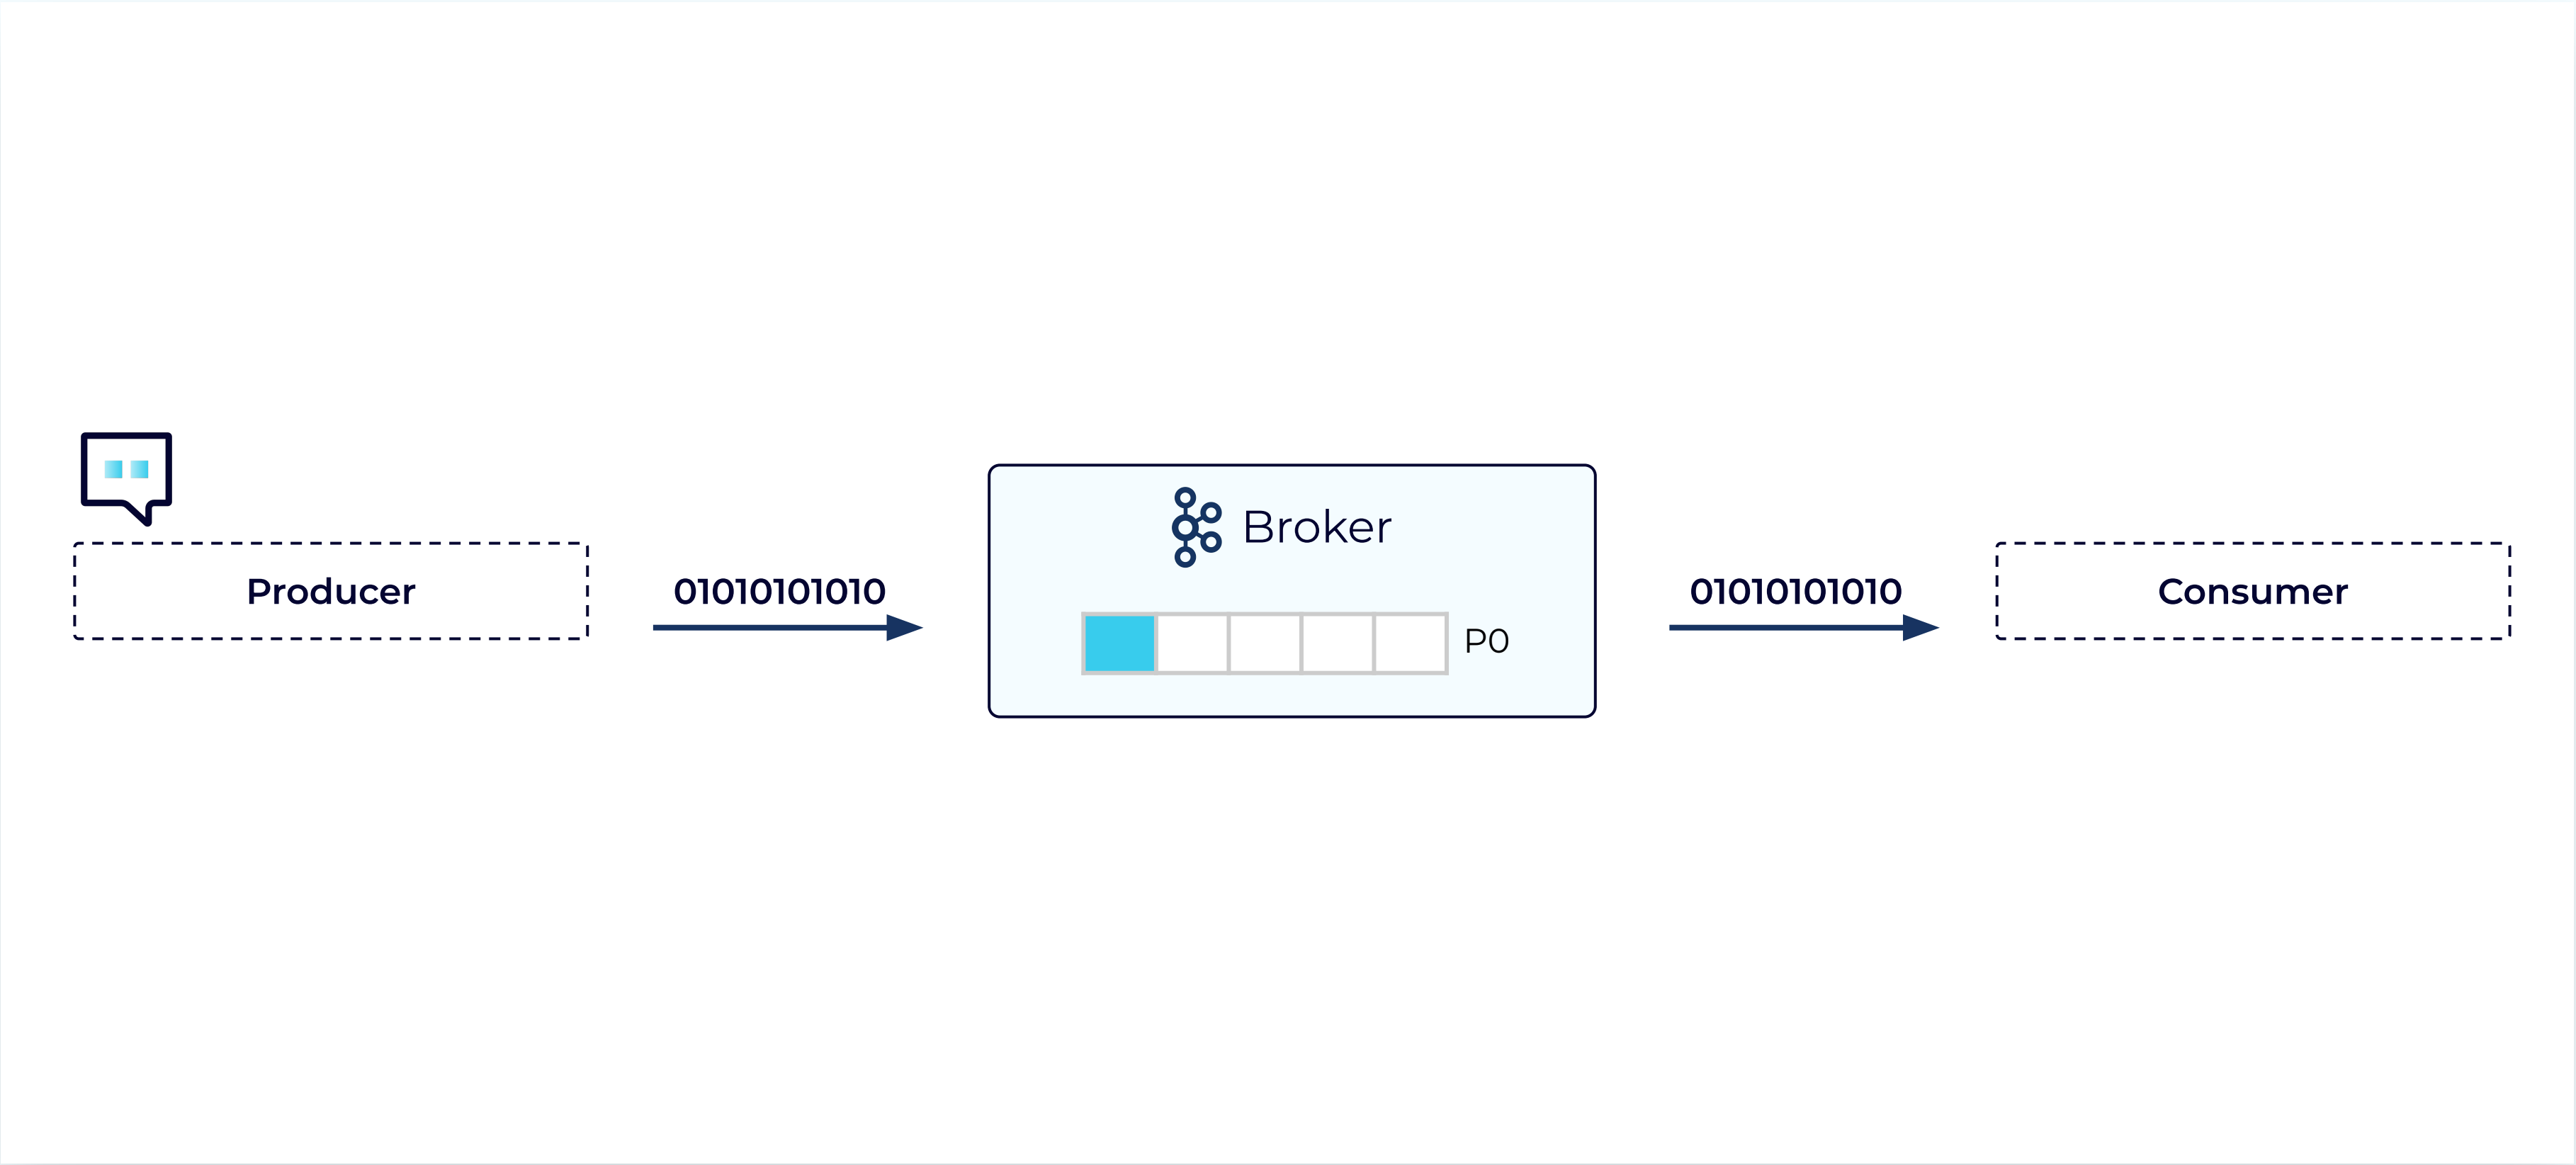
End-to-End Encryption



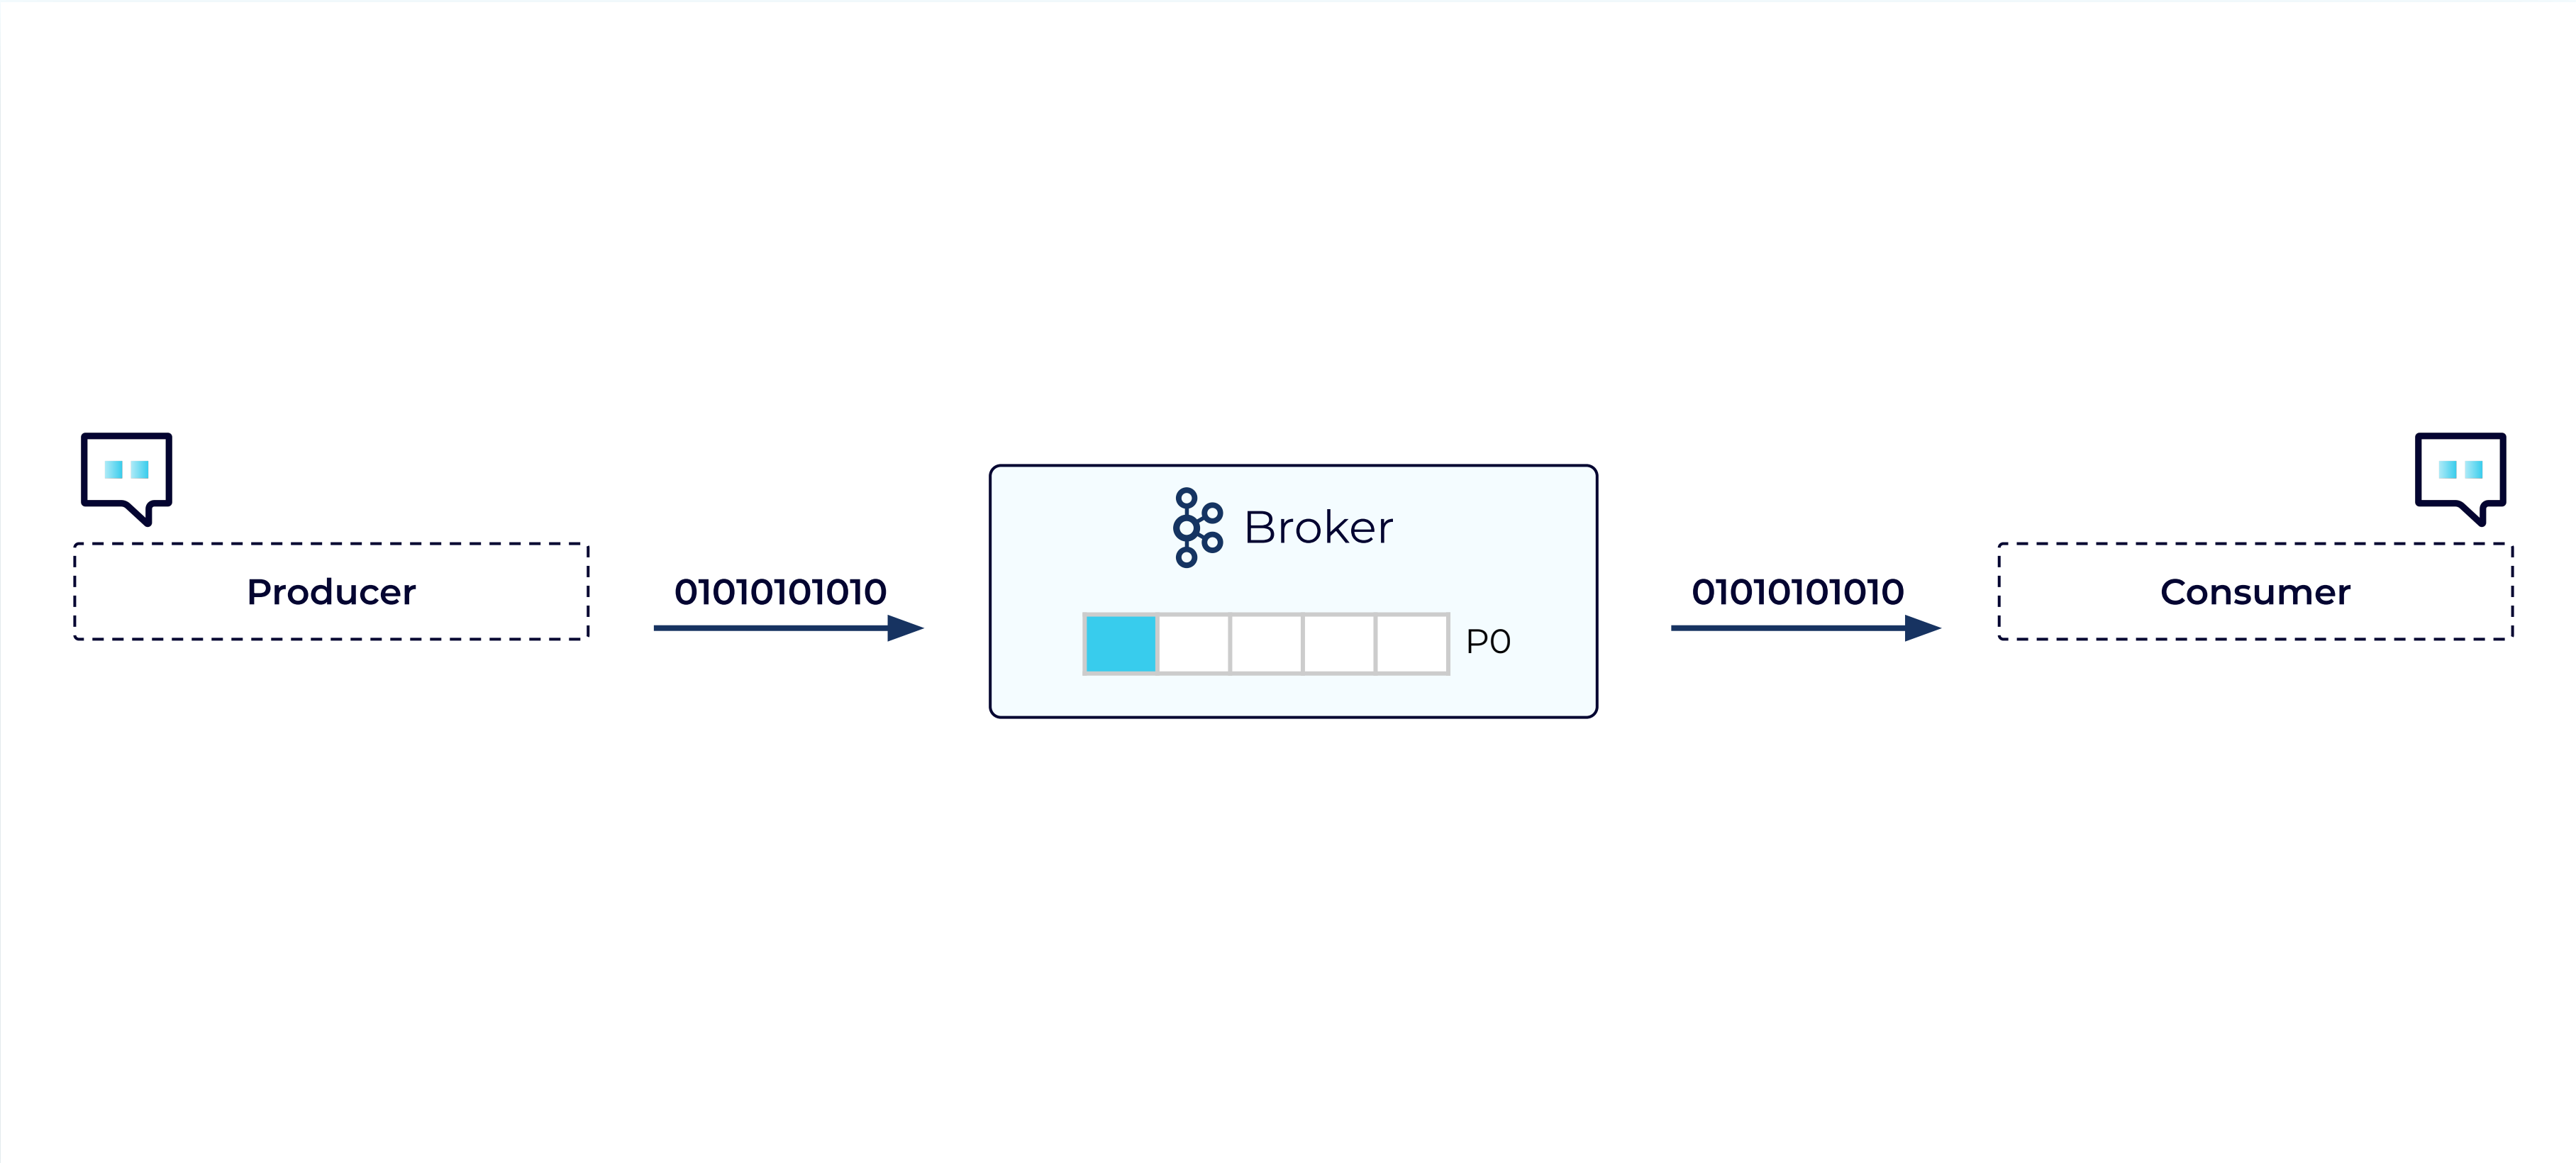
End-to-End Encryption



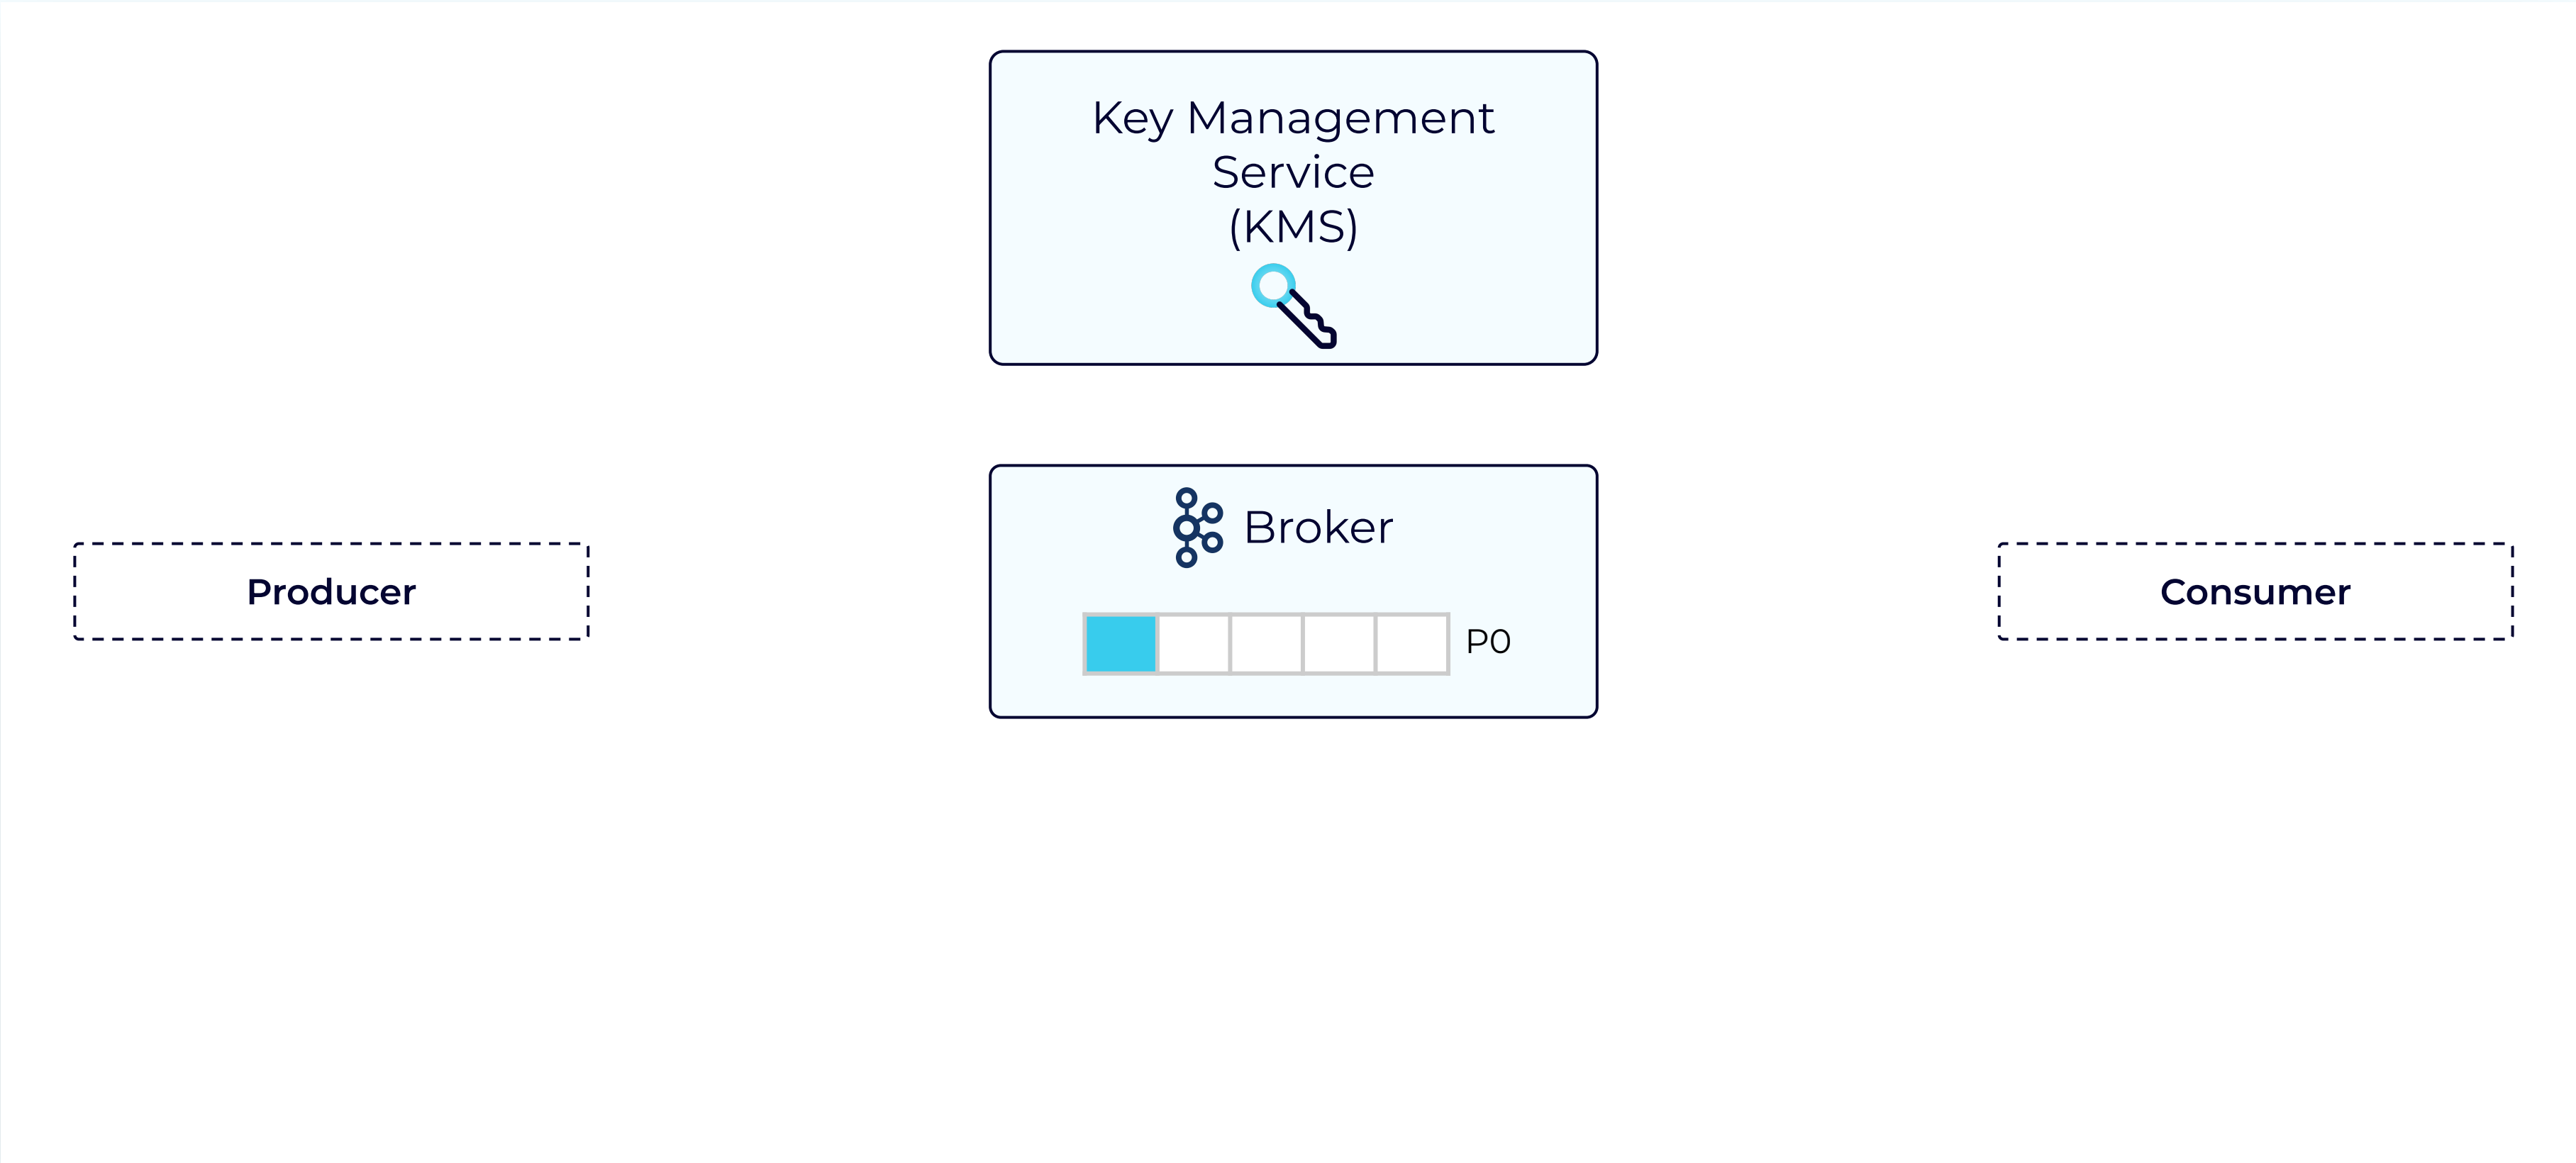
End-to-End Encryption



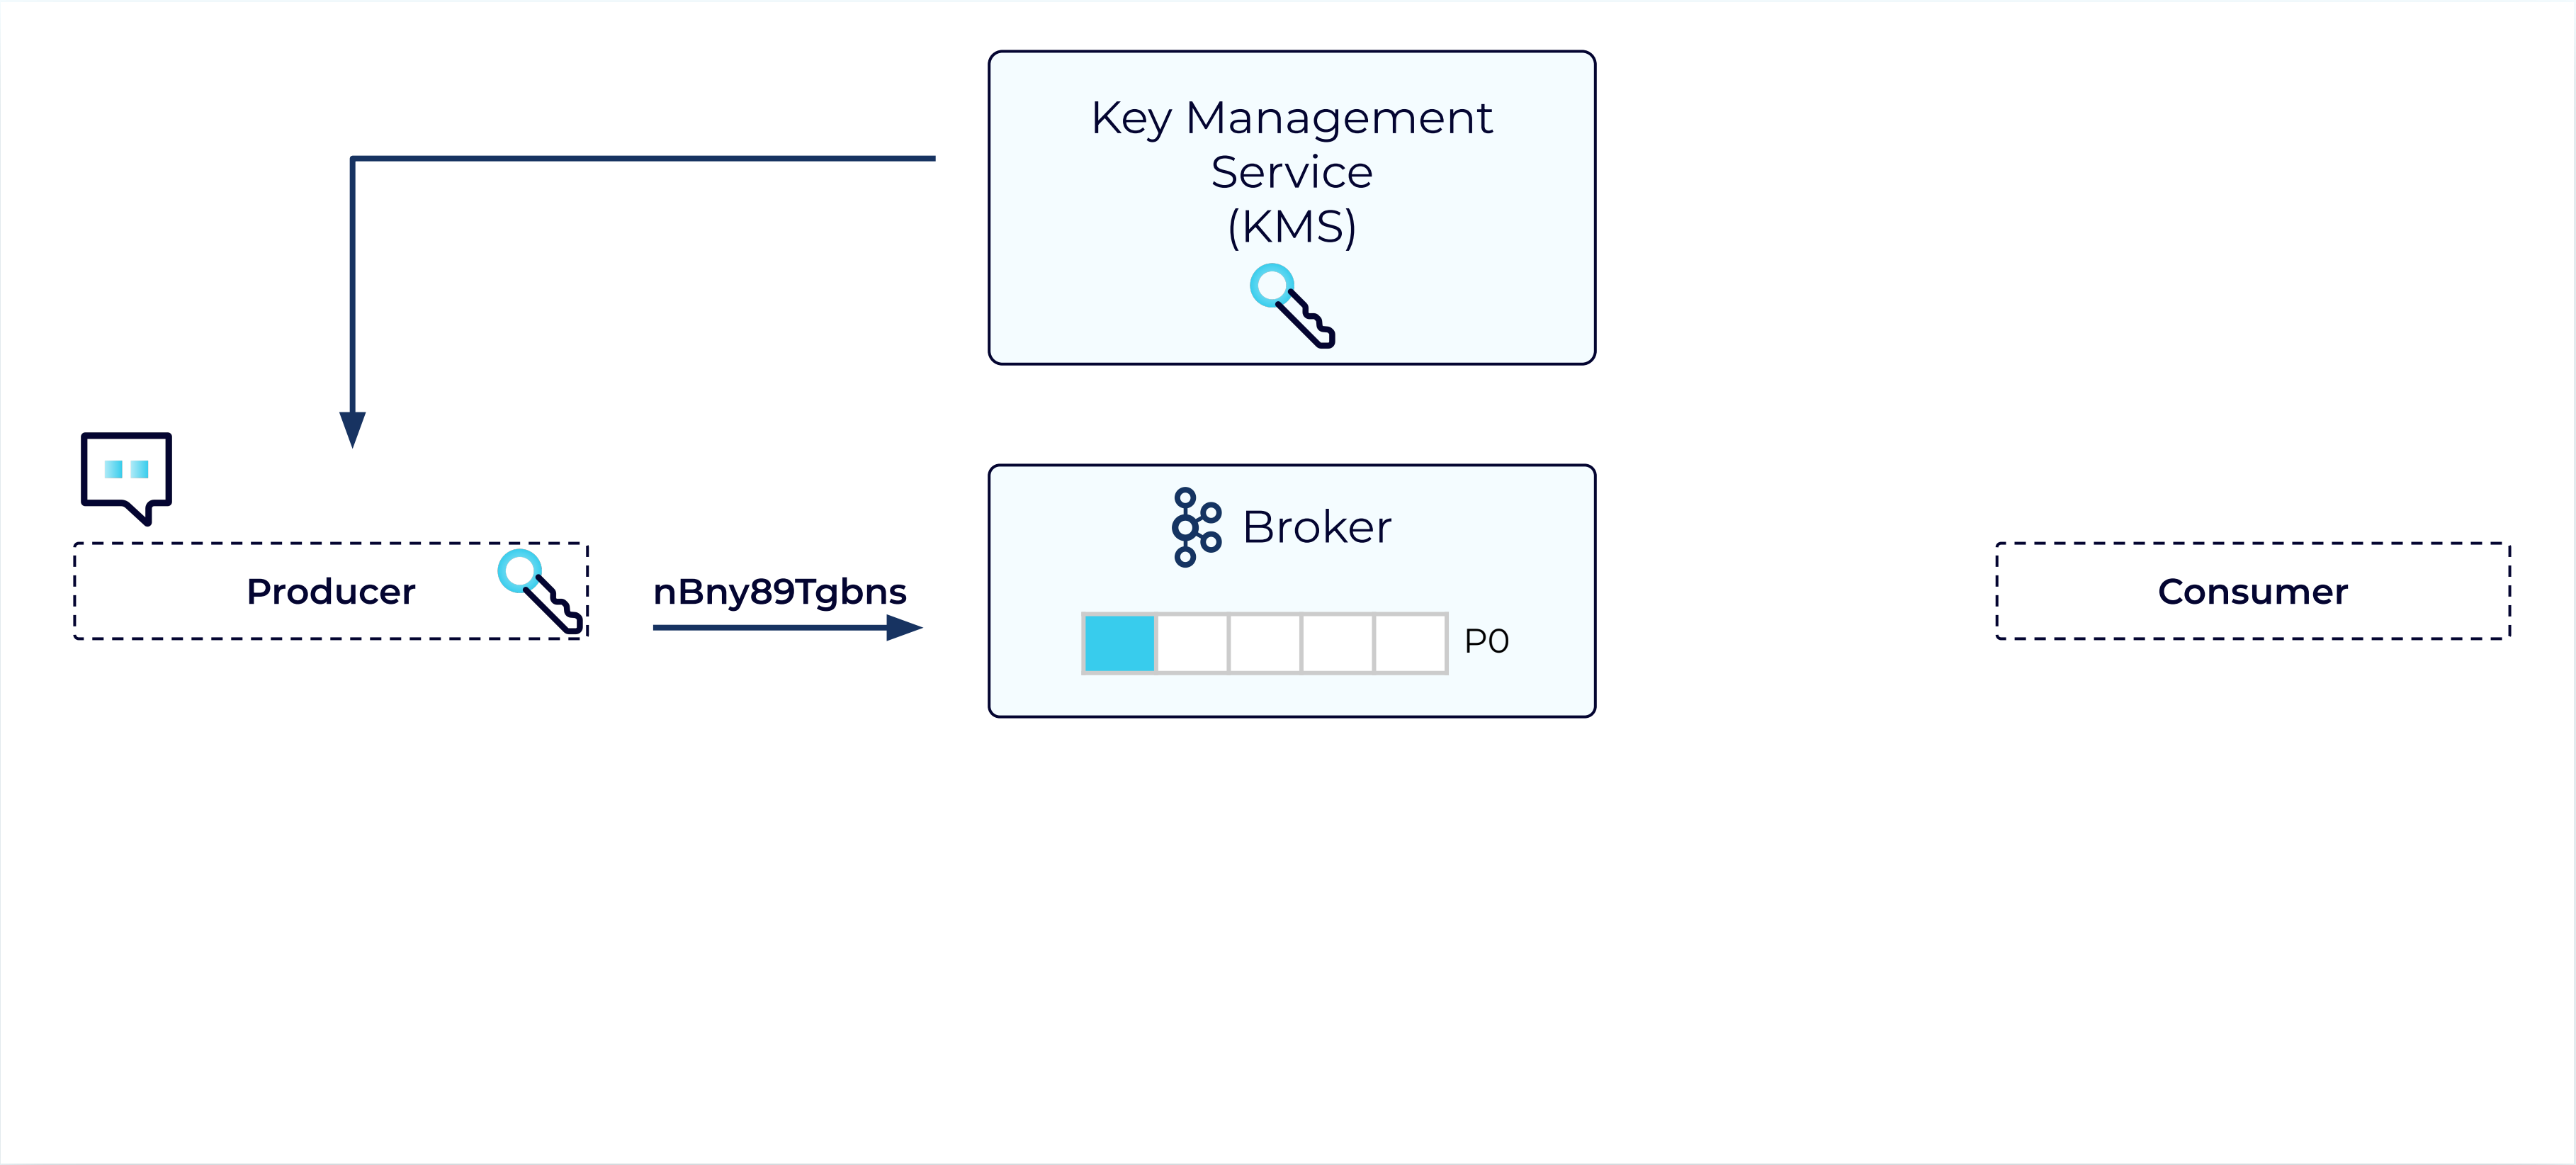
End-to-End Encryption



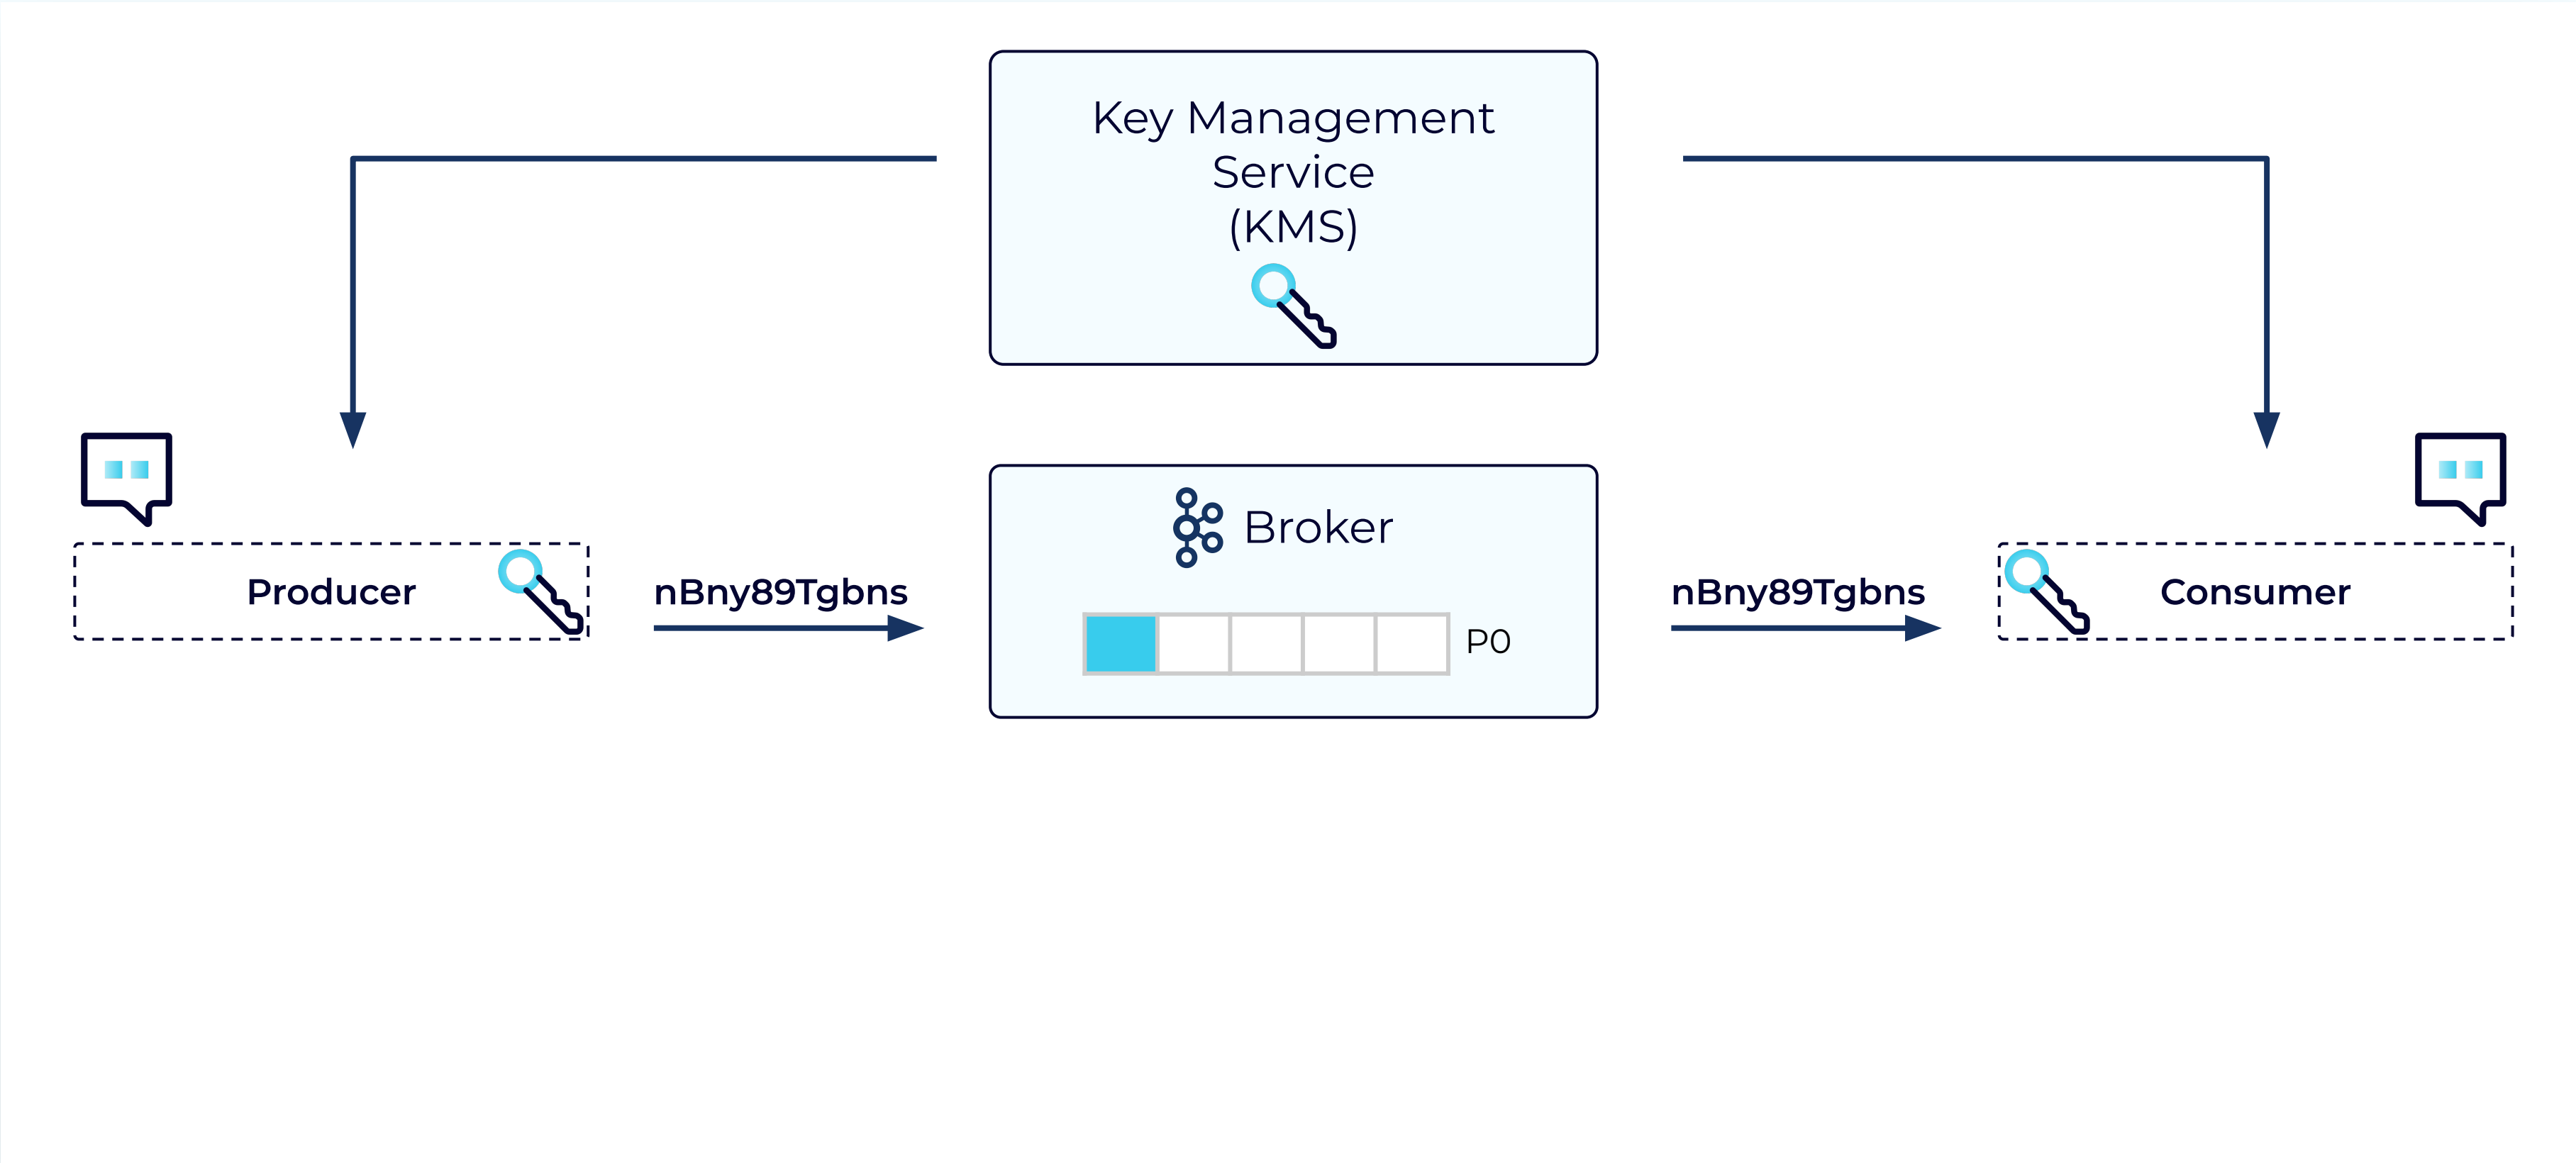
End-to-End Encryption



End-to-End Encryption



End-to-End Encryption





Securing ZooKeeper

Securing ZooKeeper

Set `zookeeper.set.acl` to `true`

SSL client authentication

SASL

Securing ZooKeeper

Set `zookeeper.set.acl` to `true`

SSL client authentication

- Each broker and CLI tool must use same Distinguished Name for *authorization*

SASL

Securing ZooKeeper

Set `zookeeper.set.acl` to `true`

SSL client authentication

- Each broker and CLI tool must use same Distinguished Name for *authorization*
- Use wildcard certificates OR Subject Alternative Name with list of broker hostnames

SASL

Securing ZooKeeper

Set `zookeeper.set.acl` to `true`

SSL client authentication

- Each broker and CLI tool must use same Distinguished Name for *authorization*
- Use wildcard certificates OR Subject Alternative Name with list of broker hostnames

SASL

- Integrate with Kerberos

Securing ZooKeeper

Set `zookeeper.set.acl` to `true`

SSL client authentication

- Each broker and CLI tool must use same Distinguished Name for *authorization*
- Use wildcard certificates OR Subject Alternative Name with list of broker hostnames

SASL

- Integrate with Kerberos
- Use TLS encryption
 - `ssl.clientAuth=none` in the ZooKeeper configuration

Securing ZooKeeper

Set `zookeeper.set.acl` to `true`

SSL client authentication

- Each broker and CLI tool must use same Distinguished Name for *authorization*
- Use wildcard certificates OR Subject Alternative Name with list of broker hostnames

SASL

- Integrate with Kerberos
- Use TLS encryption
 - `ssl.clientAuth=none` in the ZooKeeper configuration
- Configure each broker with same Kerberos principal

Securing ZooKeeper



SSL + SASL


- **Able to use either identity**
- **No need to use same distinguished name**
- **Ability to use hostnames in the distinguished name**



Audit Logs

Audit Logs

1. Insight – Logging every attempted operation



Audit Logs

1. **Insight** – Logging every attempted operation
2. **Security** – Monitoring and validating operations



Audit Logs

1. **Insight** – Logging every attempted operation
2. **Security** – Monitoring and validating operations
3. **Impact** – Debugging client interactions



Audit Logs

1. **Insight** – Logging every attempted operation
2. **Security** – Monitoring and validating operations
3. **Impact** – Debugging client interactions
4. **Compliance** – Generating audit reports



Audit Logs



- kafka.authorizer.logger
 - Used for authorization logs
- kafka.request.logger
 - Used for request logs

Log4j authorization logger



log4j.properties

`log4j.logger.kafka.authorizer.logger=DEBUG`

output

DEBUG Principal = User:Alice is Allowed Operation = Write from host = 127.0.0.1 on
resource = Topic:LITERAL:customerOrders for request = Produce with resourceRefCount =
1 (kafka.authorizer.logger)

INFO Principal = User:Mallory is Denied Operation = Describe from host = 10.0.0.13 on
resource = Topic:LITERAL:customerOrders for request = Metadata with resourceRefCount
= 1 (kafka.authorizer.logger)

Log4j request logger



log4j.properties

```
log4j.logger.kafka.request.logger=DEBUG
```

output

```
DEBUG Completed request:RequestHeader(apiKey=PRODUCE, apiVersion=8,
clientId=producer-1, correlationId=6) --
{acks=-1,timeout=30000,partitionSizes=[customerOrders-0=15514]},response:
{responses=[{topic=customerOrders,partition_responses=[{partition=0,error_code=0
,base_offset=13,log_append_time=-1,log_start_offset=0,record_errors=[],error_mes
sage=null}]]],throttle_time_ms=0} from connection
127.0.0.1:9094-127.0.0.1:61040-0;totalTime:2.42,requestQueueTime:0.112,local-Time:2.15,re
moteTime:0.0,throttleTime:0,responseQueueTime:0.04,sendTime:
0.118,securityProtocol:SASL_SSL,principal:User:Alice,listener:SASL_SSL,clientInf
ormation:ClientInformation(softwareName=apache-kafka-java,
softwareVersion=2.7.0-SNAPSHOT) (kafka.request.logger)
```

Considerations



- Ensure you have sufficient disk space for each broker

Considerations



- Ensure you have sufficient disk space for each broker
- Set a retention policy for logs

Considerations



- Ensure you have sufficient disk space for each broker
- Set a retention policy for logs
- Capture and consolidate logs to view holistic issues

Considerations



- Ensure you have sufficient disk space for each broker
- Set a retention policy for logs
- Capture and consolidate logs to view holistic issues
- Use a visualization tool (e.g., ELK stack, Elasticsearch, Logstash, and Kibana)



Security Recommendations

Education

Start with security in mind

Encrypt the filesystem

Secure data in transit

Set up a system for administering ACLs

Rotate your keys

Dynamically update certificates

Enable Reauthentication

Protect ZooKeeper

Set up and monitor audit logs

Play, tinker, break things



***Your Apache Kafka
journey begins here***

developer.confluent.io